

# The Processor Architecture Based on Hennessey/Patterson RISC-V edition

# Introduction

2

- CPU performance factors
  - ▣ Instruction count
    - Determined by ISA and compiler
  - ▣ CPI and Cycle time
    - Determined by CPU hardware
- We will examine two RISC-V implementations
  - ▣ A simplified version
  - ▣ A more realistic pipelined version
- Simple subset, shows most aspects
  - ▣ Memory reference: ld, sd
  - ▣ Arithmetic/logical: add, sub, and, or
  - ▣ Control transfer: beq

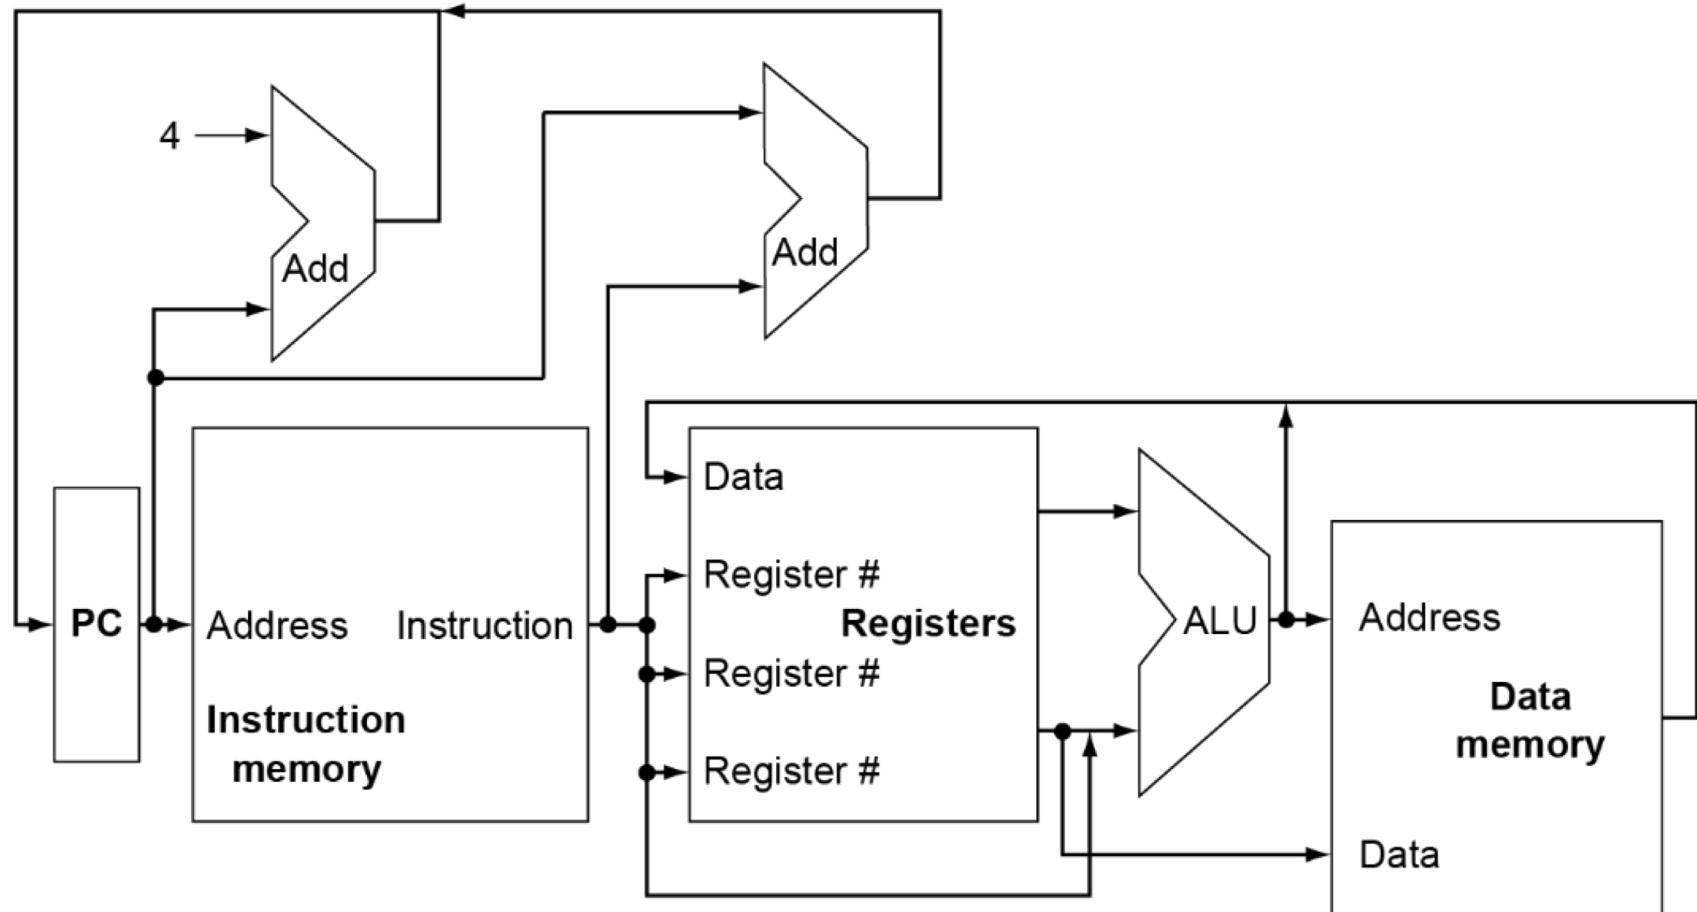
# Instruction Execution

3

- PC → instruction memory, fetch instruction
- Register numbers → register file, read registers
- Depending on instruction class
  - ▣ Use ALU to calculate
    - Arithmetic result
    - Memory address for load/store
    - Branch comparison
  - ▣ Access data memory for load/store
  - ▣ PC ← target address or PC + 4

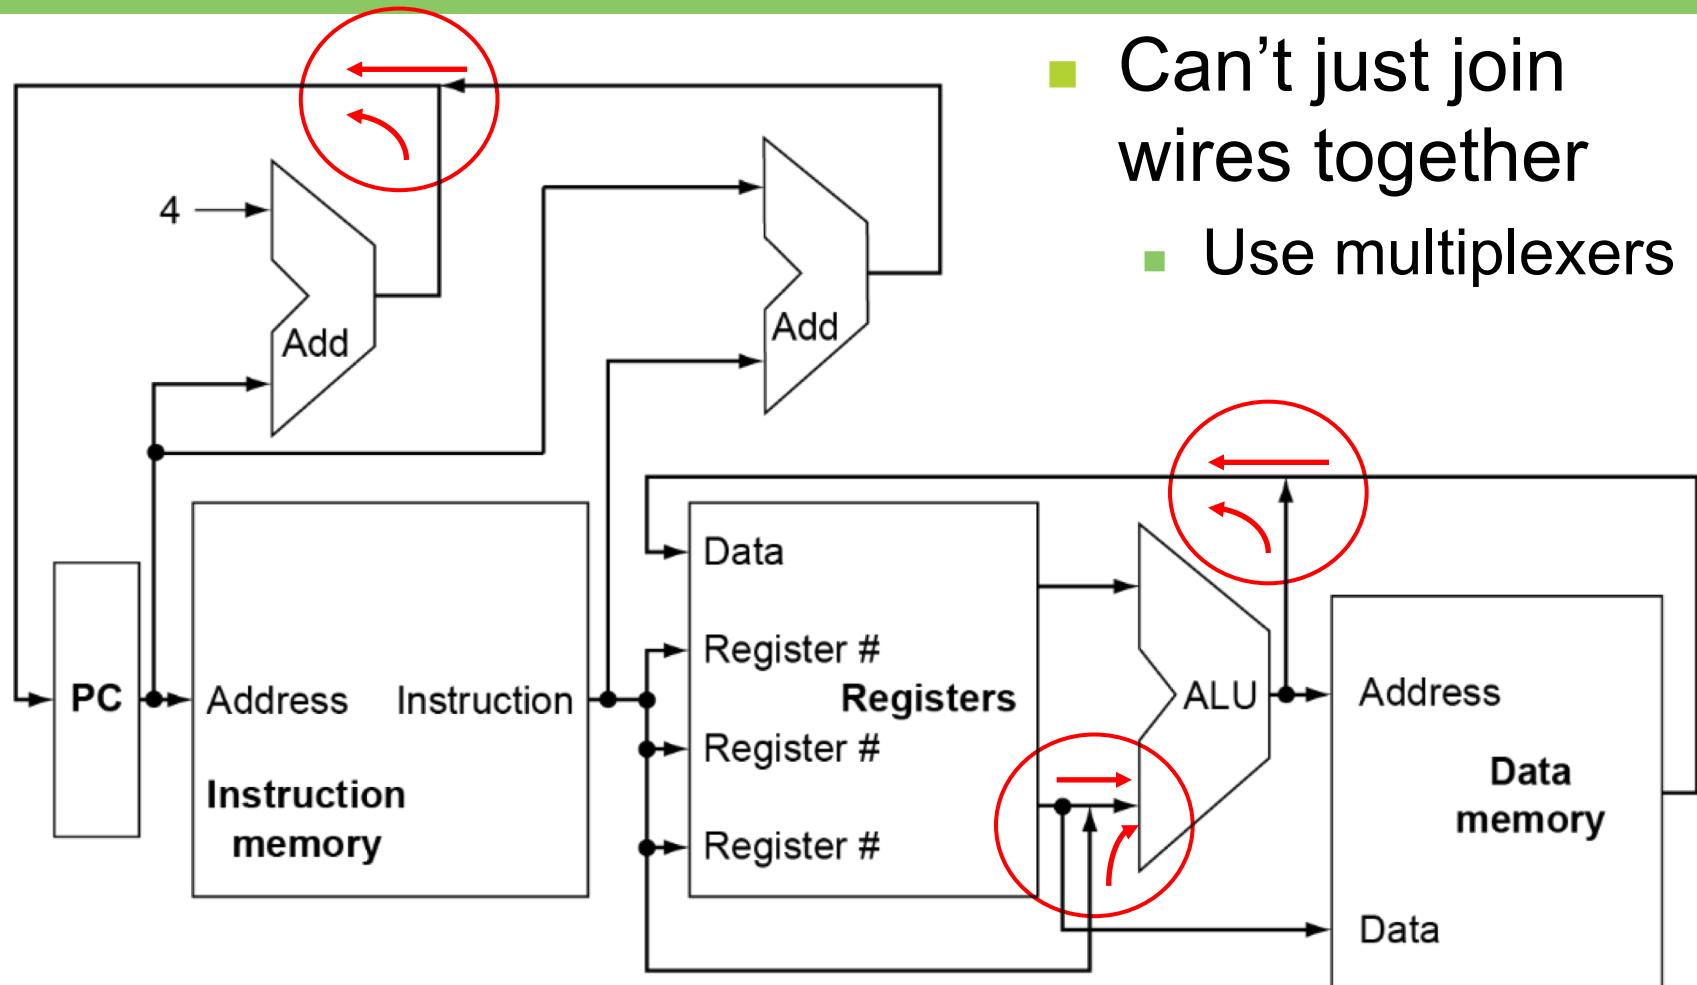
# CPU Overview

4



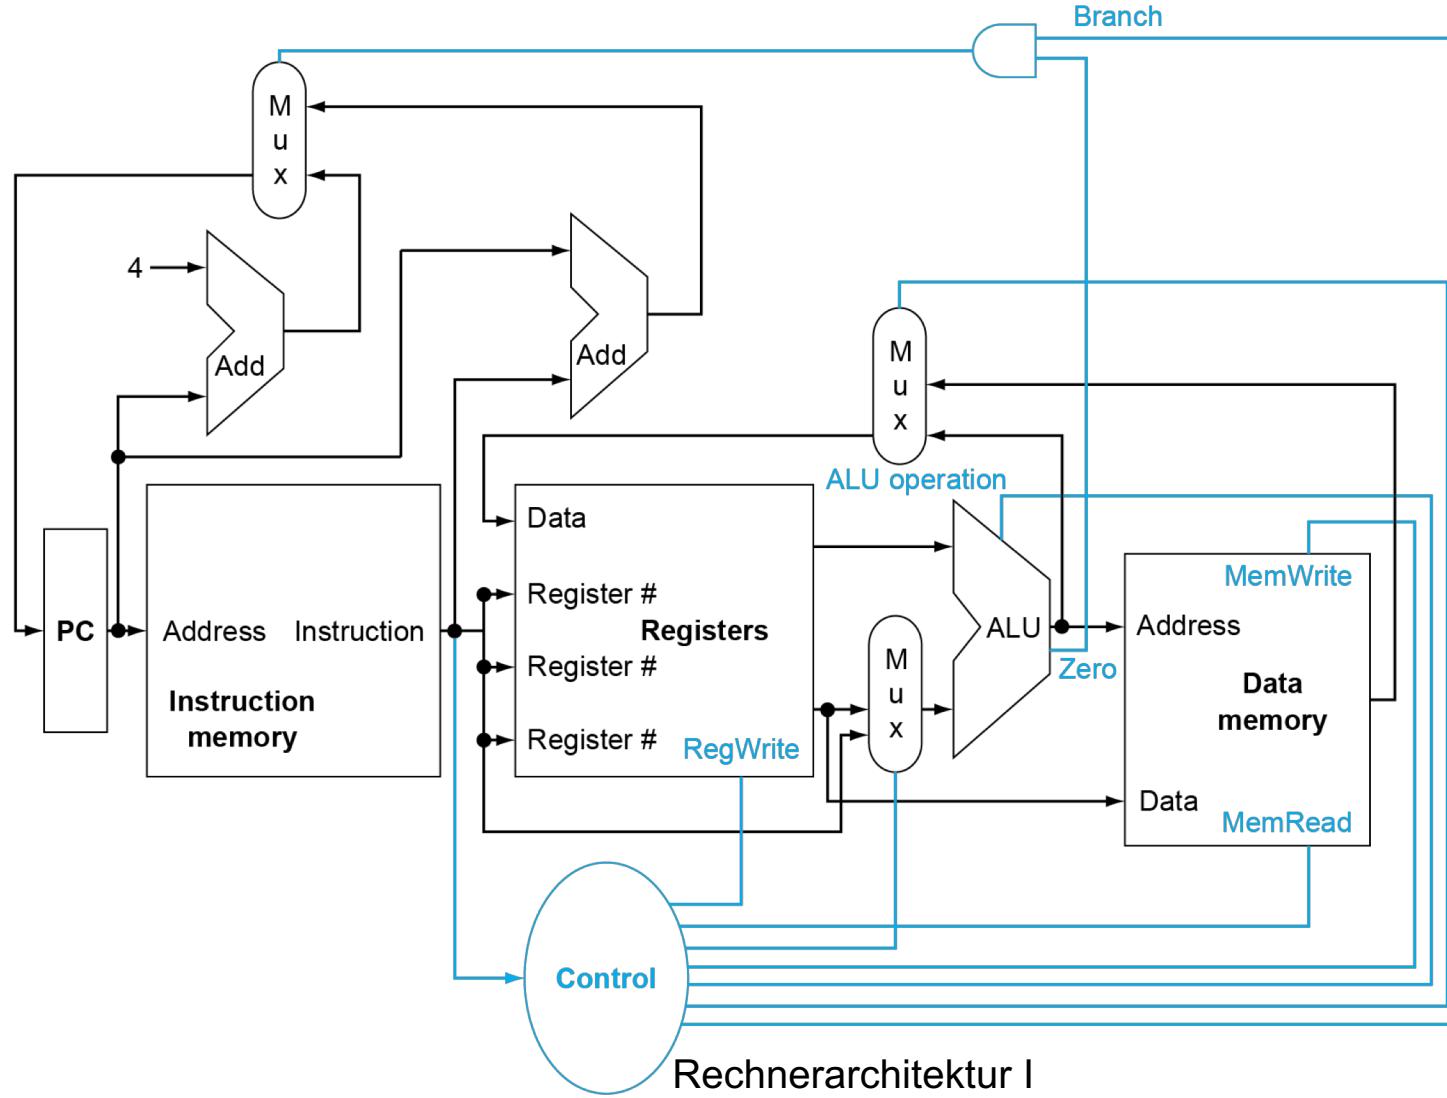
# Multiplexers

5



# Control

6



# Logic Design Basics

7

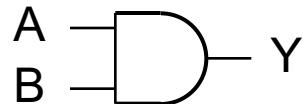
- Information encoded in binary
  - ▣ Low voltage = 0, High voltage = 1
  - ▣ One wire per bit
  - ▣ Multi-bit data encoded on multi-wire buses
- Combinational element
  - ▣ Operate on data
  - ▣ Output is a function of input
- State (sequential) elements
  - ▣ Store information

# Combinational Elements

8

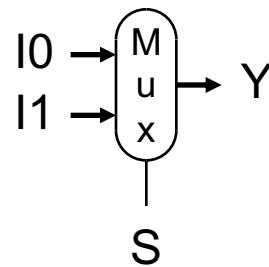
- AND-gate

- $Y = A \& B$



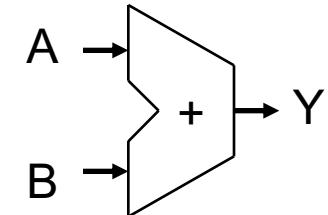
- Multiplexer

- $Y = S ? I_1 : I_0$



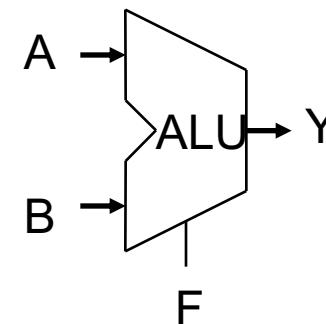
- Adder

- $Y = A + B$



- Arithmetic/Logic Unit

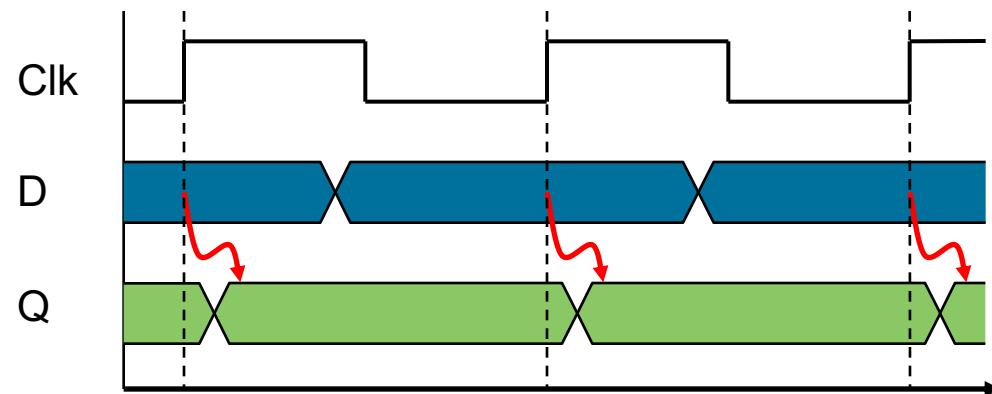
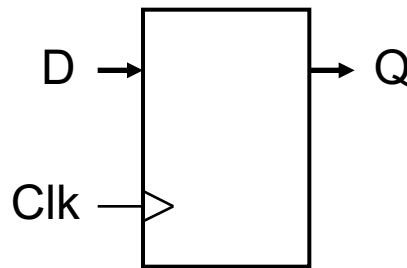
- $Y = F(A, B)$



# Sequential Elements

9

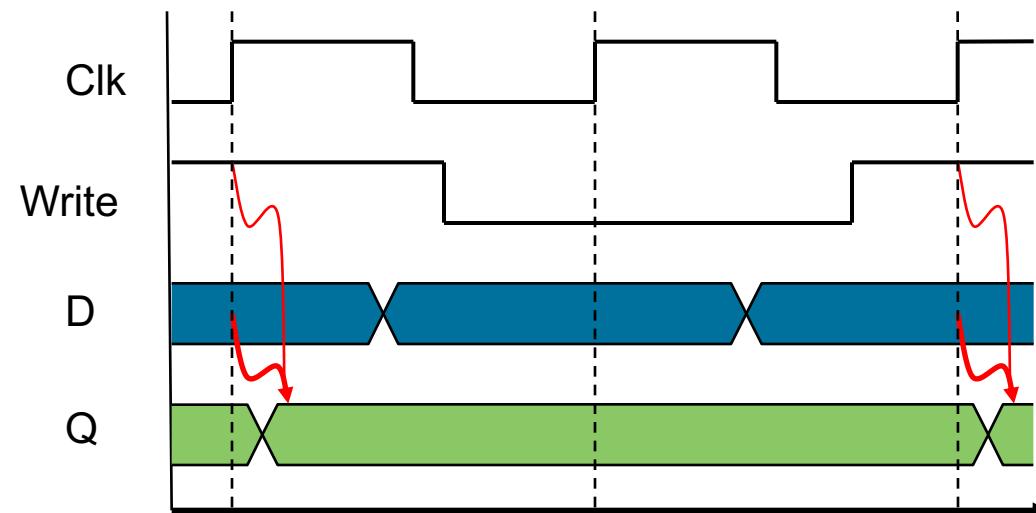
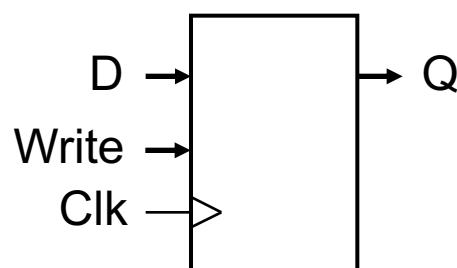
- Register: stores data in a circuit
  - Uses a clock signal to determine when to update the stored value
  - Edge-triggered: update when Clk changes from 0 to 1



# Sequential Elements

10

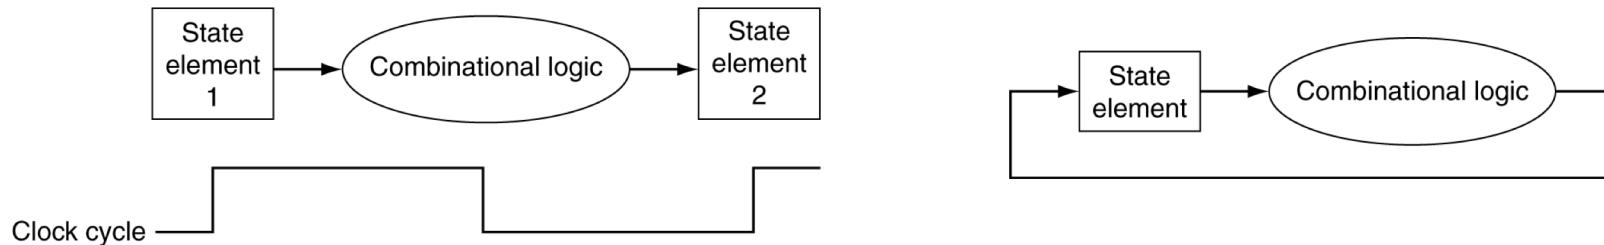
- Register with write control
  - Only updates on clock edge when write control input is 1
  - Used when stored value is required later



# Clocking Methodology

11

- Combinational logic transforms data during clock cycles
  - Between clock edges
  - Input from state elements, output to state element
  - Longest delay determines clock period



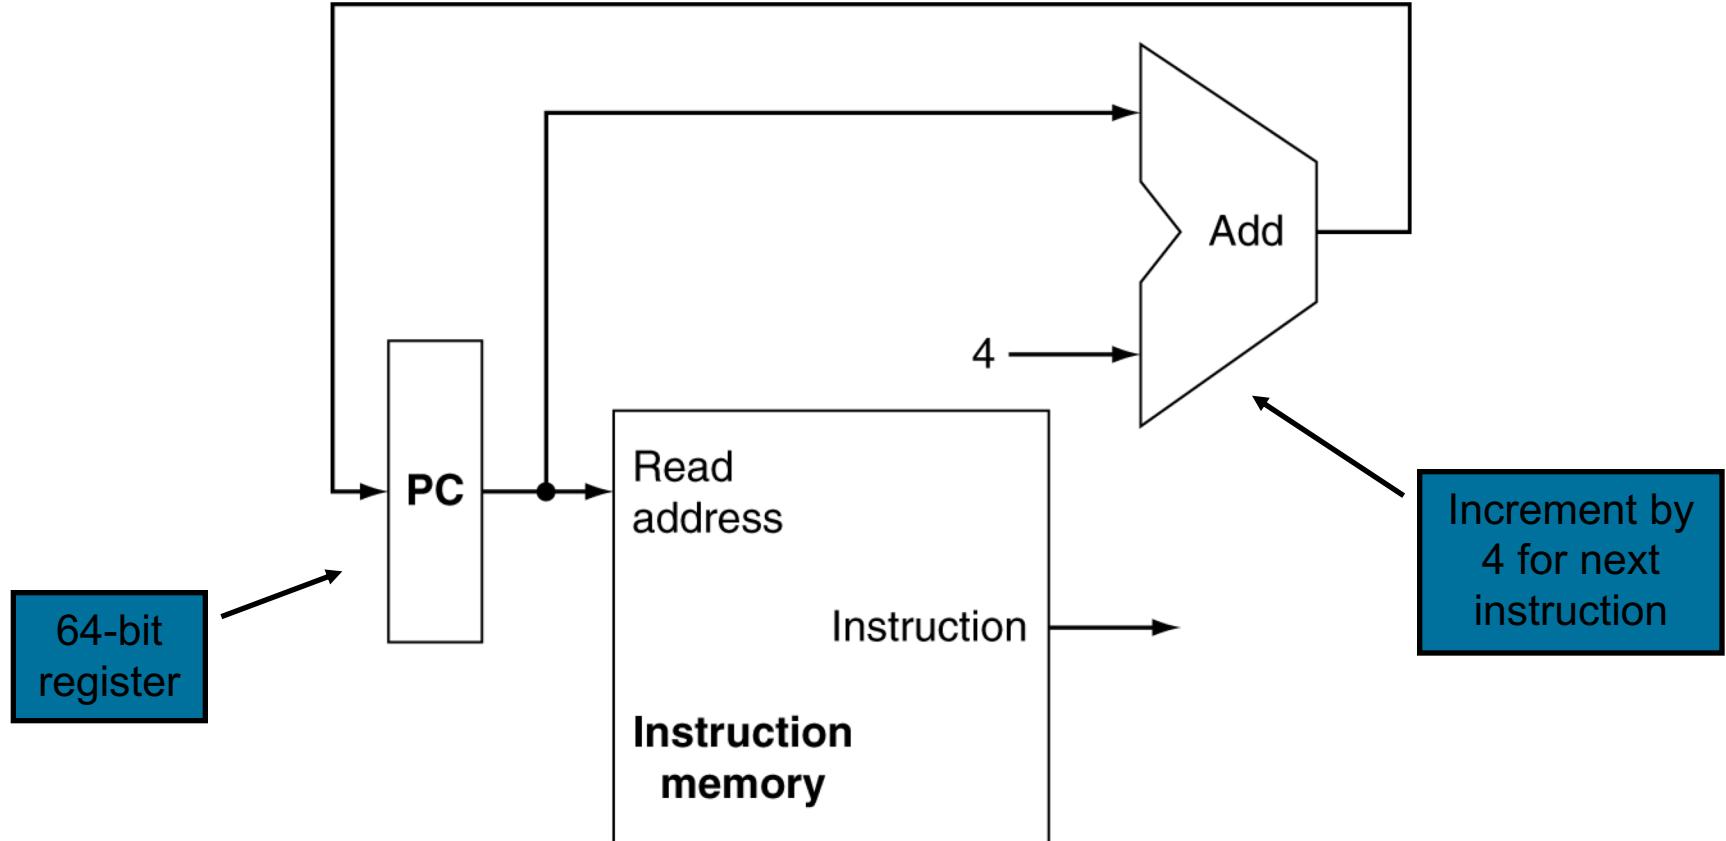
# Building a Datapath

12

- Datapath
  - Elements that process data and addresses in the CPU
    - Registers, ALUs, mux's, memories, ...
- We will build a RISC-V datapath incrementally
  - Refining the overview design

# Instruction Fetch

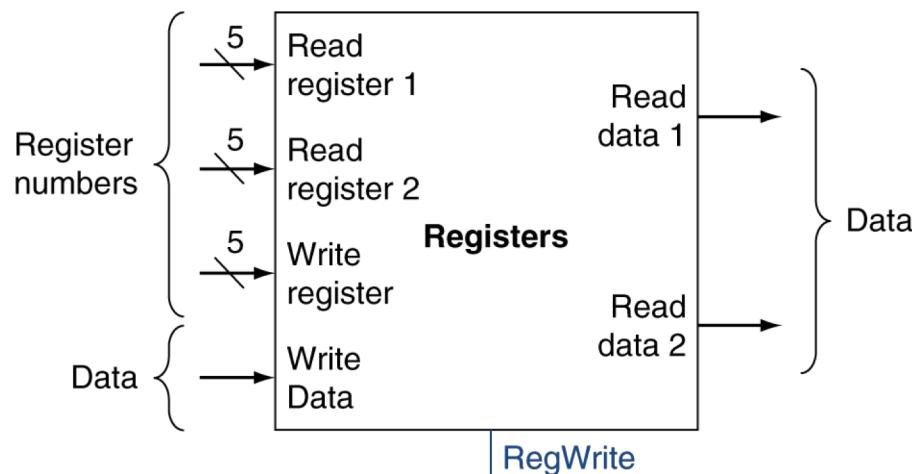
13



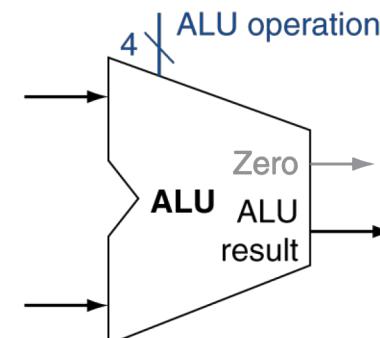
# R-Format Instructions

14

- Read two register operands
- Perform arithmetic/logical operation
- Write register result



a. Registers

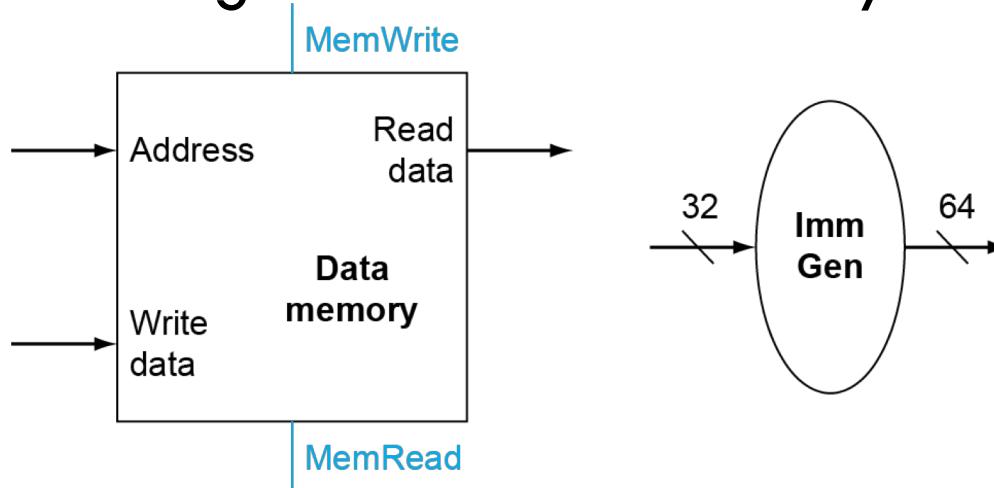


b. ALU

# Load/Store Instructions

15

- Read register operands
- Calculate address using 12-bit offset
  - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory



a. Data memory unit

b. Immediate generation unit

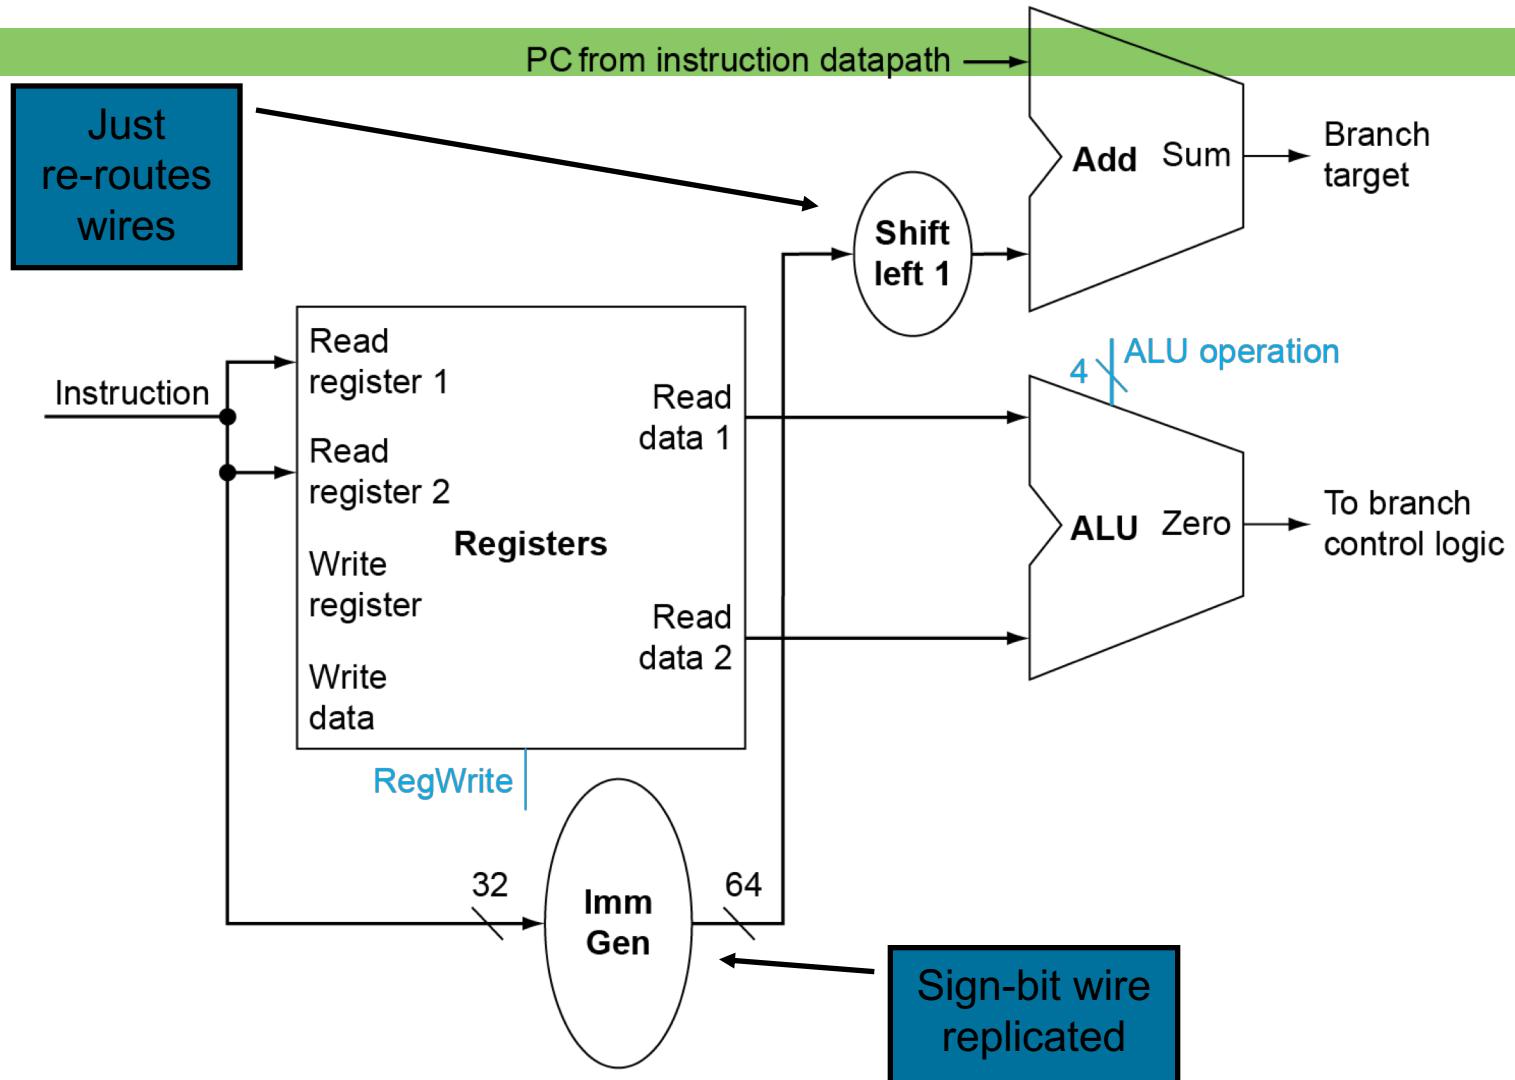
# Branch Instructions

16

- Read register operands
- Compare operands
  - ▣ Use ALU, subtract and check Zero output
- Calculate target address
  - ▣ Sign-extend displacement
  - ▣ Shift left 1 place (halfword displacement)
  - ▣ Add to PC value

# Branch Instructions

17



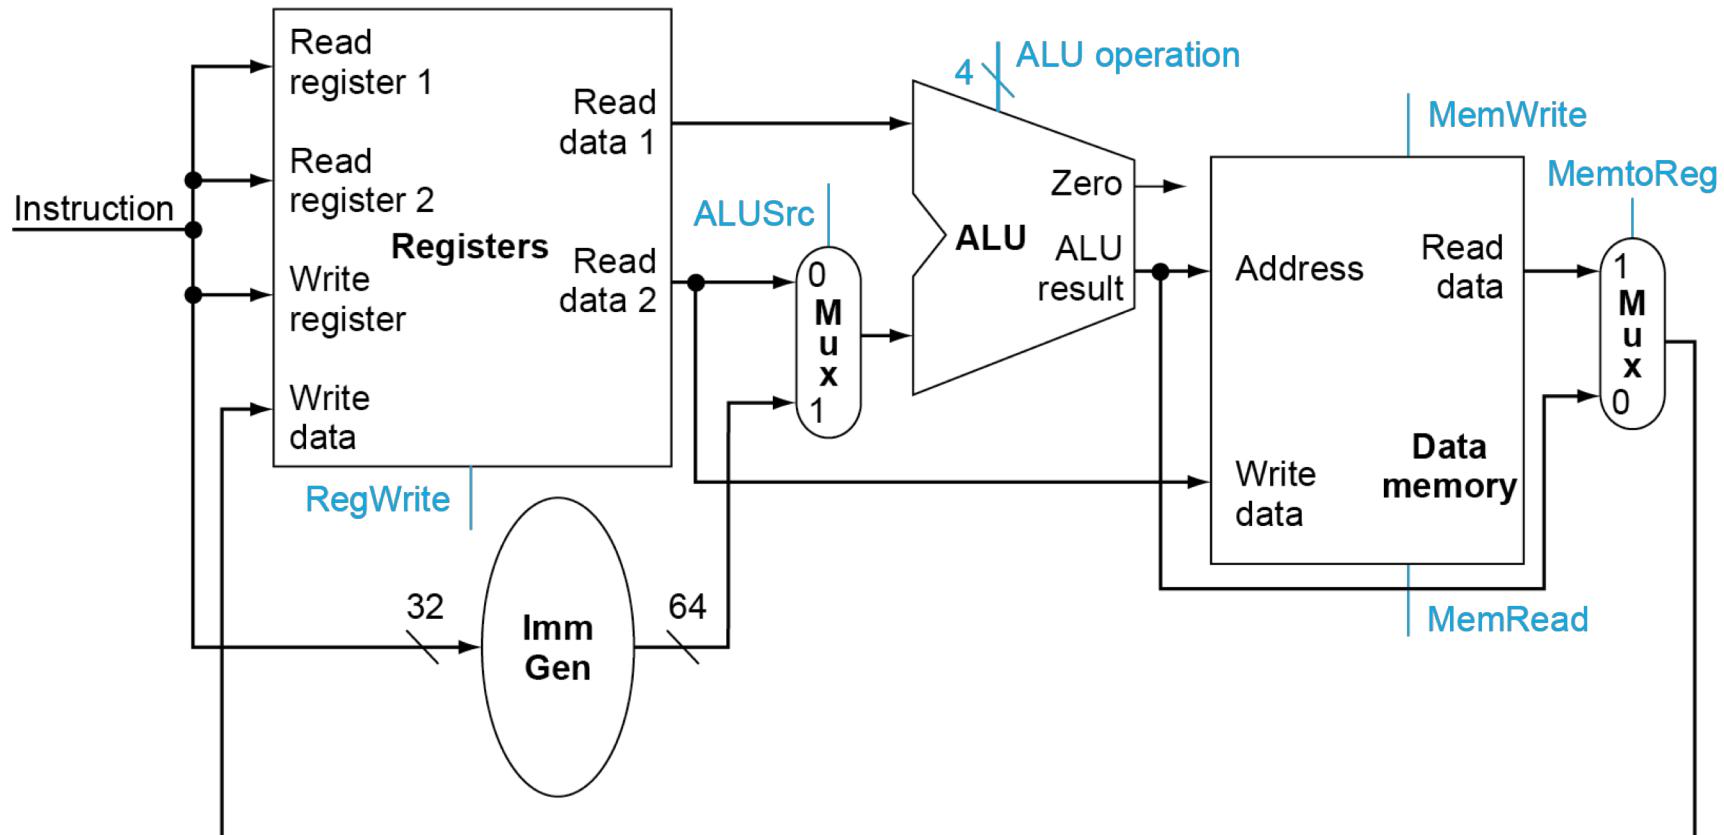
# Composing the Elements

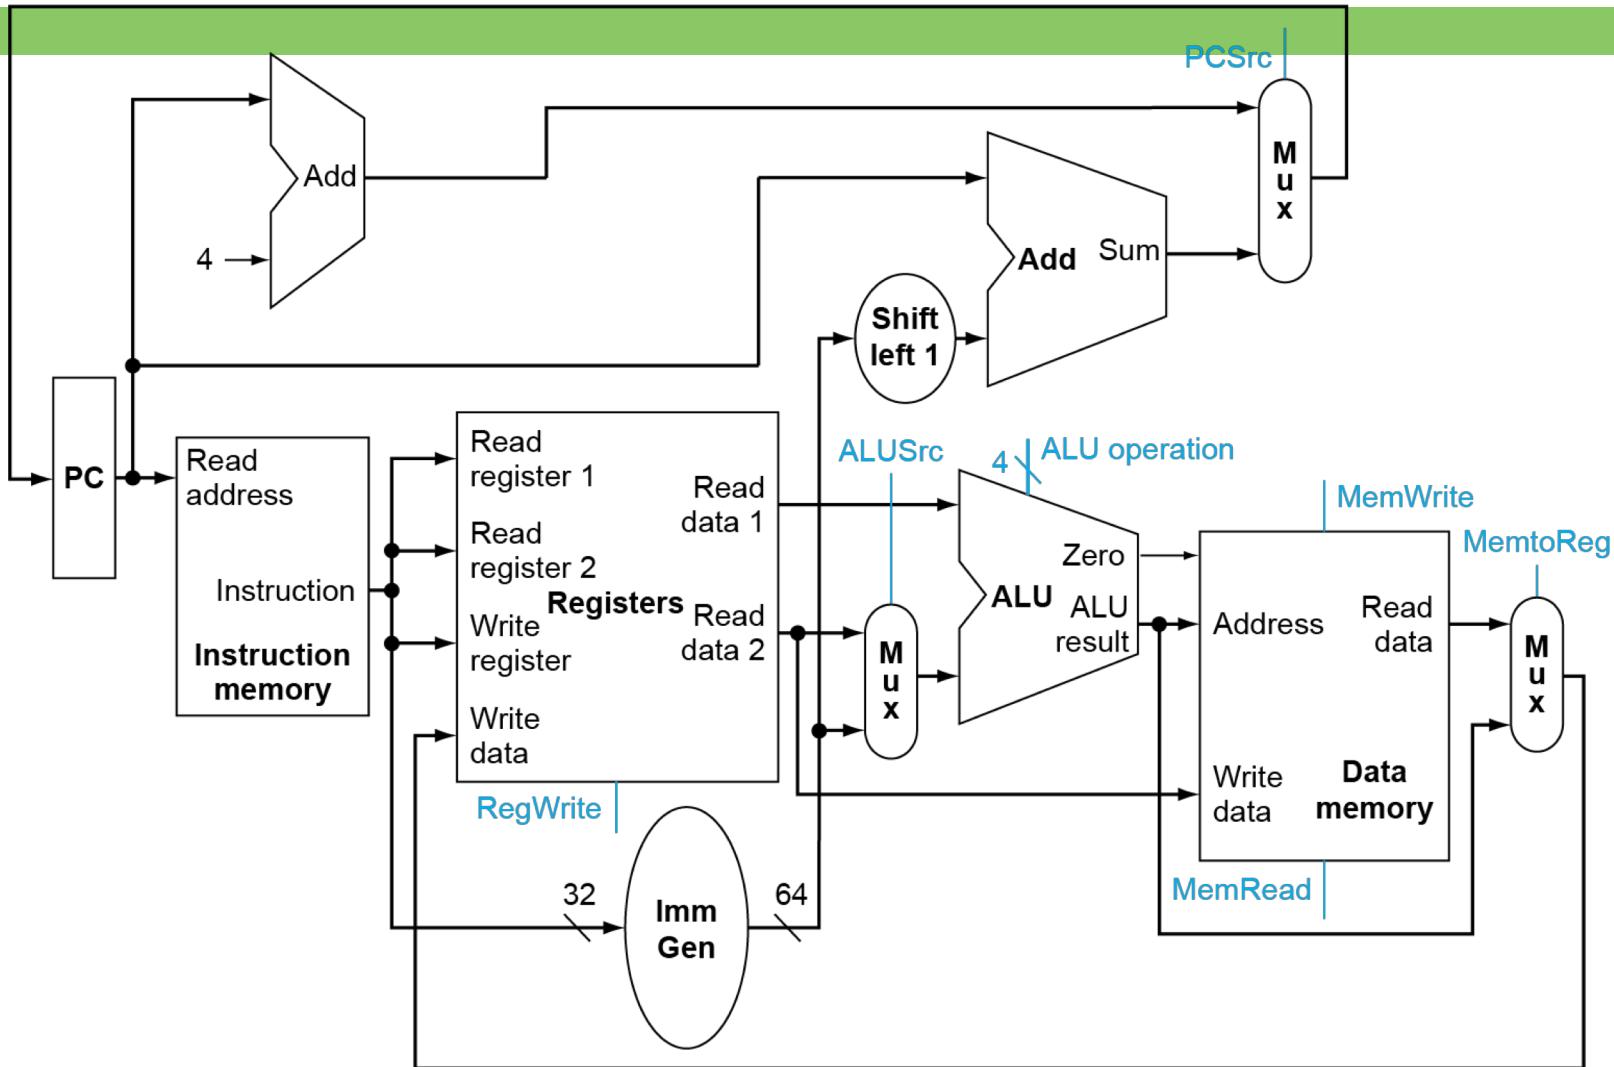
18

- First-cut data path does an instruction in one clock cycle
  - Each datapath element can only do one function at a time
  - Hence, we need separate instruction and data memories
- Use multiplexers where alternate data sources are used for different instructions

# R-Type/Load/Store Datapath

19





# ALU Control

21

- ALU used for
  - Load/Store: F = add
  - Branch: F = subtract
  - R-type: F depends on opcode

ALU control	Function
0000	AND
0001	OR
0010	add
0110	subtract

# ALU Control

22

- Assume 2-bit ALUOp derived from opcode
  - ▣ Combinational logic derives ALU control

opcode	ALUOp	Operation	Opcode field	ALU function	ALU control
ld	00	load register	XXXXXXXXXXXX	add	0010
sd	00	store register	XXXXXXXXXXXX	add	0010
beq	01	branch on equal	XXXXXXXXXXXX	subtract	0110
R-type	10	add	100000	add	0010
		subtract	100010	subtract	0110
		AND	100100	AND	0000
		OR	100101	OR	0001

# The Main Control Unit

23

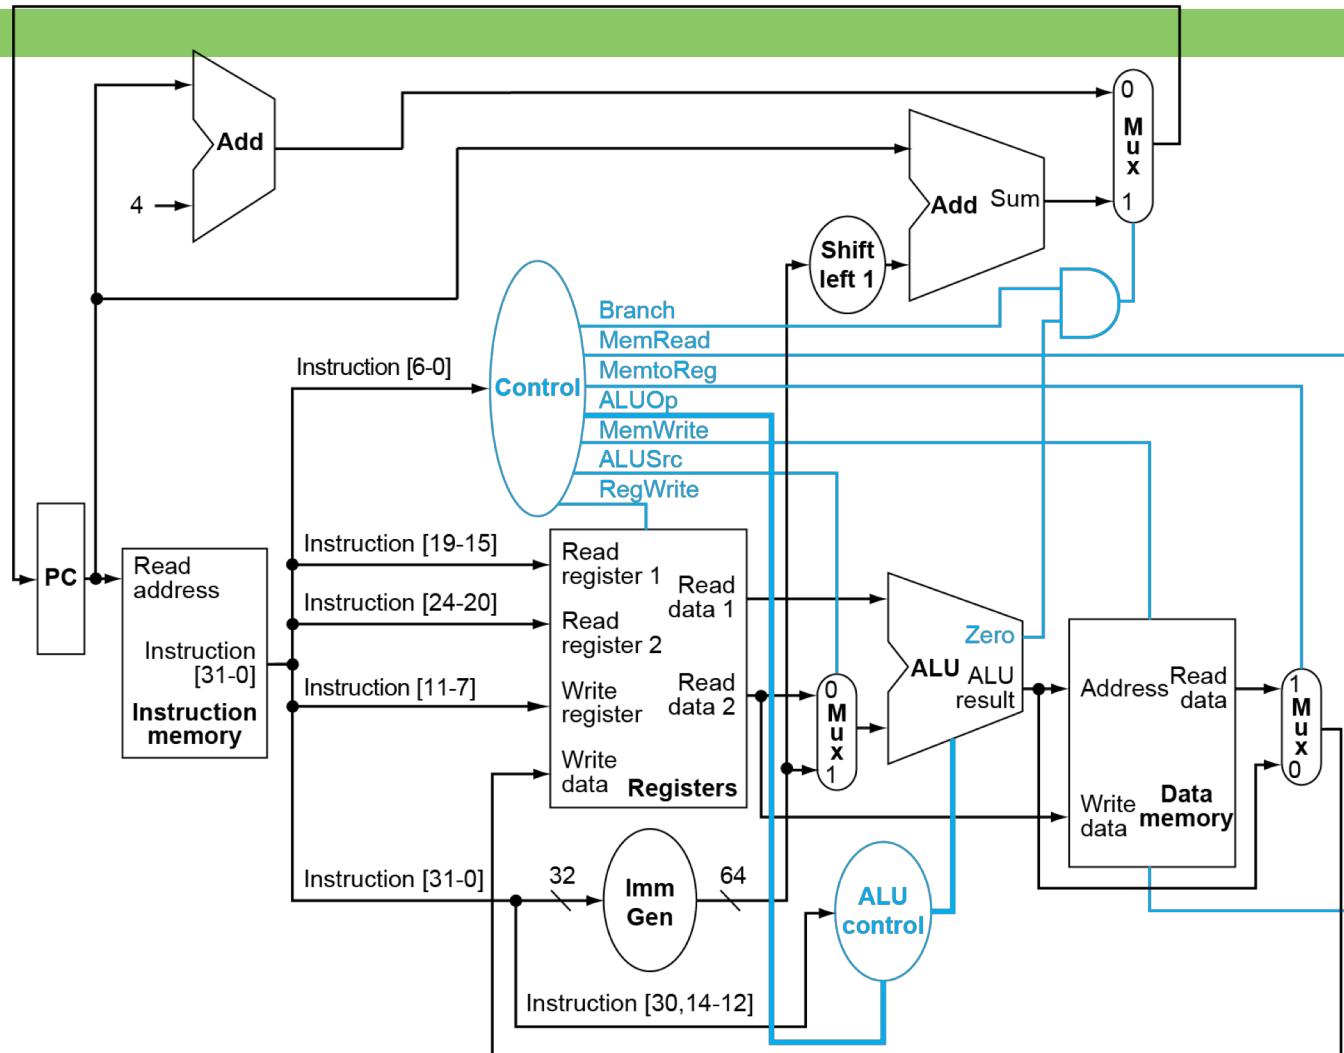
## □ Control signals derived from instruction

	Name (Bit position)	31:25	24:20	19:15	14:12	11:7	6:0
(a) R-type		funct7	rs2	rs1	funct3	rd	opcode
(b) I-type		immediate[11:0]		rs1	funct3	rd	opcode
(c) S-type		immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode
(d) SB-type		immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode

ALUOp	Funct7 field												Funct3 field	Operation
	ALUOp1	ALUOp0	I[31]	I[30]	I[29]	I[28]	I[27]	I[26]	I[25]	I[14]	I[13]	I[12]		
0	0	X	X	X	X	X	X	X	X	X	X	X	0010	
X	1	X	X	X	X	X	X	X	X	X	X	X	0110	
1	X	0	0	0	0	0	0	0	0	0	0	0	0010	
1	X	0	1	0	0	0	0	0	0	0	0	0	0110	
1	X	0	0	0	0	0	0	0	0	1	1	1	0000	
1	X	0	0	0	0	0	0	0	0	1	1	0	0001	

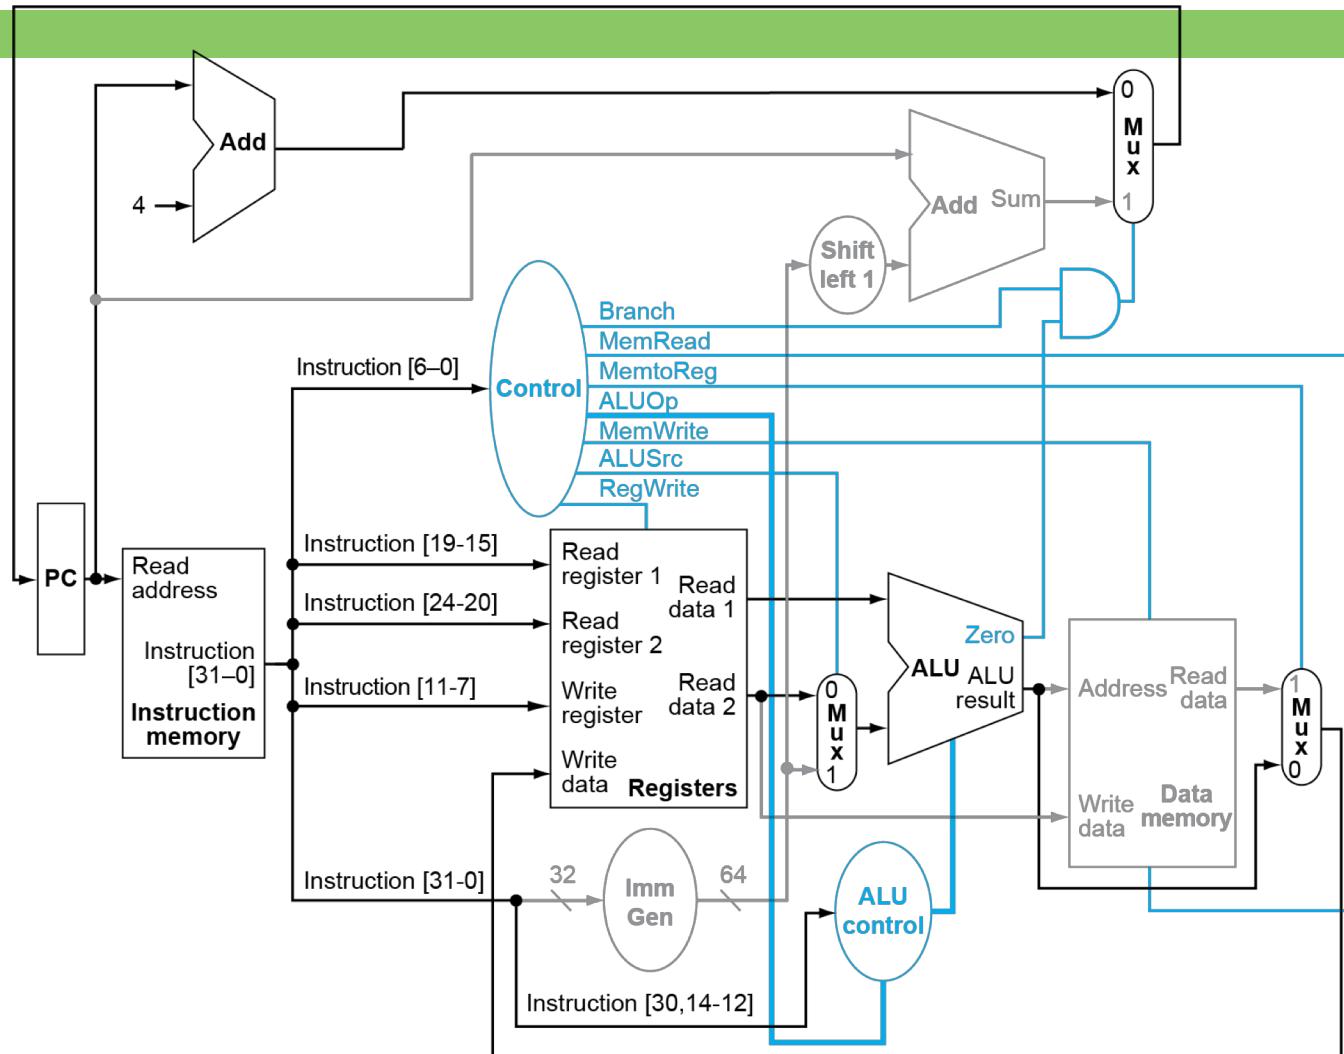
# Datapath With Control

24



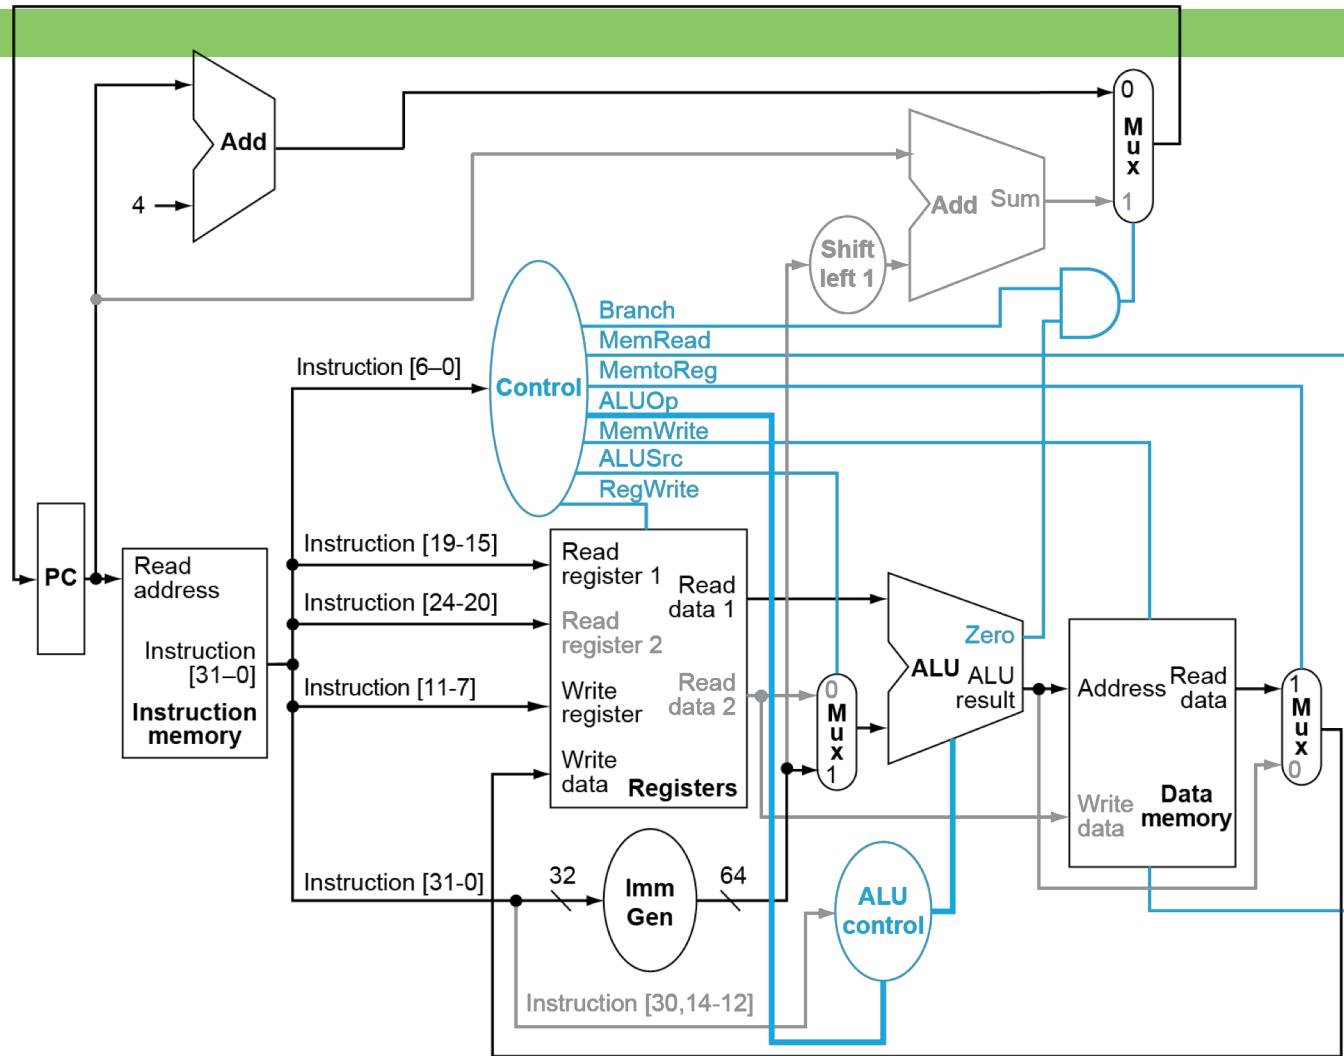
# R-Type Instruction

25



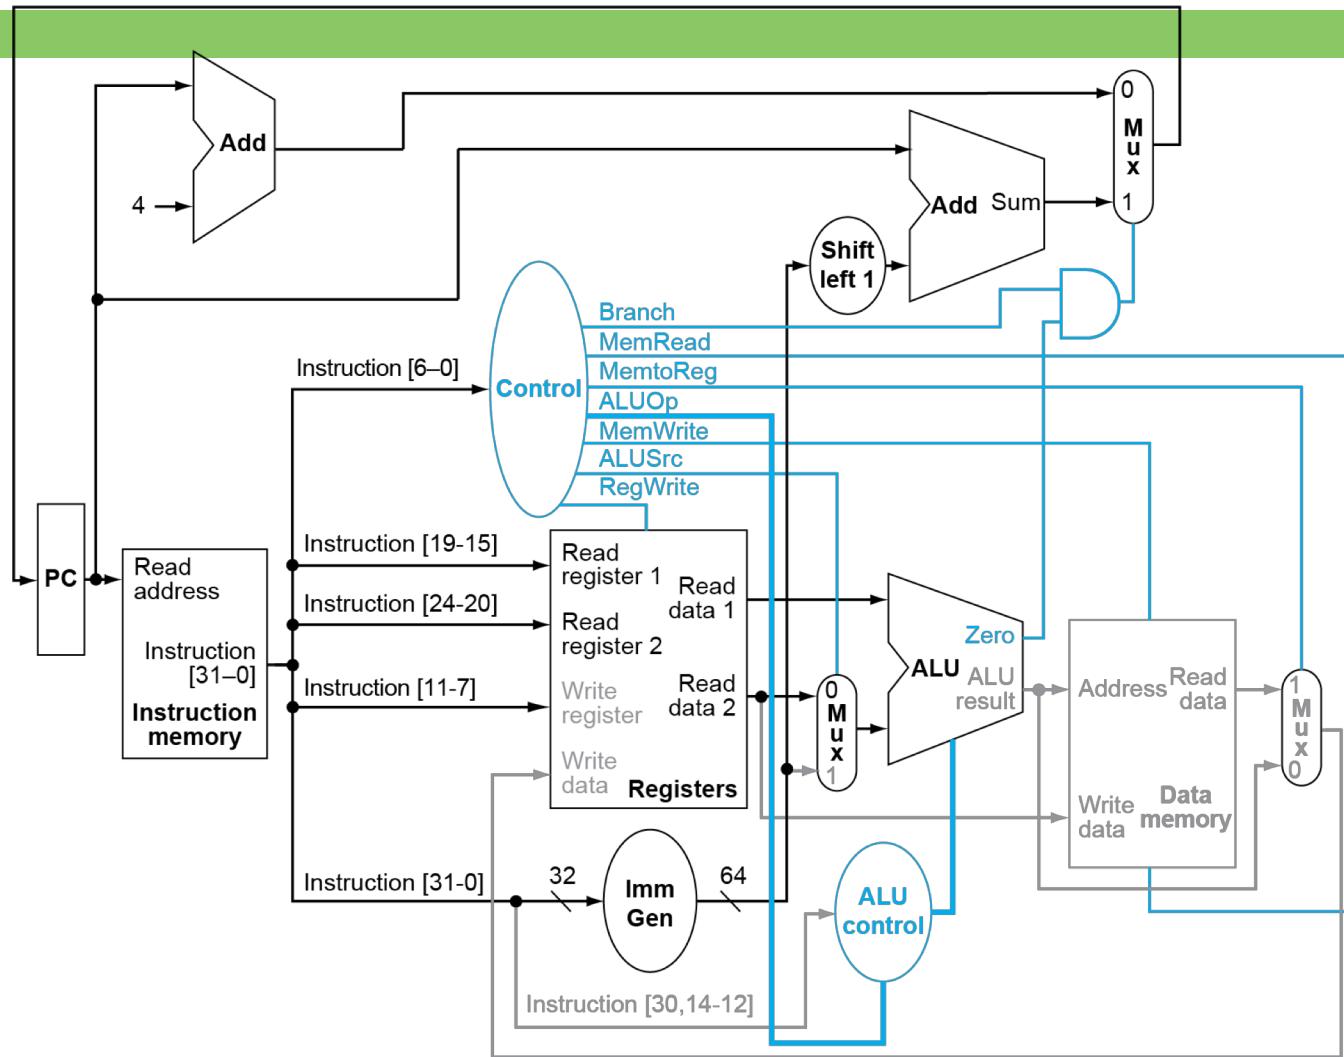
# Load Instruction

26



# BEQ Instruction

27



# Performance Issues

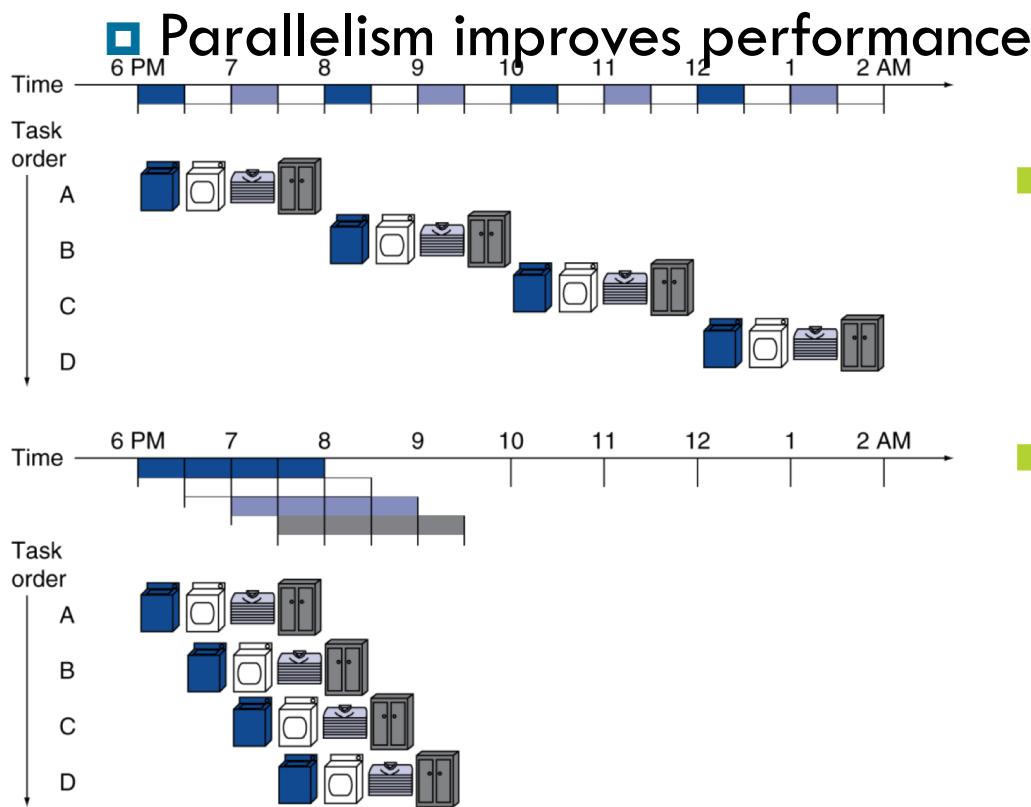
28

- Longest delay determines clock period
  - ▣ Critical path: load instruction
  - ▣ Instruction memory → register file → ALU → data memory  
→ register file
- Not feasible to vary period for different instructions
- Violates design principle
  - ▣ Making the common case fast
- We will improve performance by pipelining

# Pipelining Analogy

29

## □ Pipelined laundry: overlapping execution



- Four loads:
  - Speedup  
 $= 8/3.5 = 2.3$
- Non-stop:
  - Speedup  
 $= 2n/0.5n + 1.5 \approx 4$   
= number of stages

# RISC-V Pipeline

30

- Five stages, one step per stage
  1. IF: Instruction fetch from memory
  2. ID: Instruction decode & register read
  3. EX: Execute operation or calculate address
  4. MEM: Access memory operand
  5. WB: Write result back to register

# Pipeline Performance

31

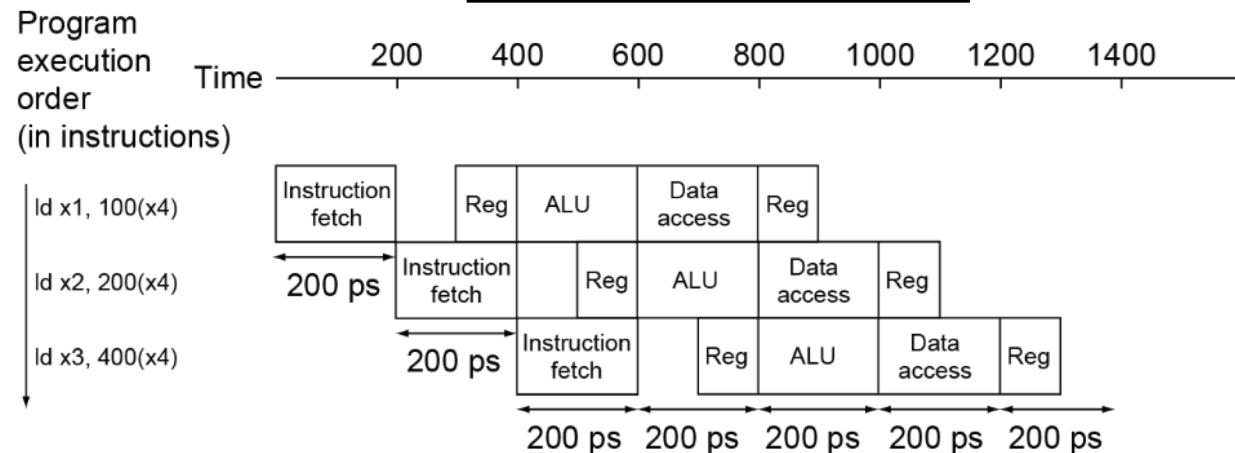
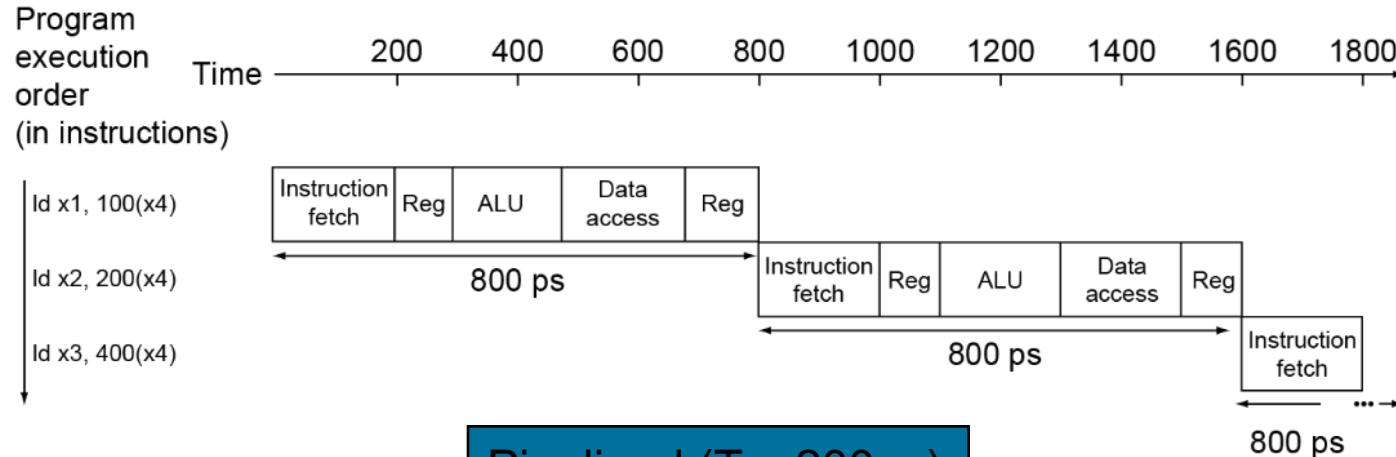
- Assume time for stages is
  - 100ps for register read or write
  - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
ld	200ps	100 ps	200ps	200ps	100 ps	800ps
sd	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

# Pipeline Performance

32

Single-cycle ( $T_c = 800\text{ps}$ )



# Pipeline Speedup

33

- If all stages are balanced
  - ▣ i.e., all take the same time
  - ▣ Time between instructions<sub>pipelined</sub>  
= Time between instructions<sub>nonpipelined</sub>  

---

Number of stages
- If not balanced, speedup is less
- Speedup due to increased throughput
  - ▣ Latency (time for each instruction) does not decrease

# Pipelining and ISA Design

34

- RISC-V ISA designed for pipelining
  - All instructions are 32-bits
    - Easier to fetch and decode in one cycle
    - c.f. x86: 1- to 17-byte instructions
  - Few and regular instruction formats
    - Can decode and read registers in one step
  - Load/store addressing
    - Can calculate address in 3<sup>rd</sup> stage, access memory in 4<sup>th</sup> stage

# Hazards

35

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
  - ▣ A required resource is busy
- Data hazard
  - ▣ Need to wait for previous instruction to complete its data read/write
- Control hazard
  - ▣ Deciding on control action depends on previous instruction

# Structure Hazards

36

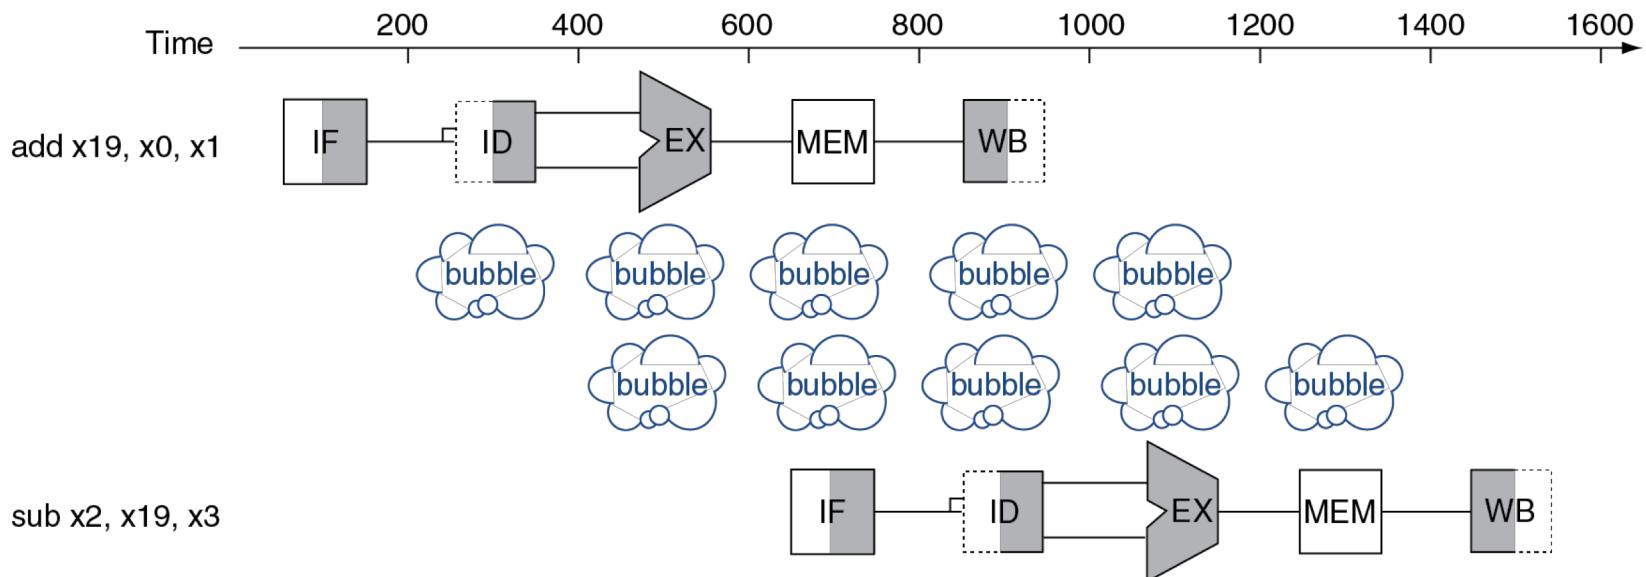
- Conflict for use of a resource
- In RISC-V pipeline with a single memory
  - ▣ Load/store requires data access
  - ▣ Instruction fetch would have to *stall* for that cycle
    - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require separate instruction/data memories
  - ▣ Or separate instruction/data caches

# Data Hazards

37

- An instruction depends on completion of data access by a previous instruction

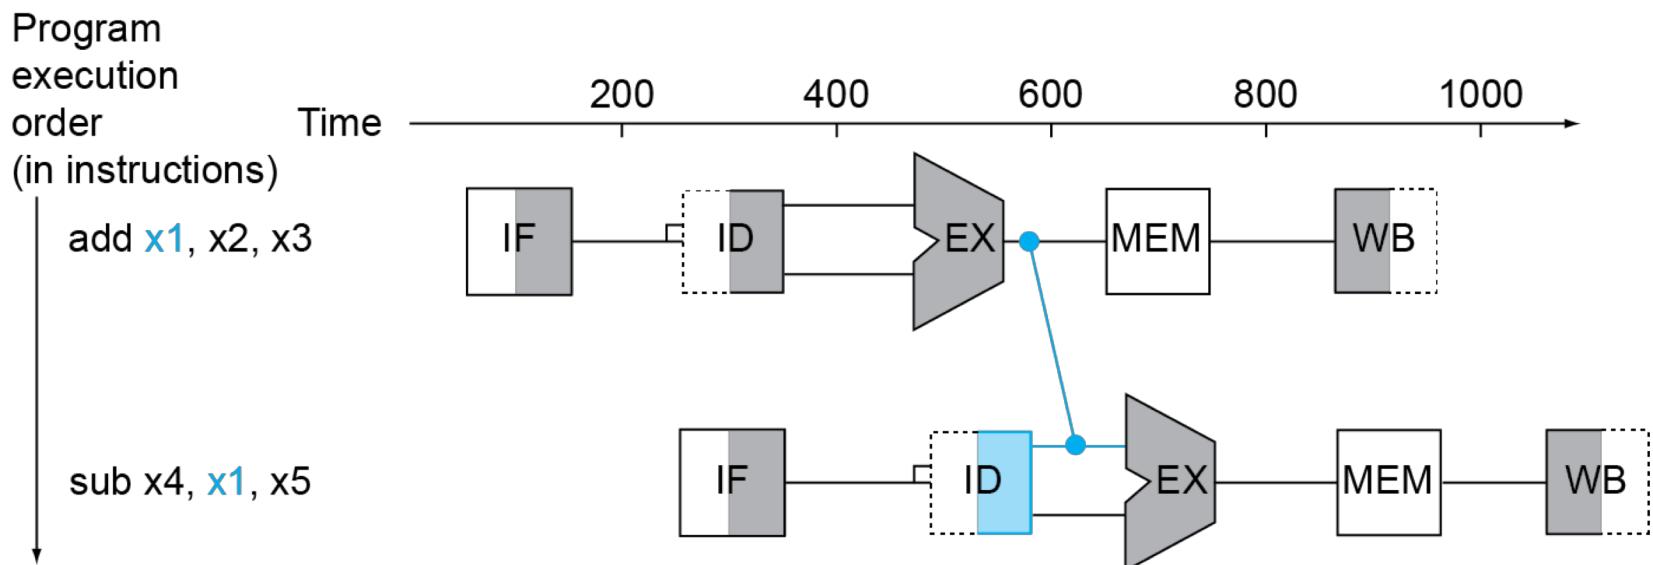
- add      **x19, x0, x1**  
sub      x2, **x19, x3**



# Forwarding (aka Bypassing)

38

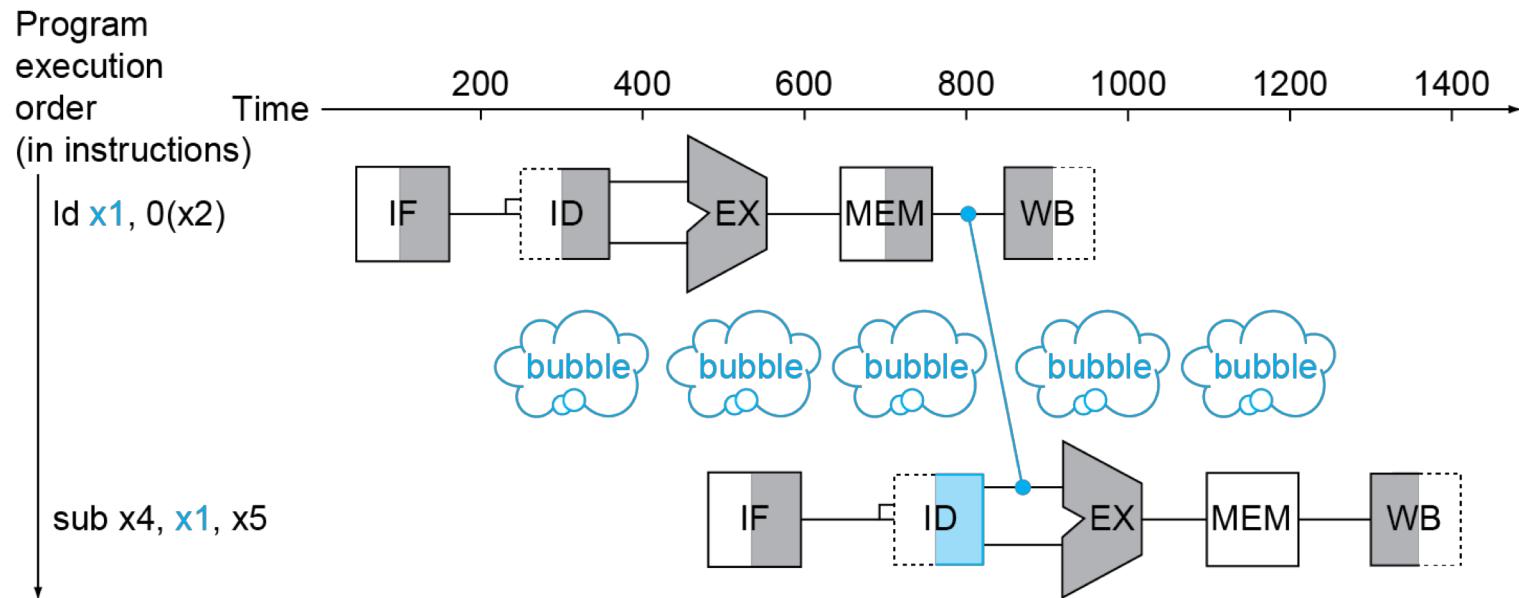
- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra connections in the datapath



# Load-Use Data Hazard

39

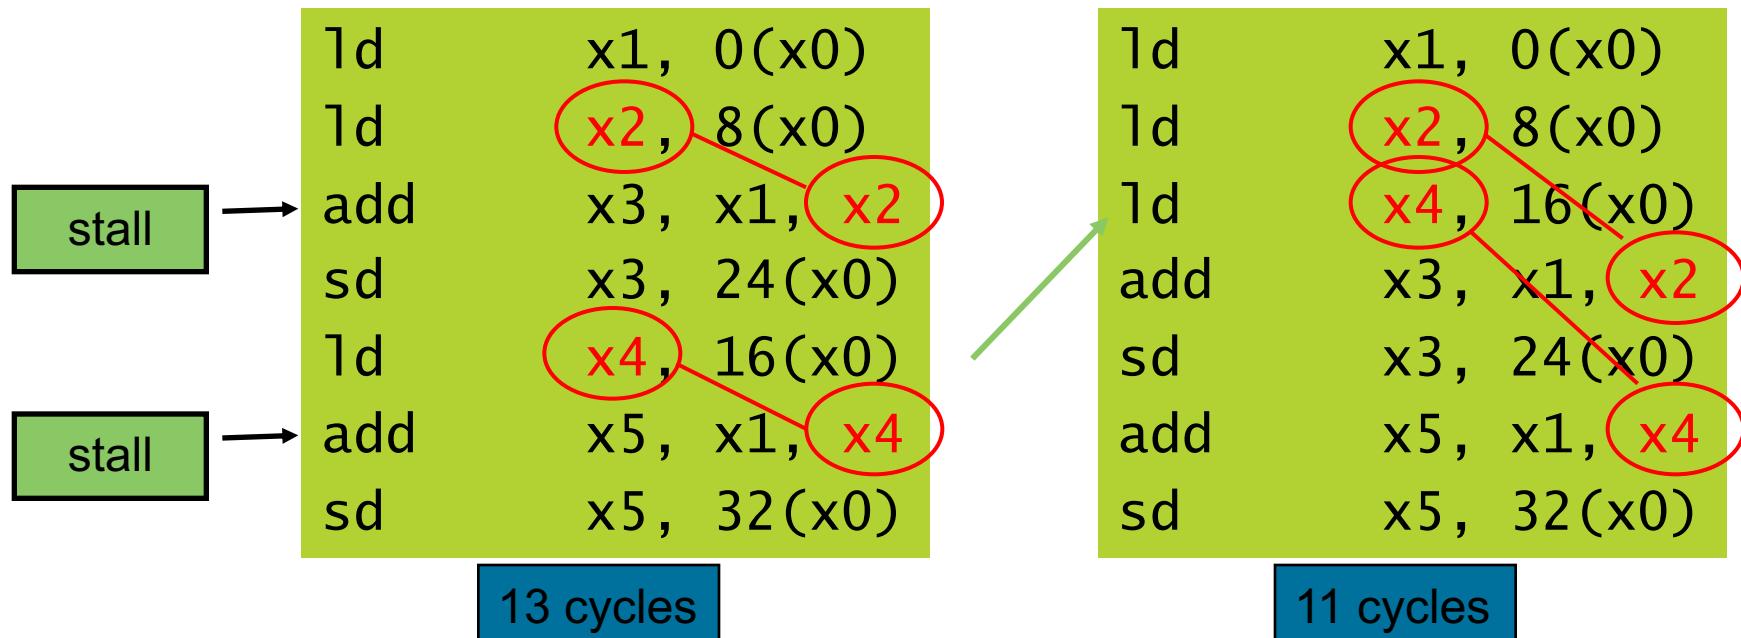
- Can't always avoid stalls by forwarding
  - If value not computed when needed
  - Can't forward backward in time!



# Code Scheduling to Avoid Stalls

40

- ❑ Reorder code to avoid use of load result in the next instruction
- ❑ C code for  $a = b + e; c = b + f;$



# Control Hazards

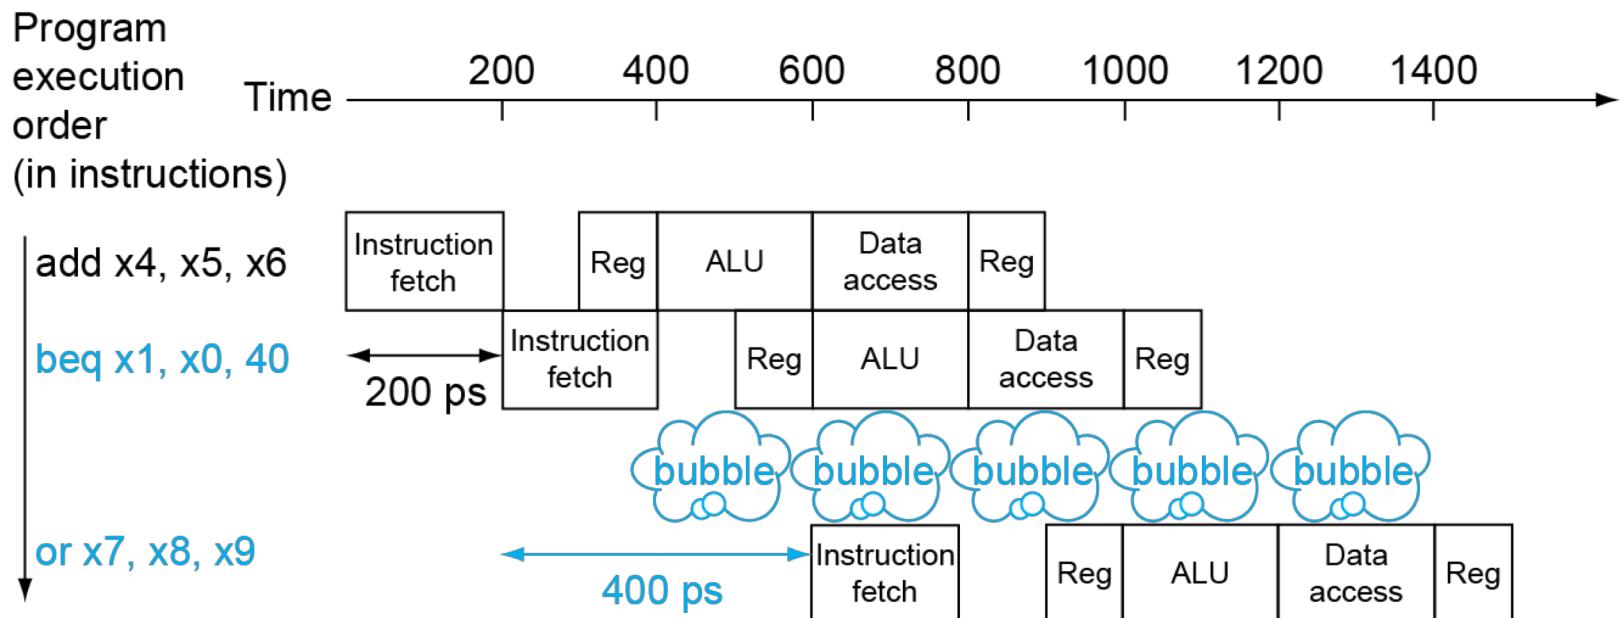
41

- Branch determines flow of control
  - Fetching next instruction depends on branch outcome
  - Pipeline can't always fetch correct instruction
    - Still working on ID stage of branch
- In RISC-V pipeline
  - Need to compare registers and compute target early in the pipeline
  - Add hardware to do it in ID stage

# Stall on Branch

42

- Wait until branch outcome determined before fetching next instruction



# Branch Prediction

43

- Longer pipelines can't readily determine branch outcome early
  - ▣ Stall penalty becomes unacceptable
- Predict outcome of branch
  - ▣ Only stall if prediction is wrong
- In RISC-V pipeline
  - ▣ Can predict branches not taken
  - ▣ Fetch instruction after branch, with no delay

# More-Realistic Branch Prediction

44

- Static branch prediction
  - ▣ Based on typical branch behavior
  - ▣ Example: loop and if-statement branches
    - Predict backward branches taken
    - Predict forward branches not taken
- Dynamic branch prediction
  - ▣ Hardware measures actual branch behavior
    - e.g., record recent history of each branch
  - ▣ Assume future behavior will continue the trend
    - When wrong, stall while re-fetching, and update history

# Pipeline Summary

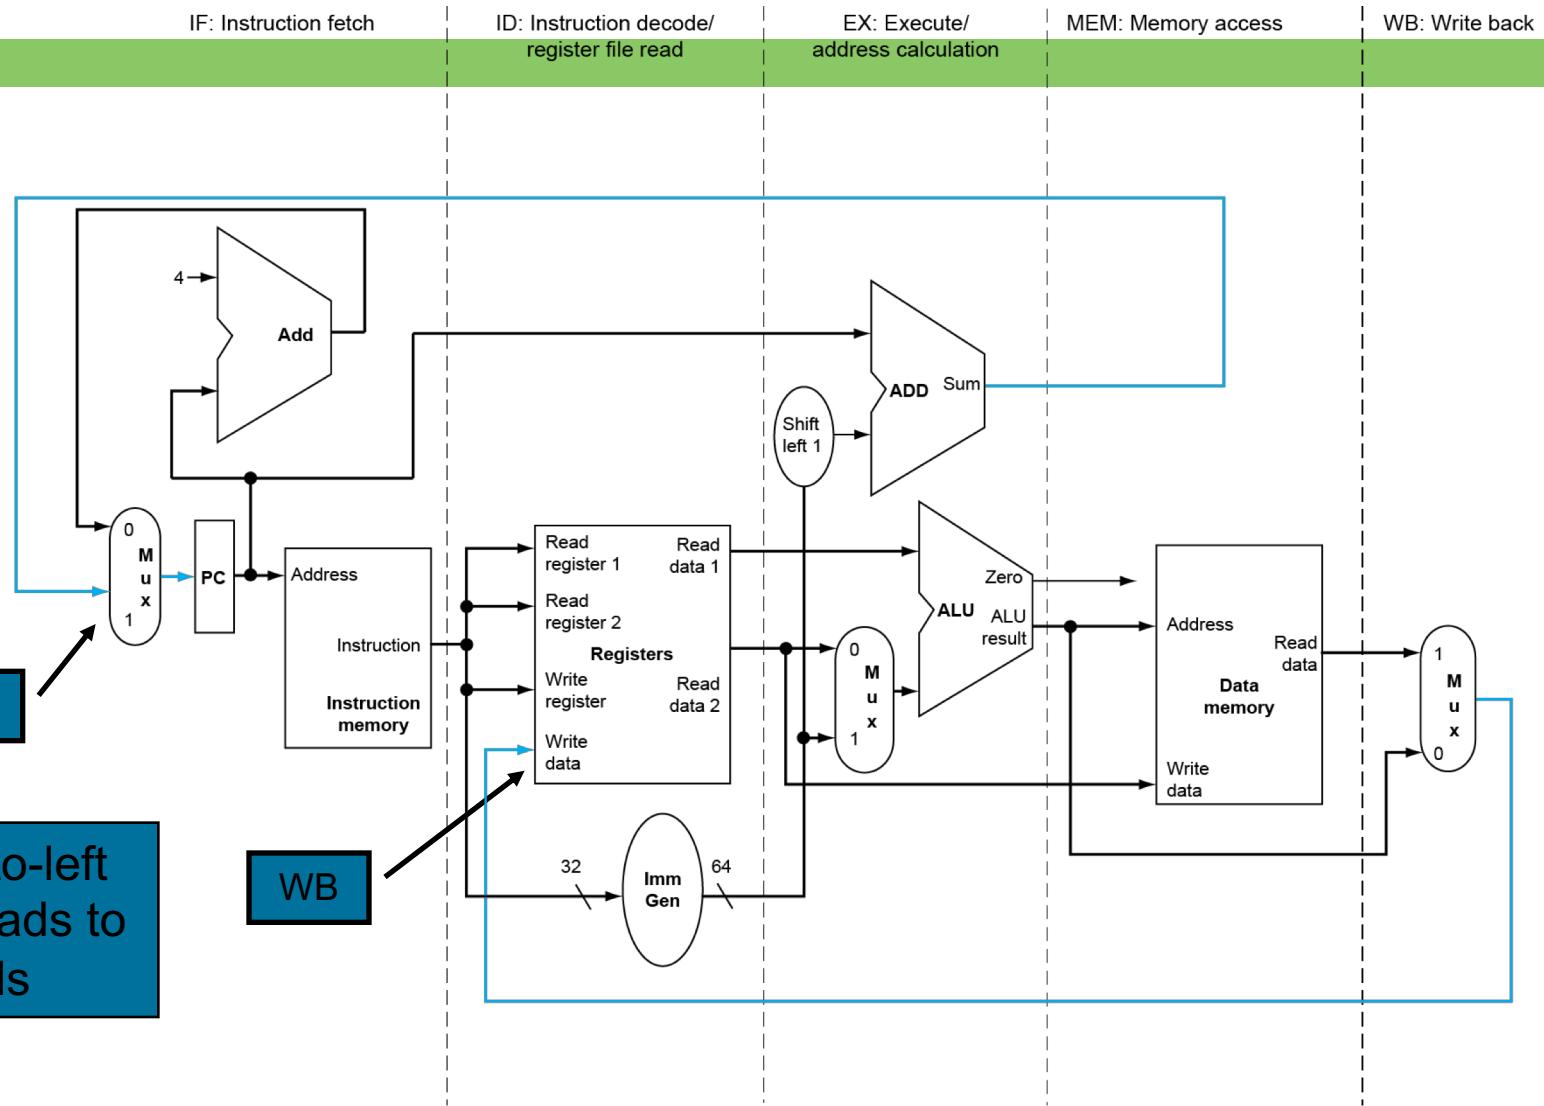
45

## The BIG Picture

- Pipelining improves performance by increasing instruction throughput
  - Executes multiple instructions in parallel
  - Each instruction has the same latency
- Subject to hazards
  - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

# RISC-V Pipelined Datapath

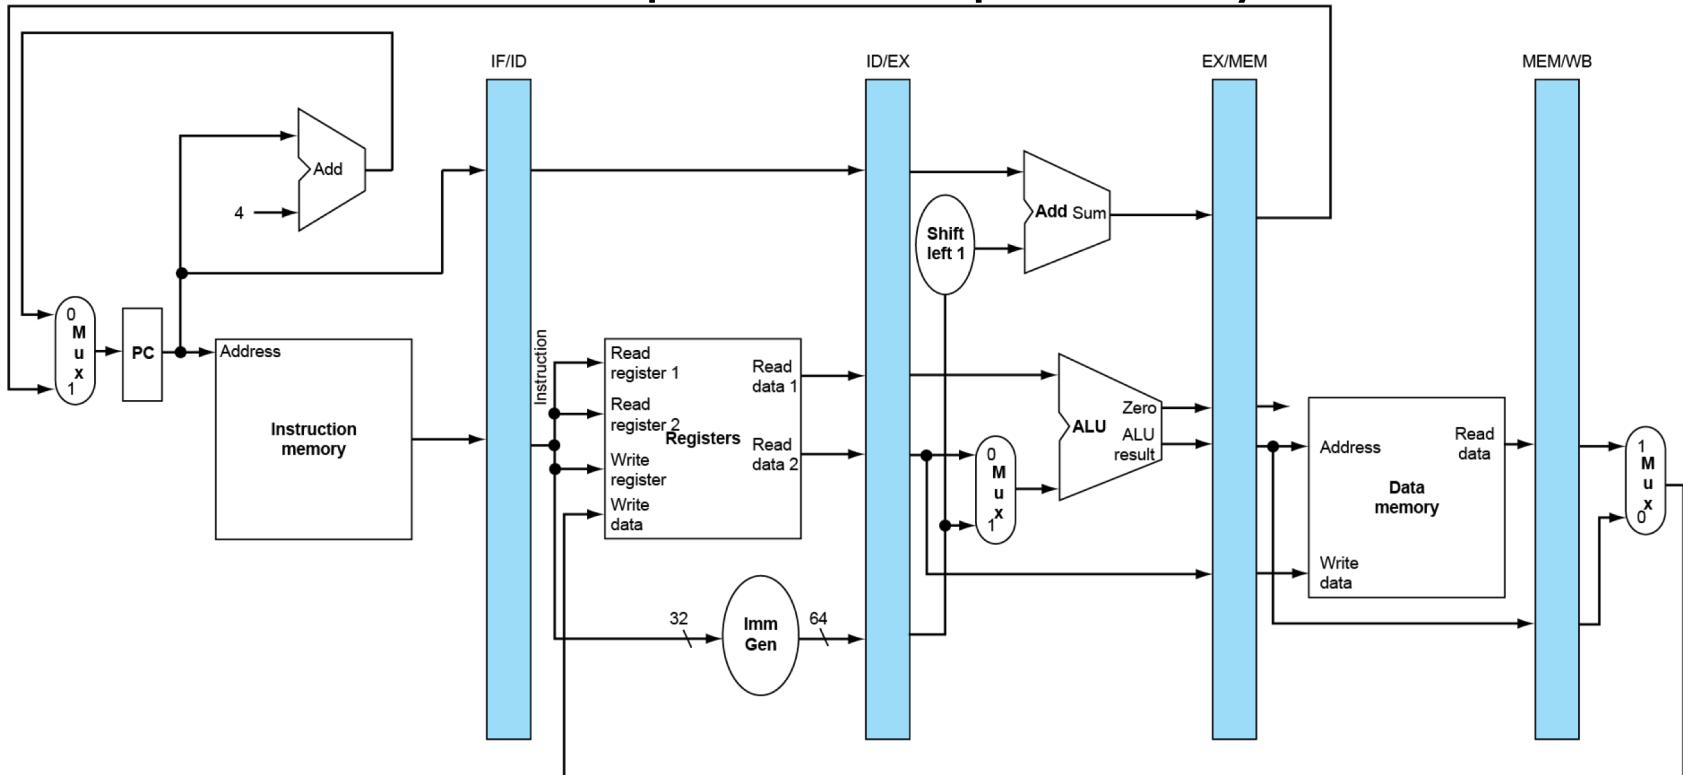
46



# Pipeline registers

47

- Need registers between stages
  - To hold information produced in previous cycle



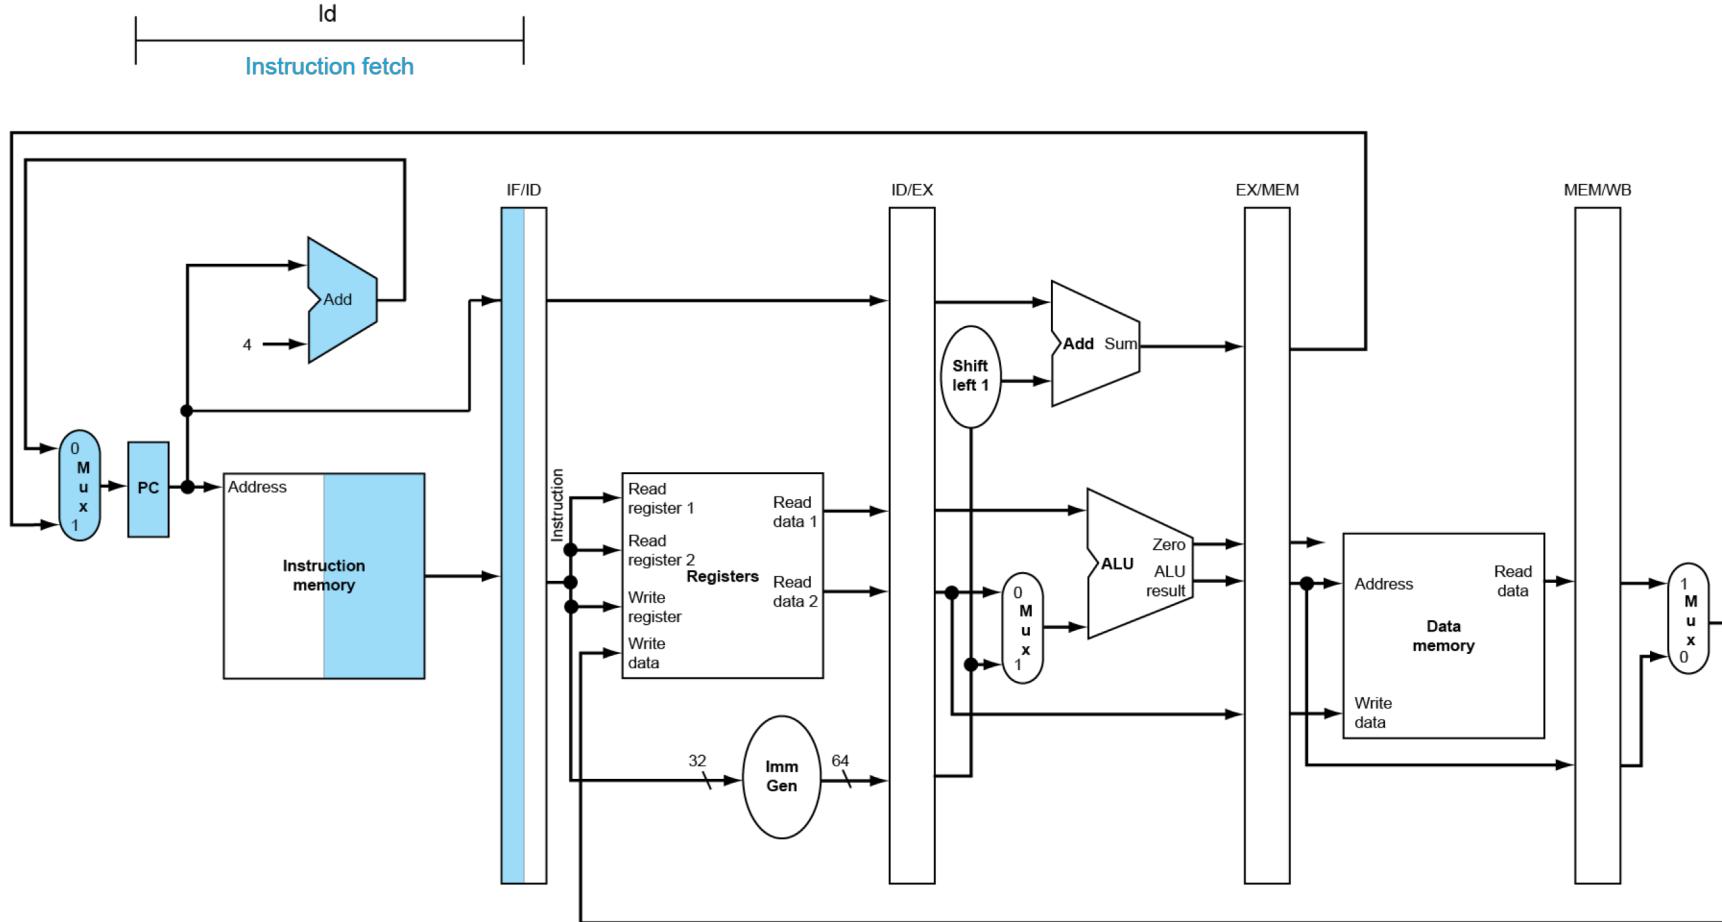
# Pipeline Operation

48

- Cycle-by-cycle flow of instructions through the pipelined datapath
  - “Single-clock-cycle” pipeline diagram
    - Shows pipeline usage in a single cycle
    - Highlight resources used
  - c.f. “multi-clock-cycle” diagram
    - Graph of operation over time
- We'll look at “single-clock-cycle” diagrams for load & store

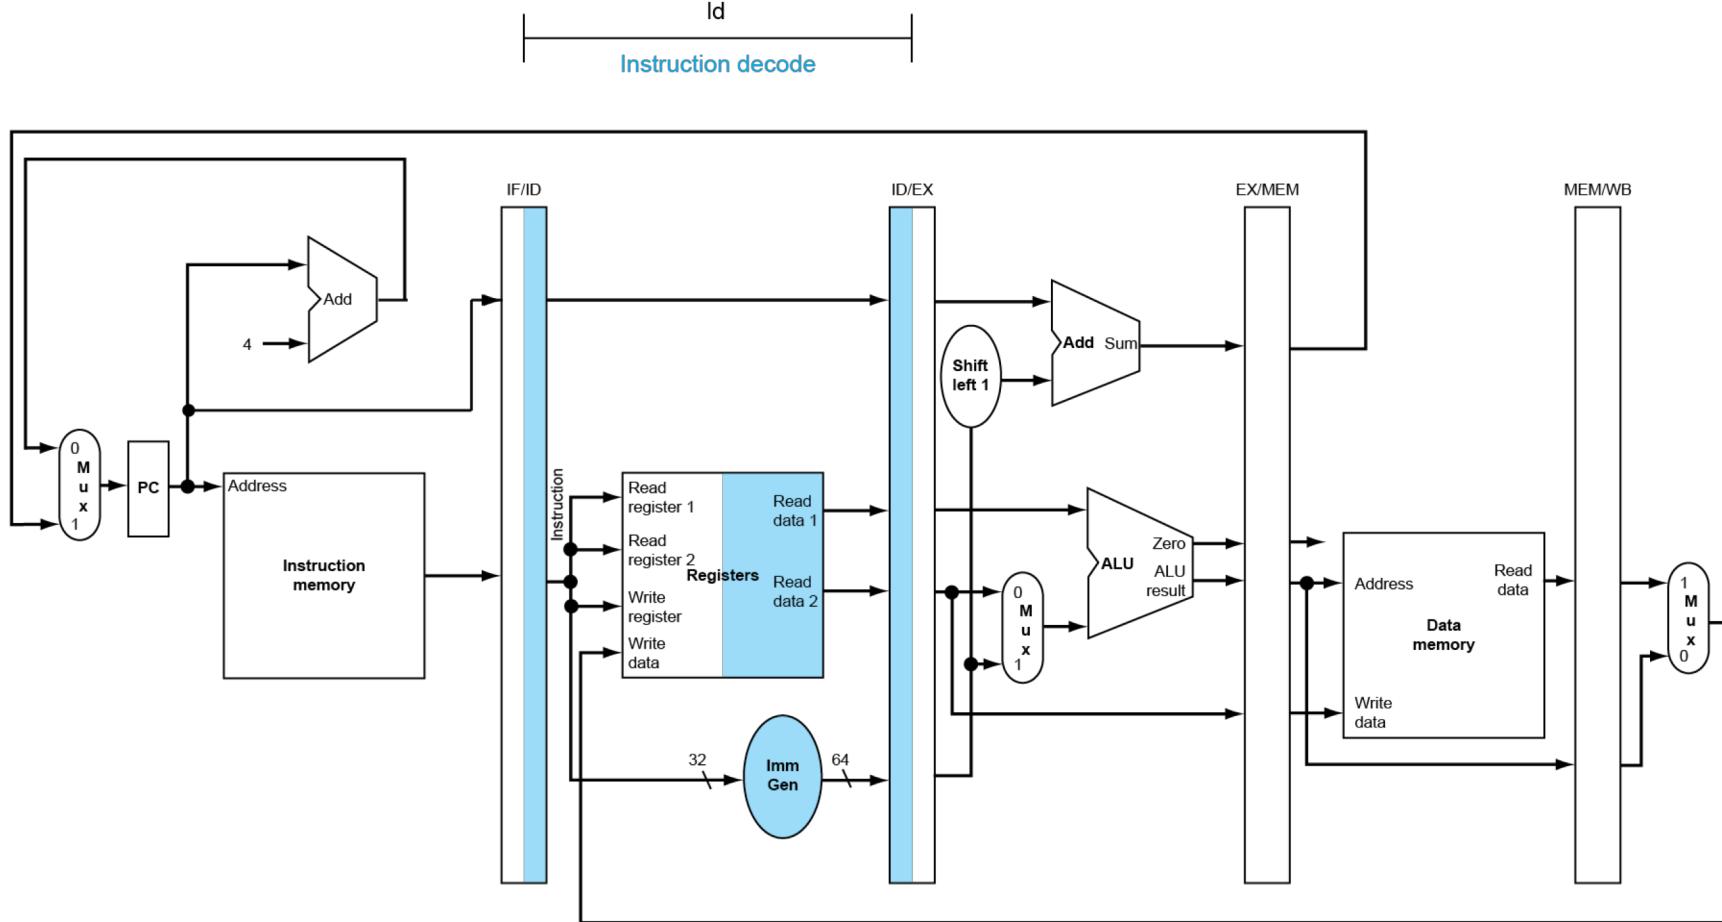
# IF for Load, Store, ...

49



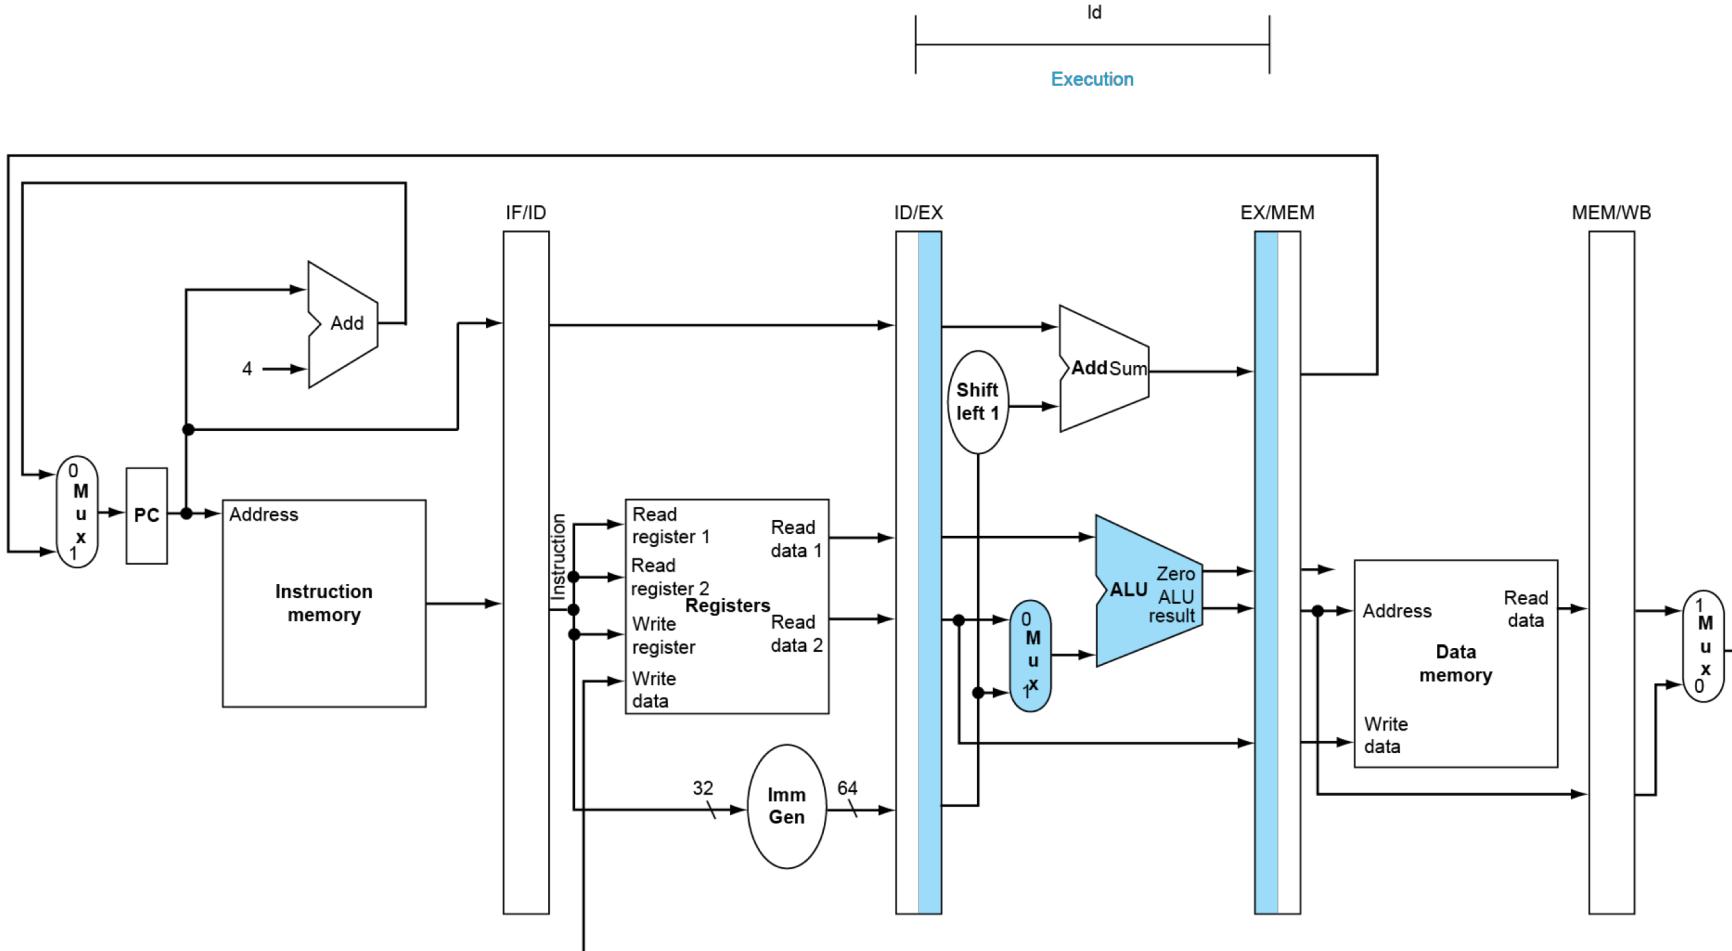
# ID for Load, Store, ...

50



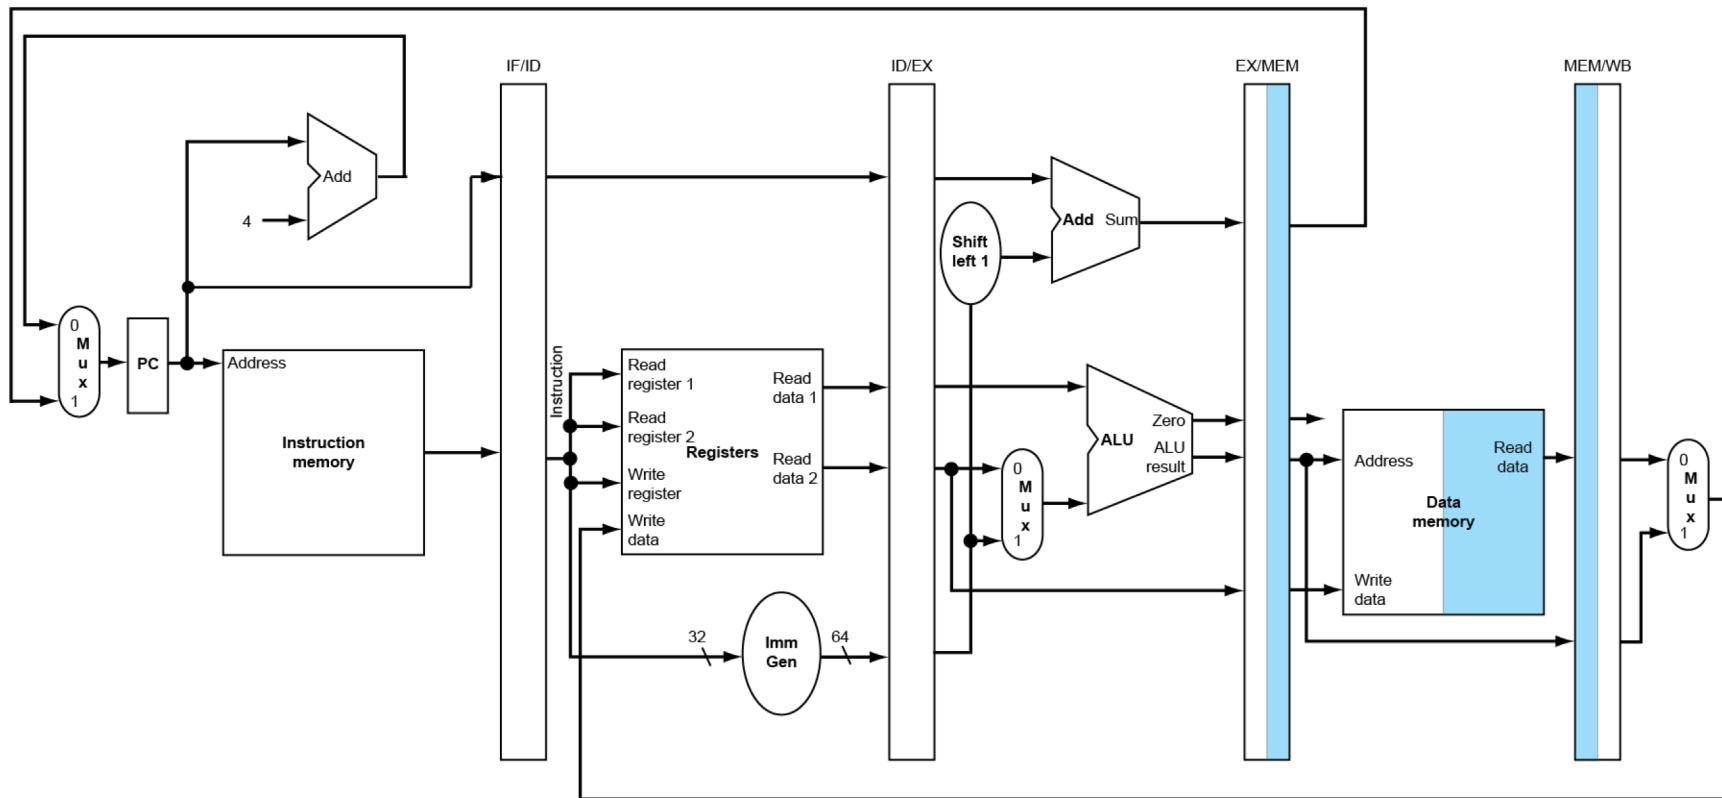
# EX for Load

51



# MEM for Load

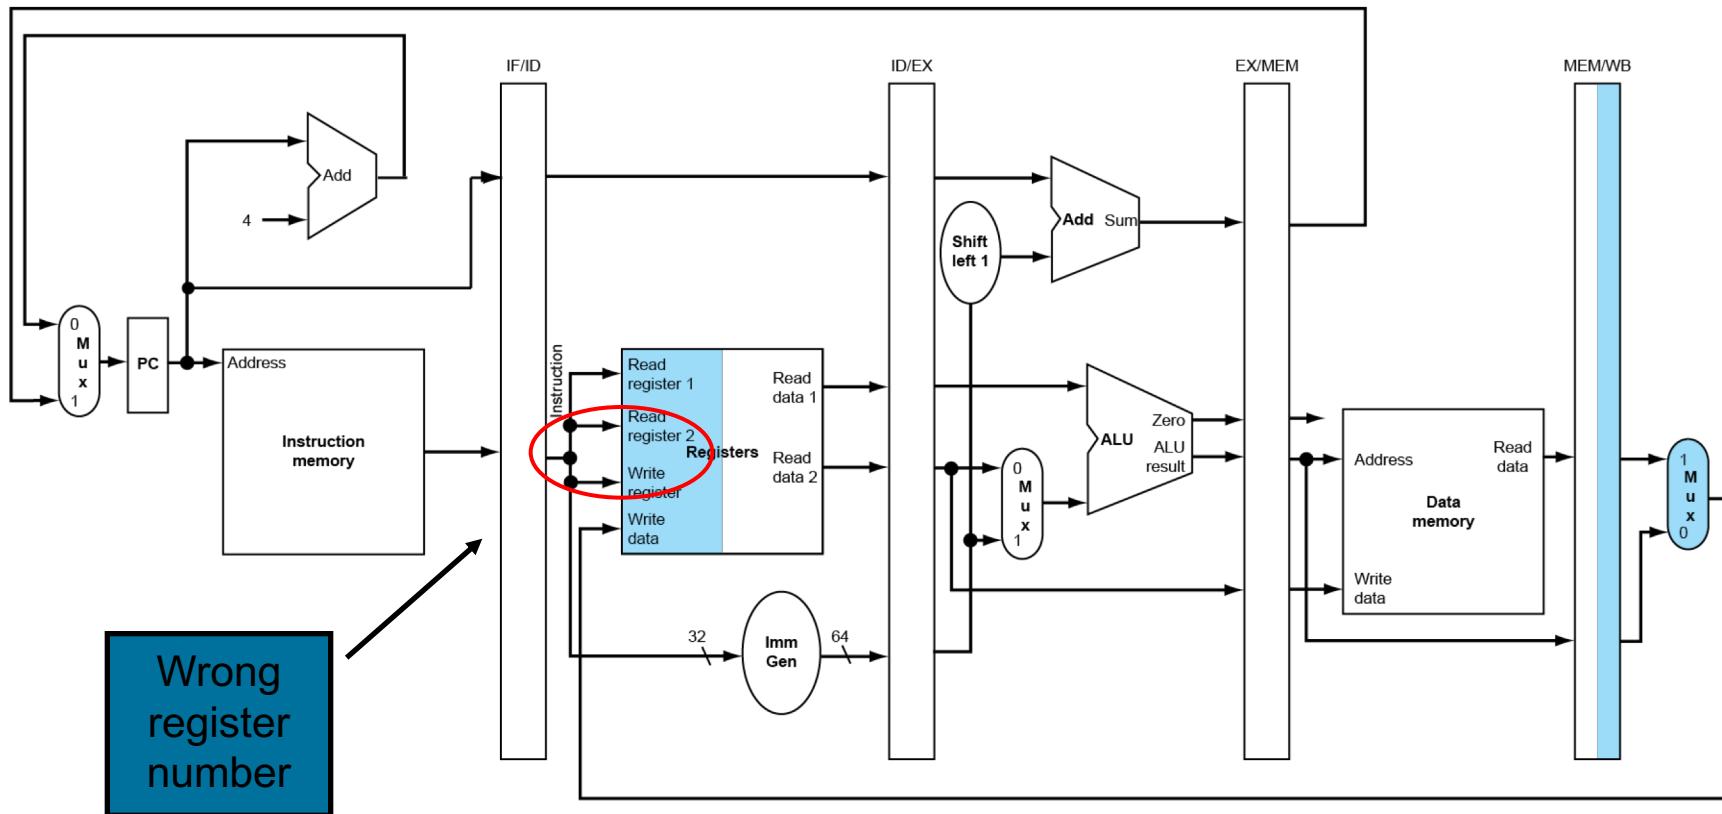
52



# WB for Load

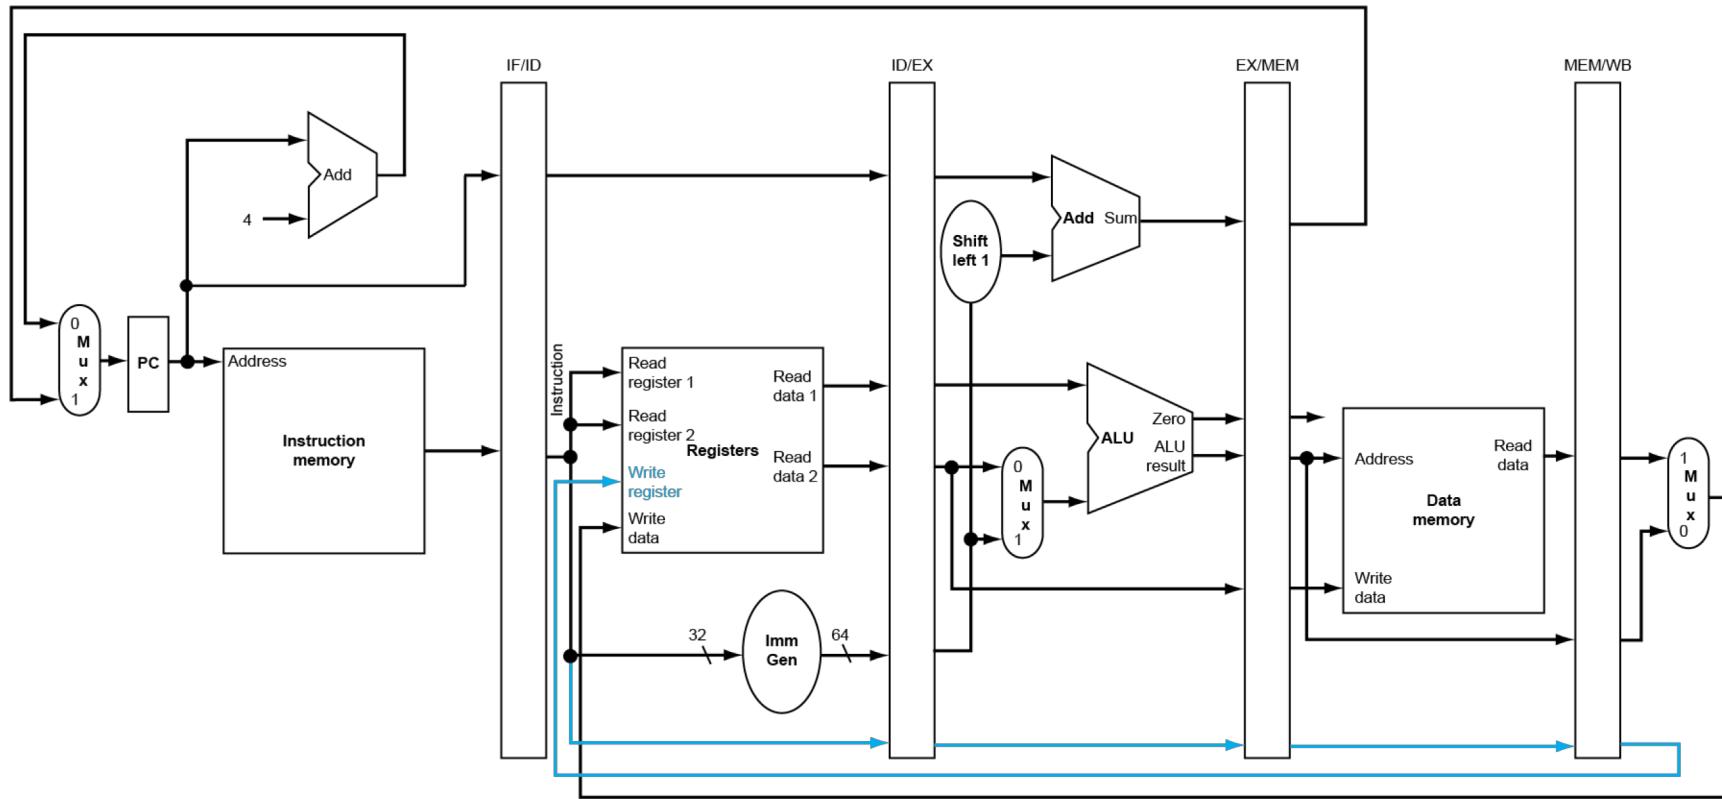
53

Id  
Write-back



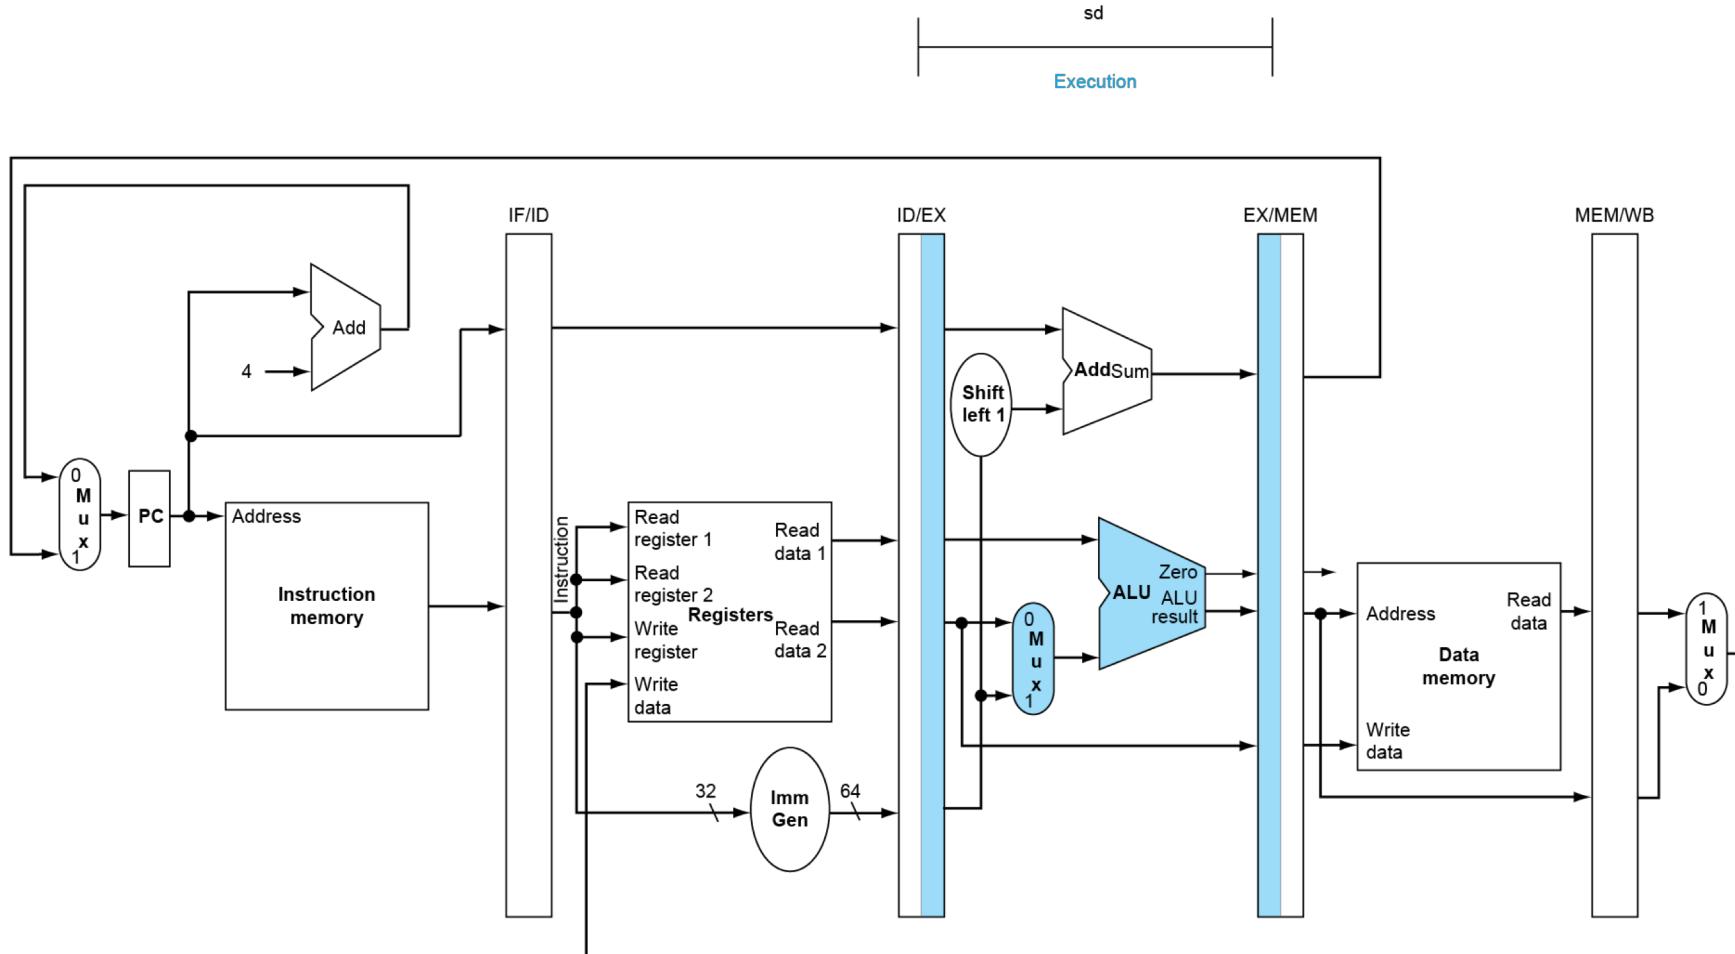
# Corrected Datapath for Load

54



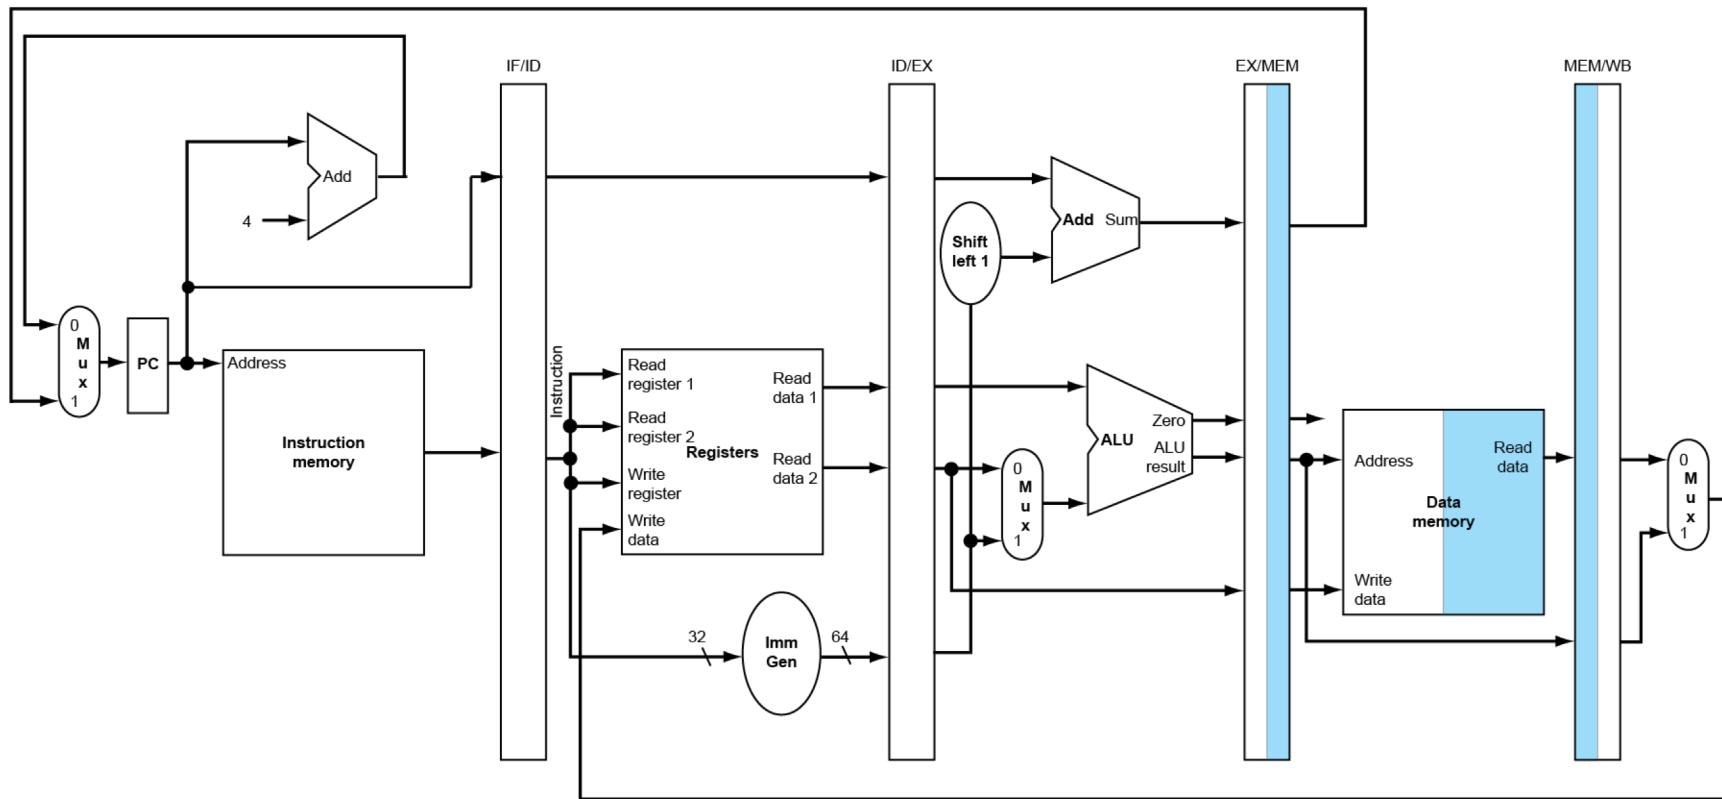
# EX for Store

55



# MEM for Store

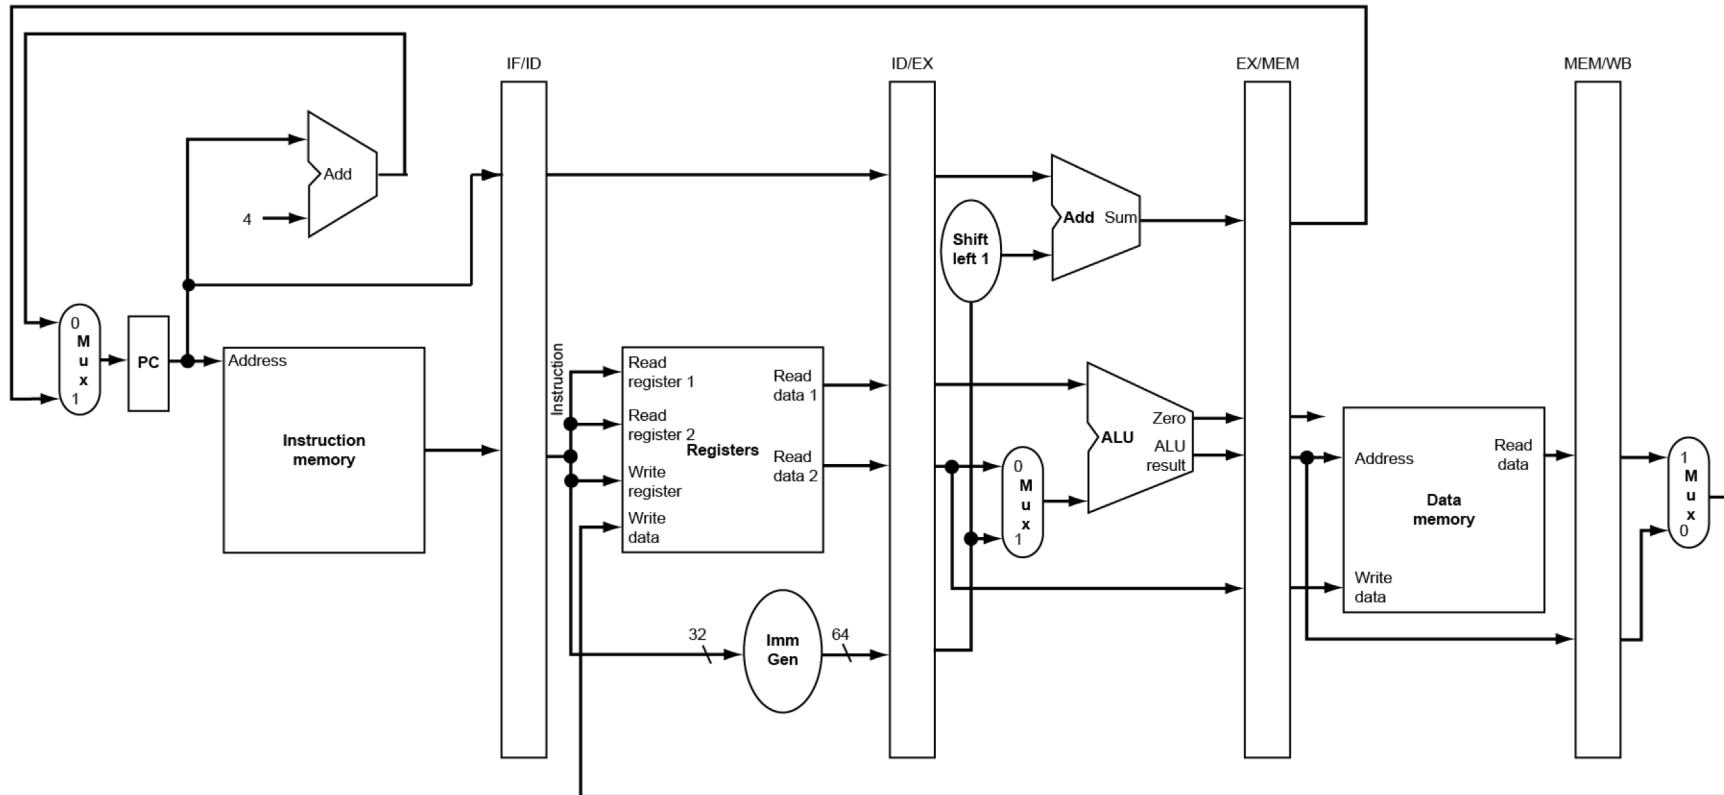
56



# WB for Store

57

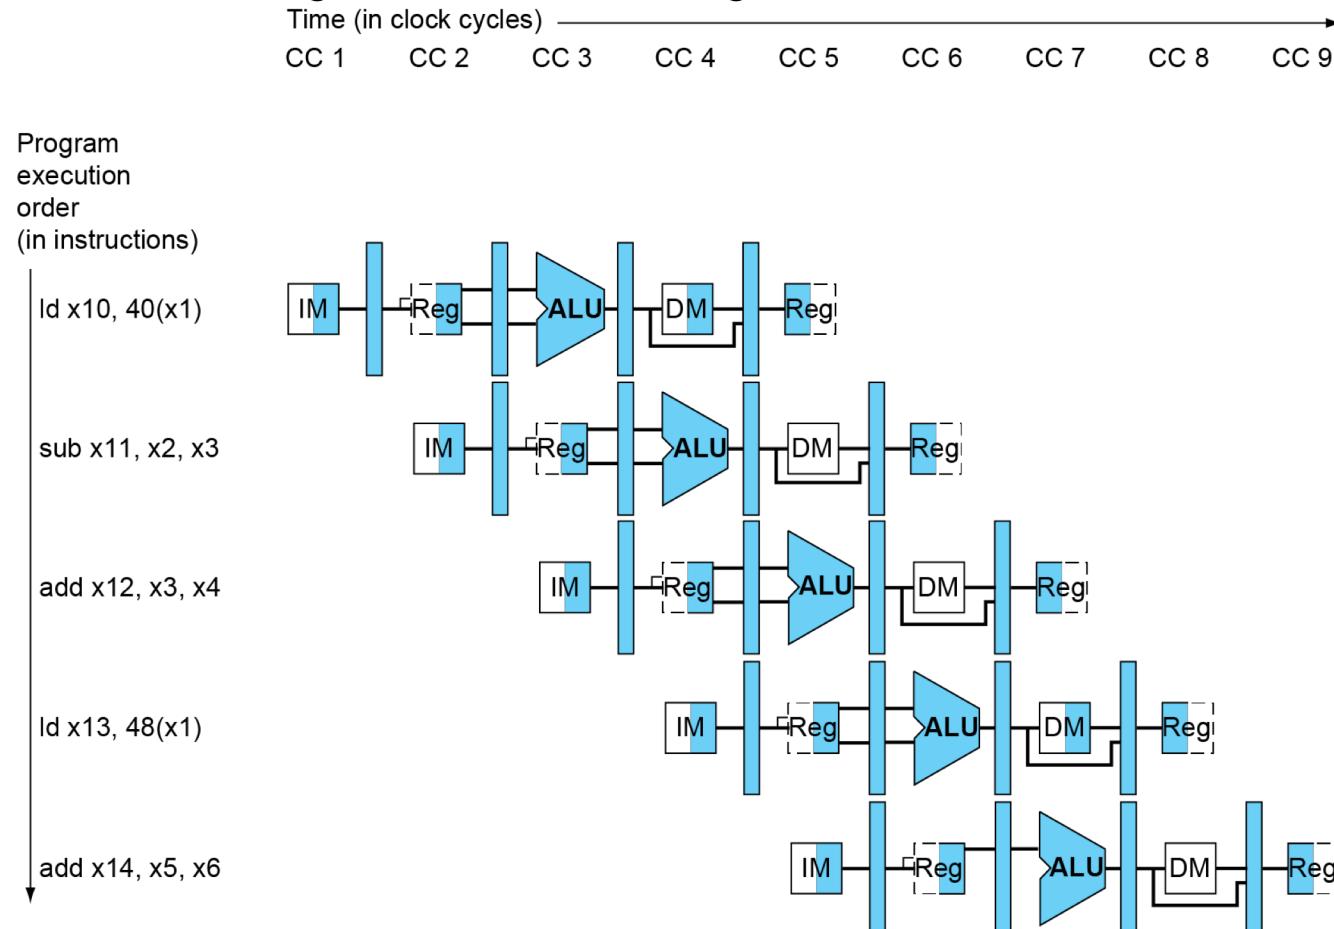
sd  
Write-back



# Multi-Cycle Pipeline Diagram

58

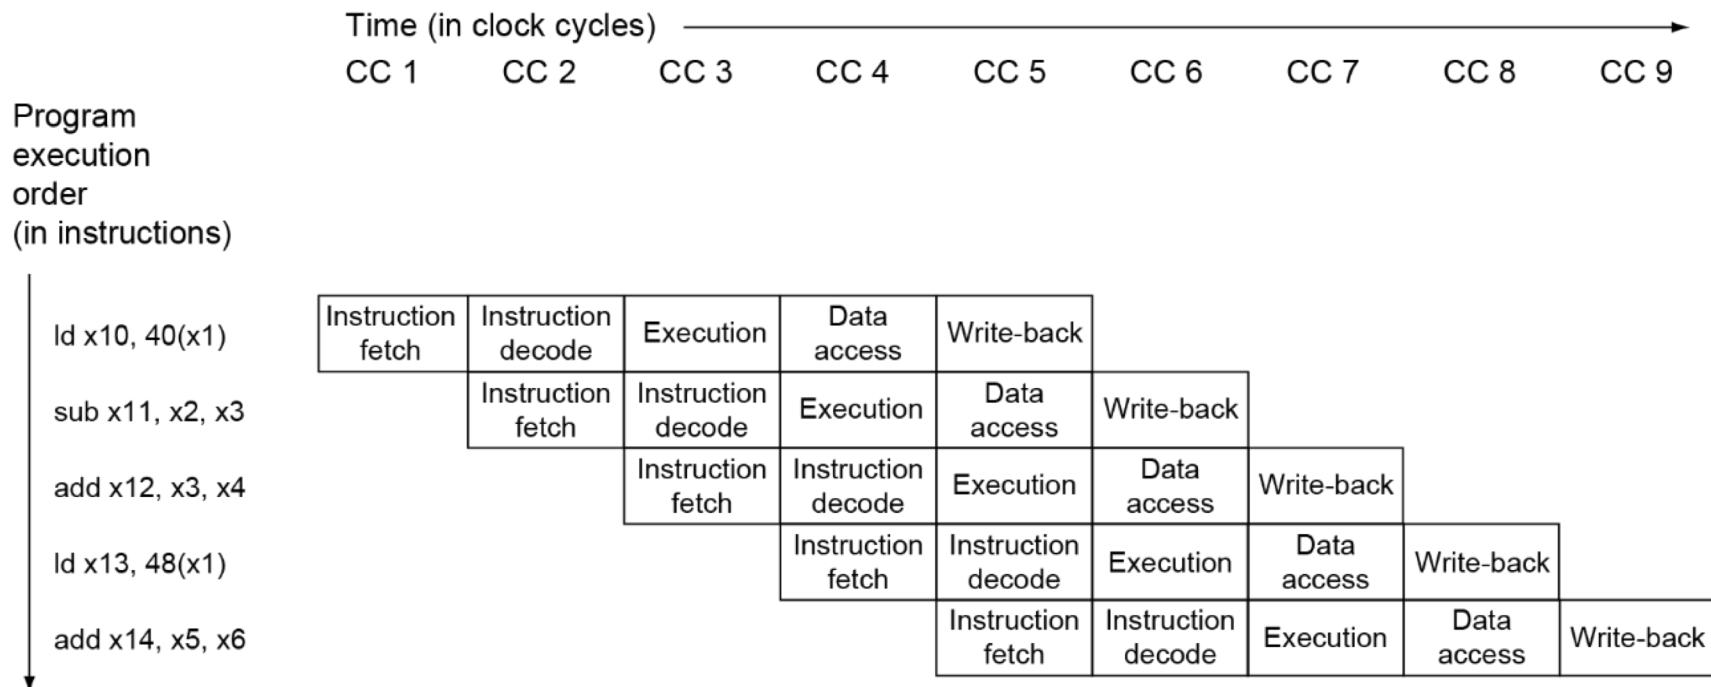
## Form showing resource usage



# Multi-Cycle Pipeline Diagram

59

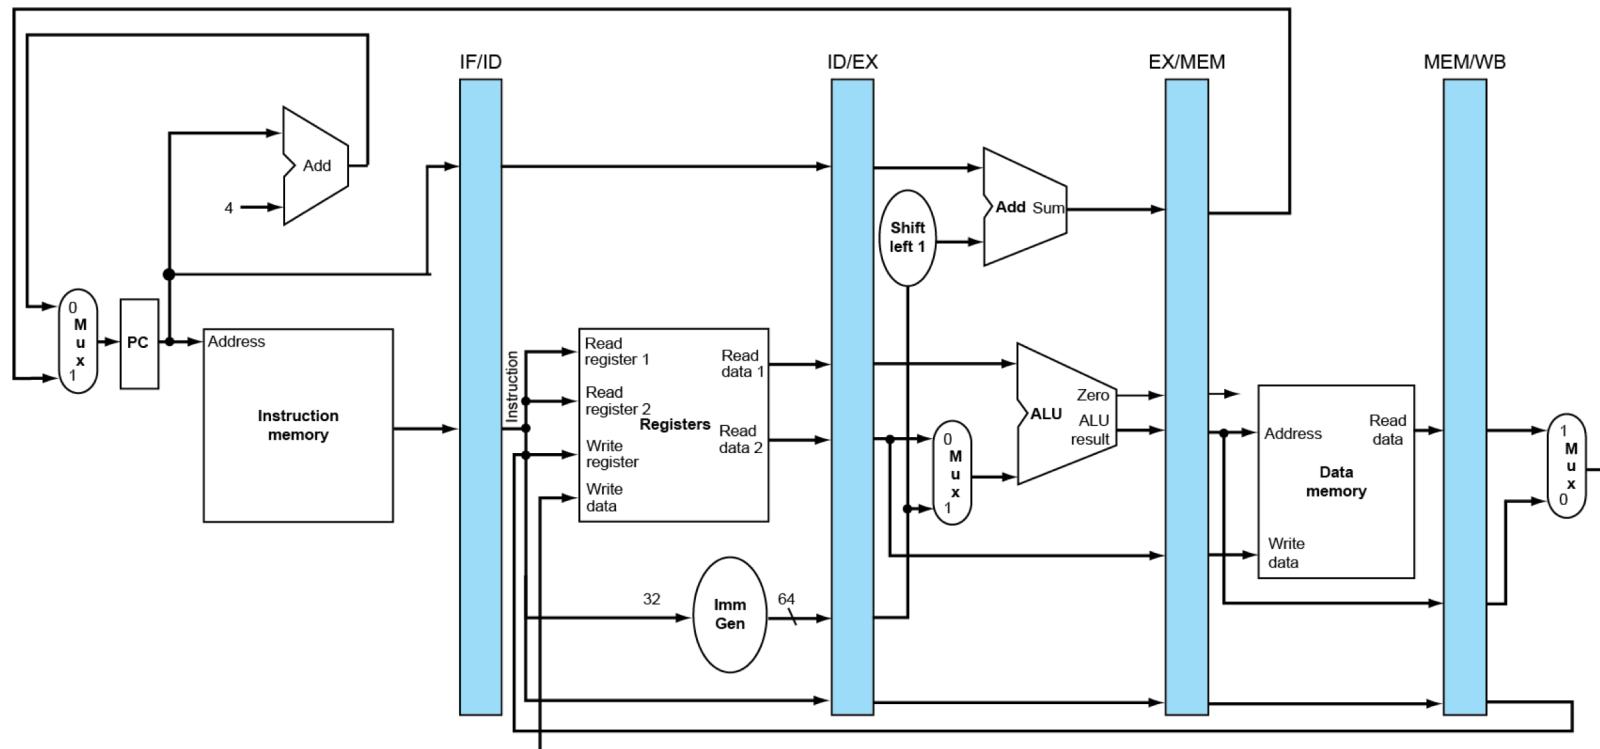
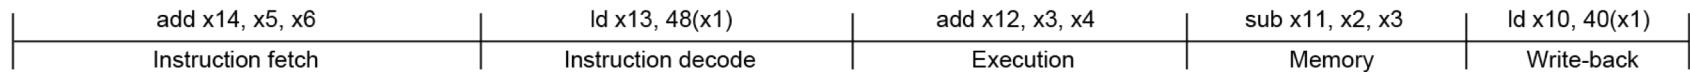
## □ Traditional form



# Single-Cycle Pipeline Diagram

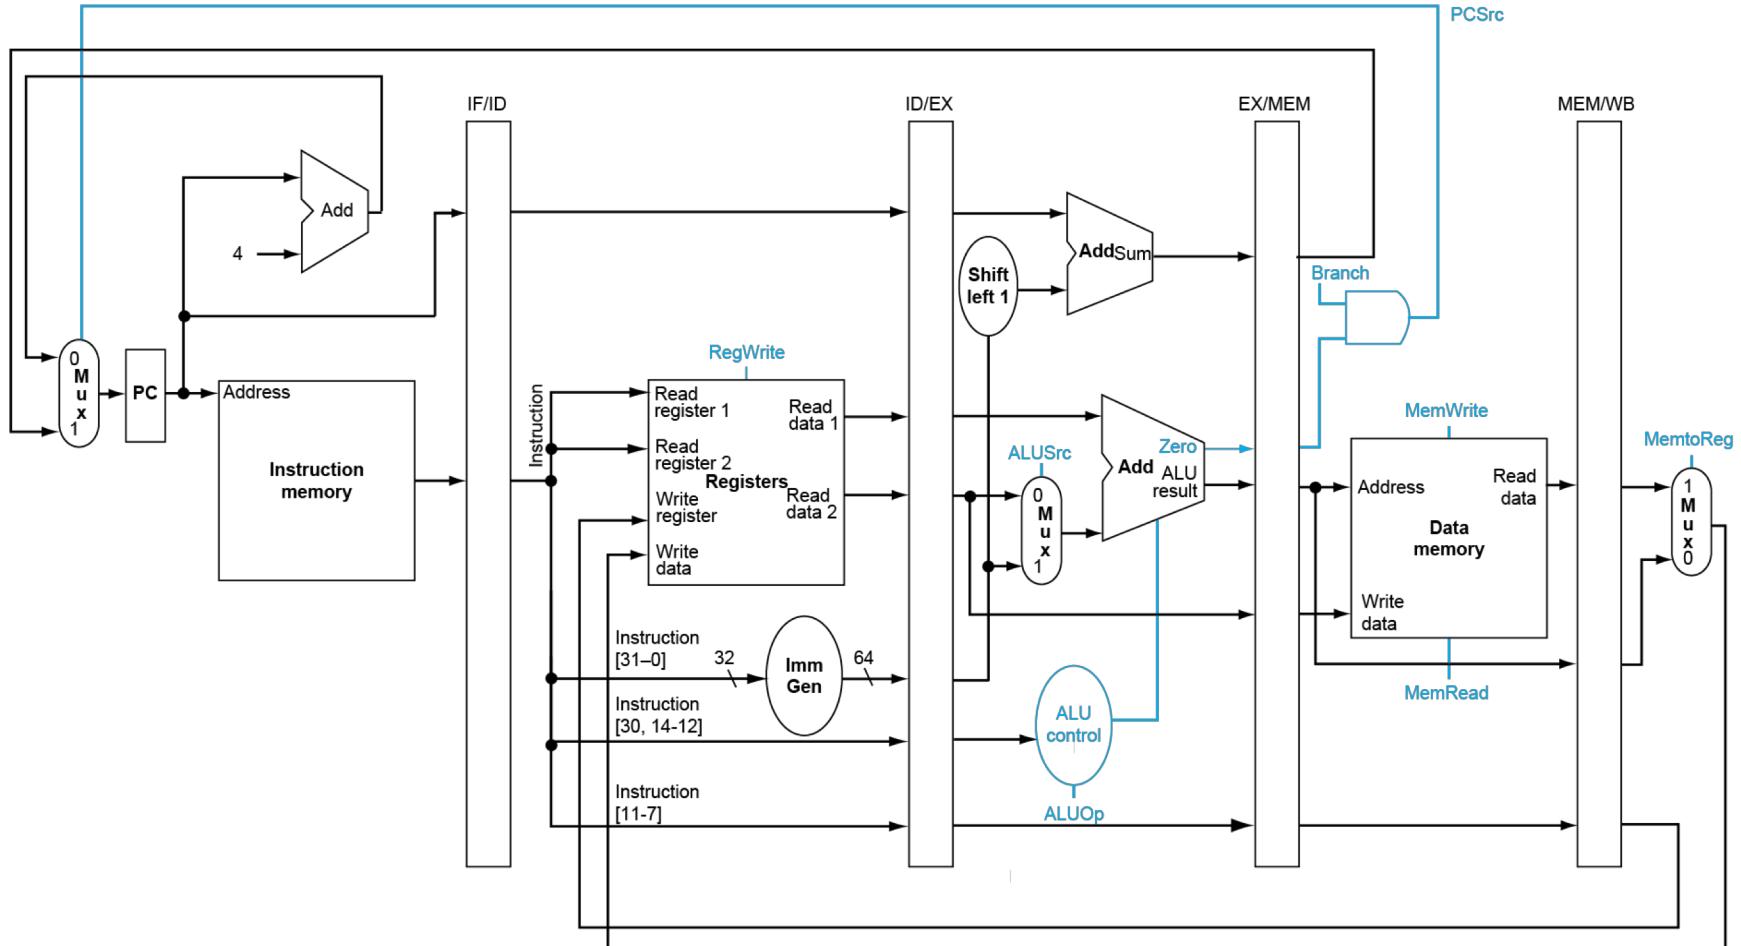
60

## State of pipeline in a given cycle



# Pipelined Control (Simplified)

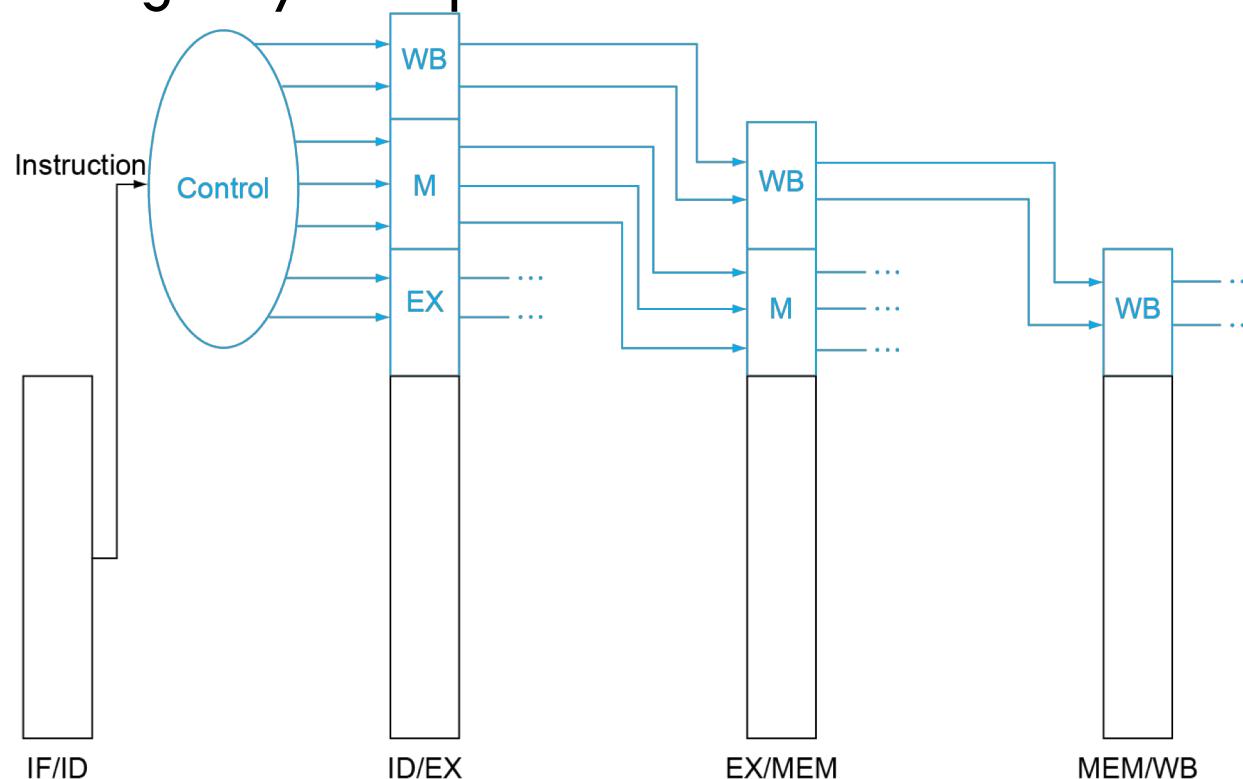
61



# Pipelined Control

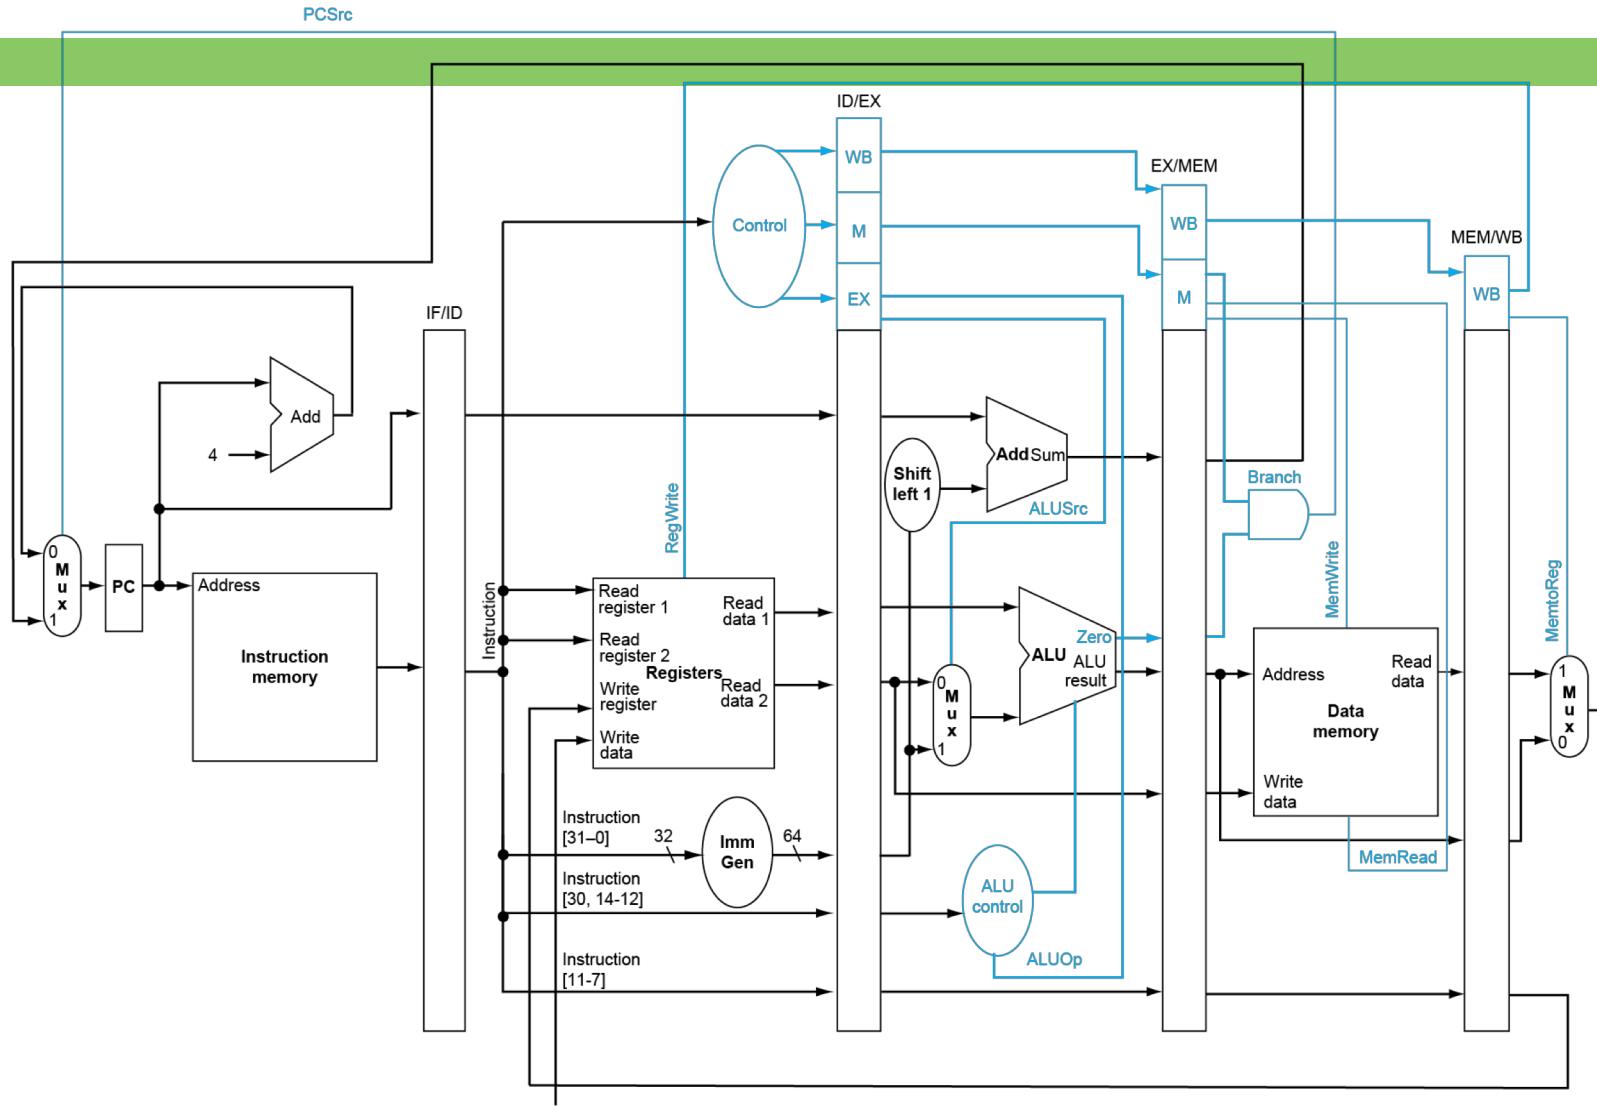
62

- Control signals derived from instruction
  - As in single-cycle implementation



# Pipelined Control

63



# Data Hazards in ALU Instructions

64

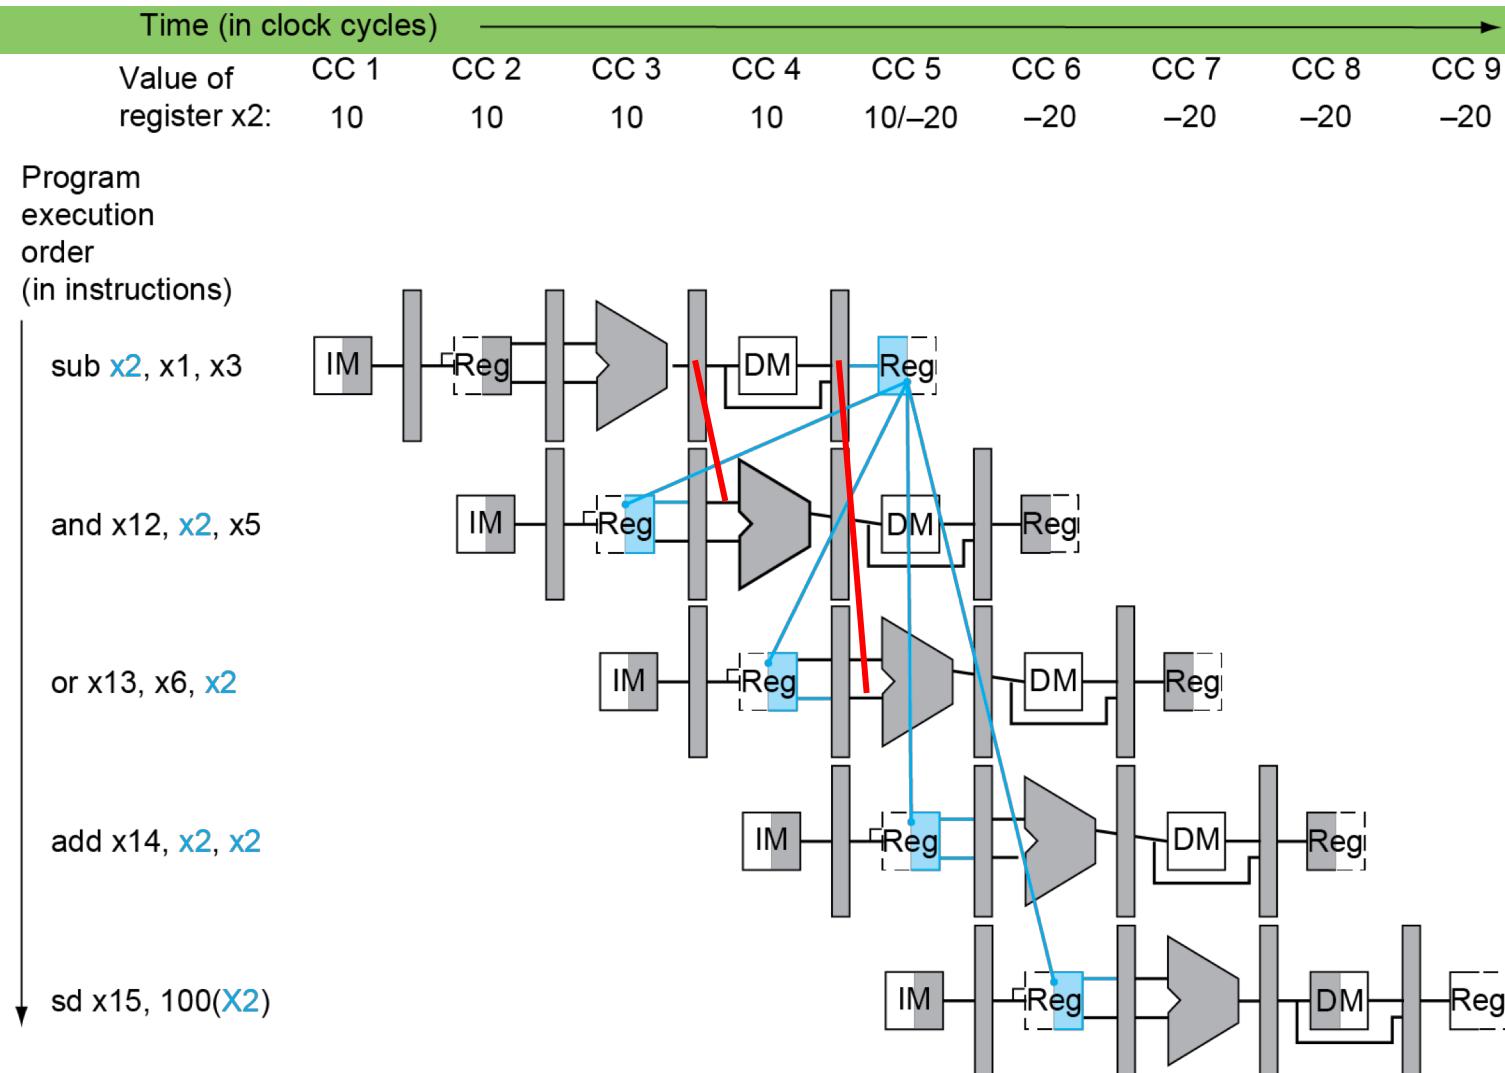
- Consider this sequence:

```
sub    x2, x1,x3  
and    x12, x2, x5  
or     x13, x6, x2  
add    x14, x2, x2  
sd     x15, 100(x2)
```

- We can resolve hazards with forwarding
  - How do we detect when to forward?

# Dependencies & Forwarding

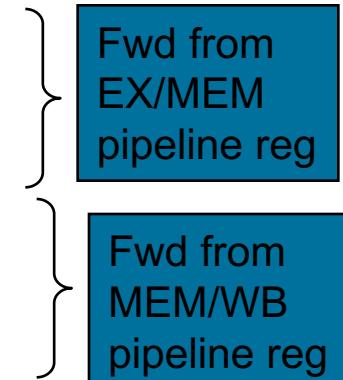
65



# Detecting the Need to Forward

66

- Pass register numbers along pipeline
  - e.g., ID/EX.RegisterRs1 = register number for Rs1 sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
  - ID/EX.RegisterRs1, ID/EX.RegisterRs2
- Data hazards when
  - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs1
  - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRs2
  - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs1
  - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRs2



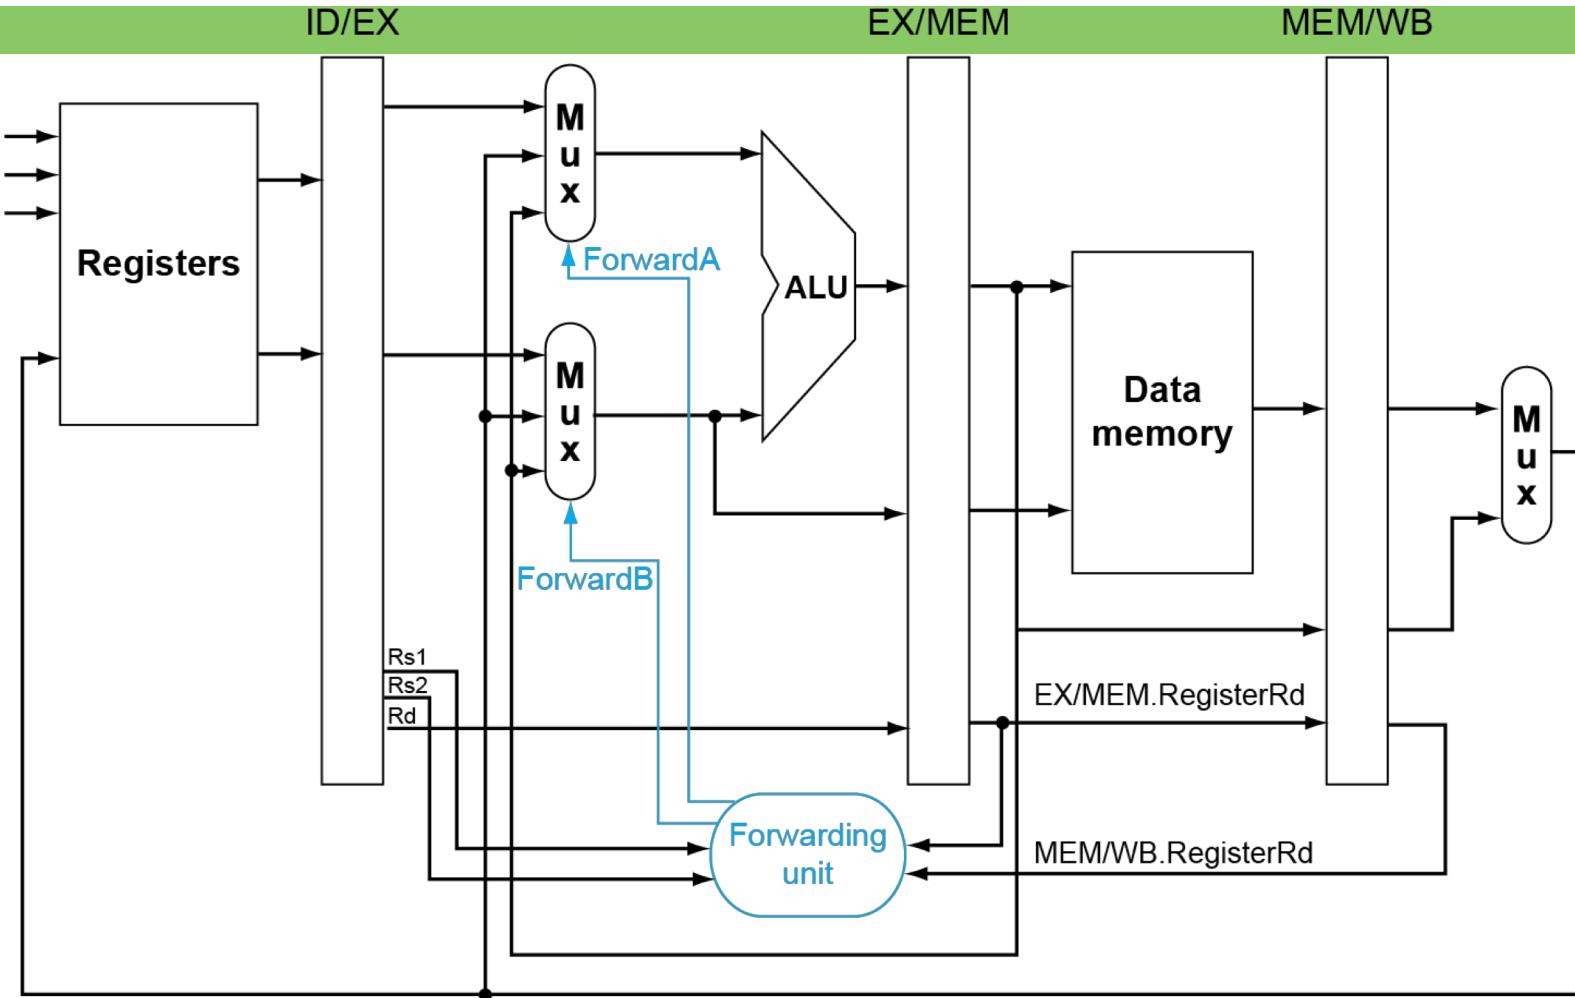
# Detecting the Need to Forward

67

- But only if forwarding instruction will write to a register!
  - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not x0
  - EX/MEM.RegisterRd  $\neq$  0,  
MEM/WB.RegisterRd  $\neq$  0

# Forwarding Paths

68



# Forwarding Conditions

69

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

# Double Data Hazard

70

- Consider the sequence:

add  $x_1, x_1, x_2$

add  $x_1, x_1, x_3$

add  $x_1, x_1, x_4$

- Both hazards occur
  - Want to use the most recent
- Revise MEM hazard condition
  - Only fwd if EX hazard condition isn't true

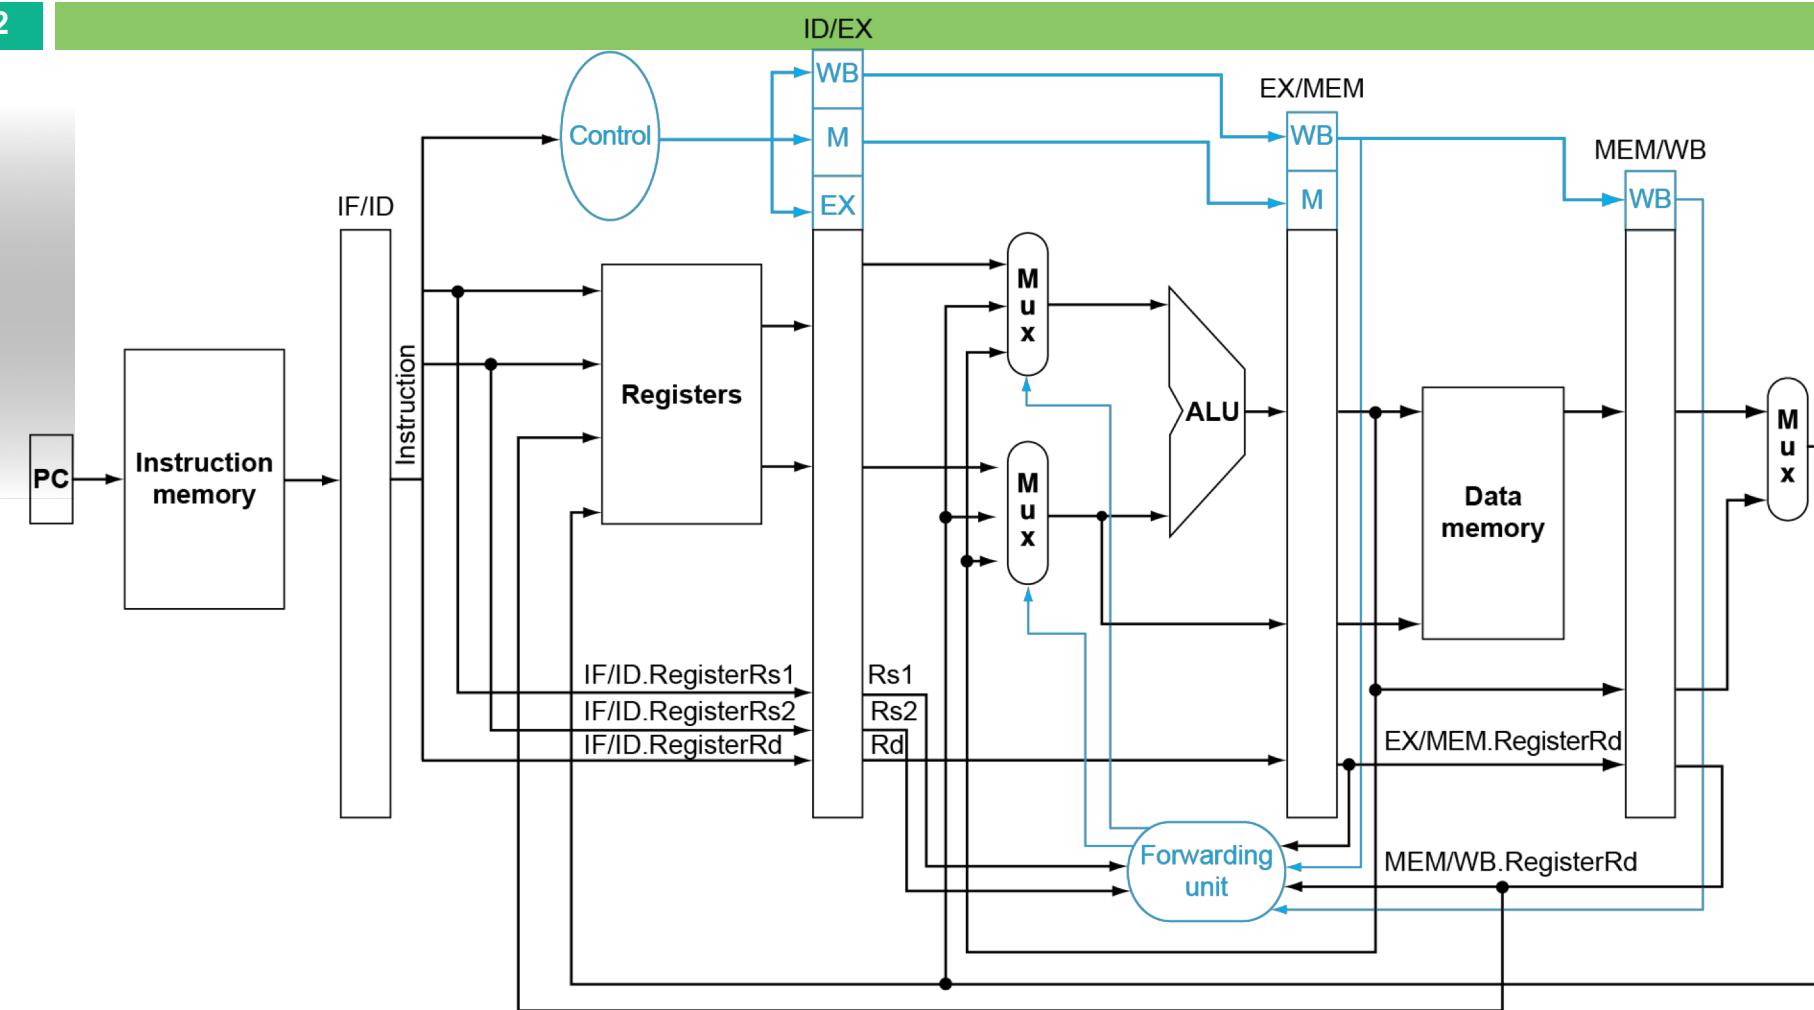
# Revised Forwarding Condition

71

- MEM hazard
  - ▣ if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd ≠ 0)  
and not(EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRs1))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs1)) ForwardA = 01
  - ▣ if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd ≠ 0)  
and not(EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRs2))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs2)) ForwardB = 01

# Datapath with Forwarding

72



# Load-Use Hazard Detection

73

- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
  - IF/ID.RegisterRs1, IF/ID.RegisterRs2
- Load-use hazard when
  - ID/EX.MemRead and
    - ((ID/EX.RegisterRd = IF/ID.RegisterRs1) or (ID/EX.RegisterRd = IF/ID.RegisterRs2))
- If detected, stall and insert bubble

# How to Stall the Pipeline

74

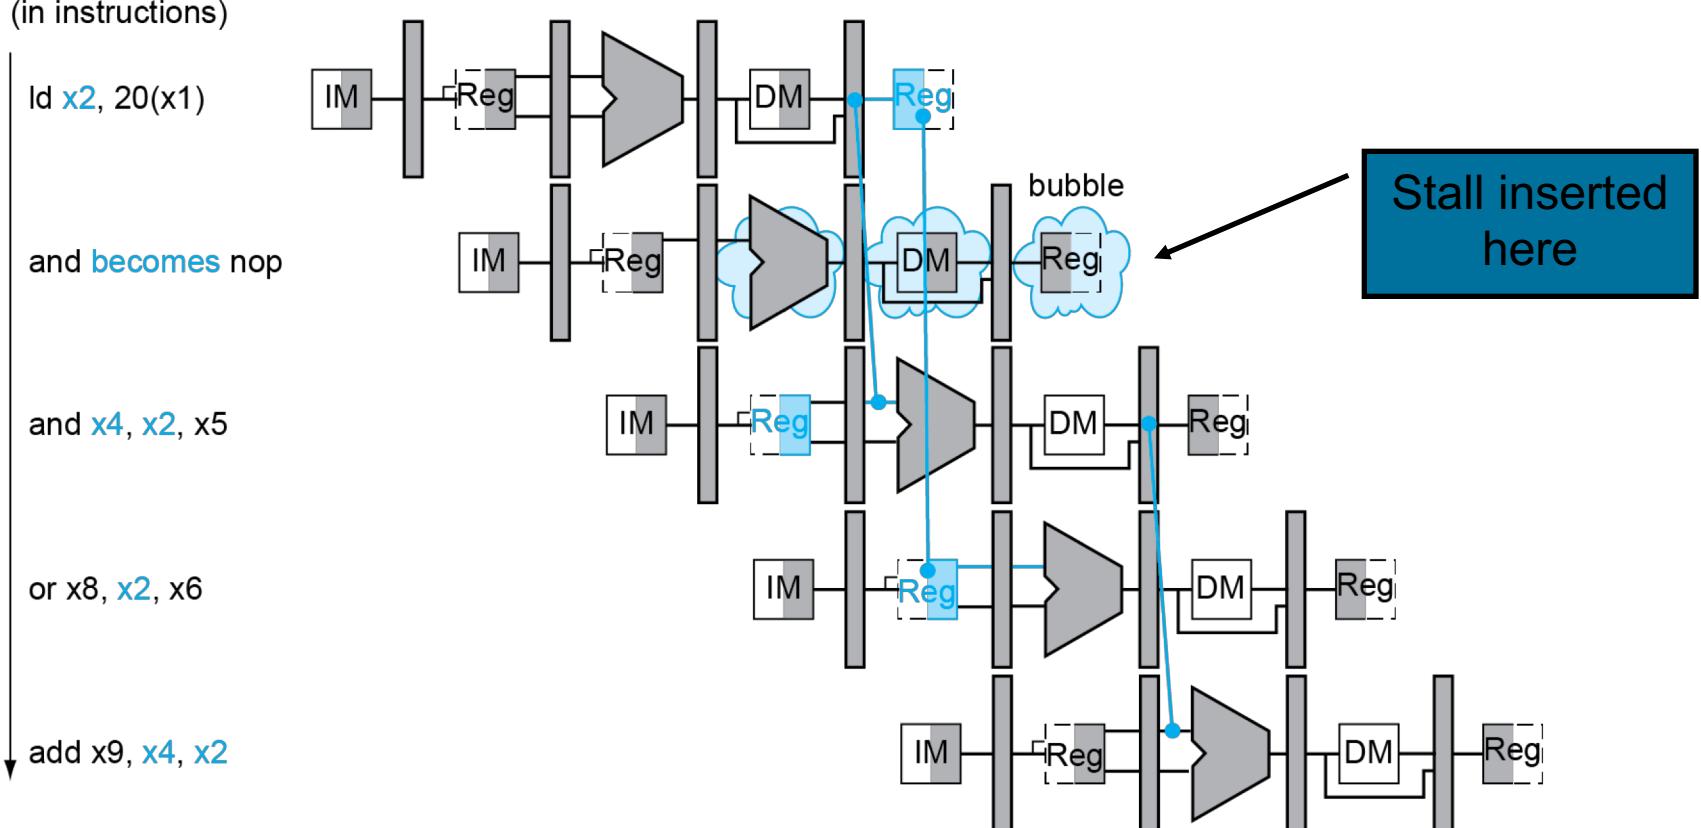
- Force control values in ID/EX register to 0
  - ▣ EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
  - ▣ Using instruction is decoded again
  - ▣ Following instruction is fetched again
  - ▣ 1-cycle stall allows MEM to read data for 1d
    - Can subsequently forward to EX stage

# Load-Use Data Hazard

75

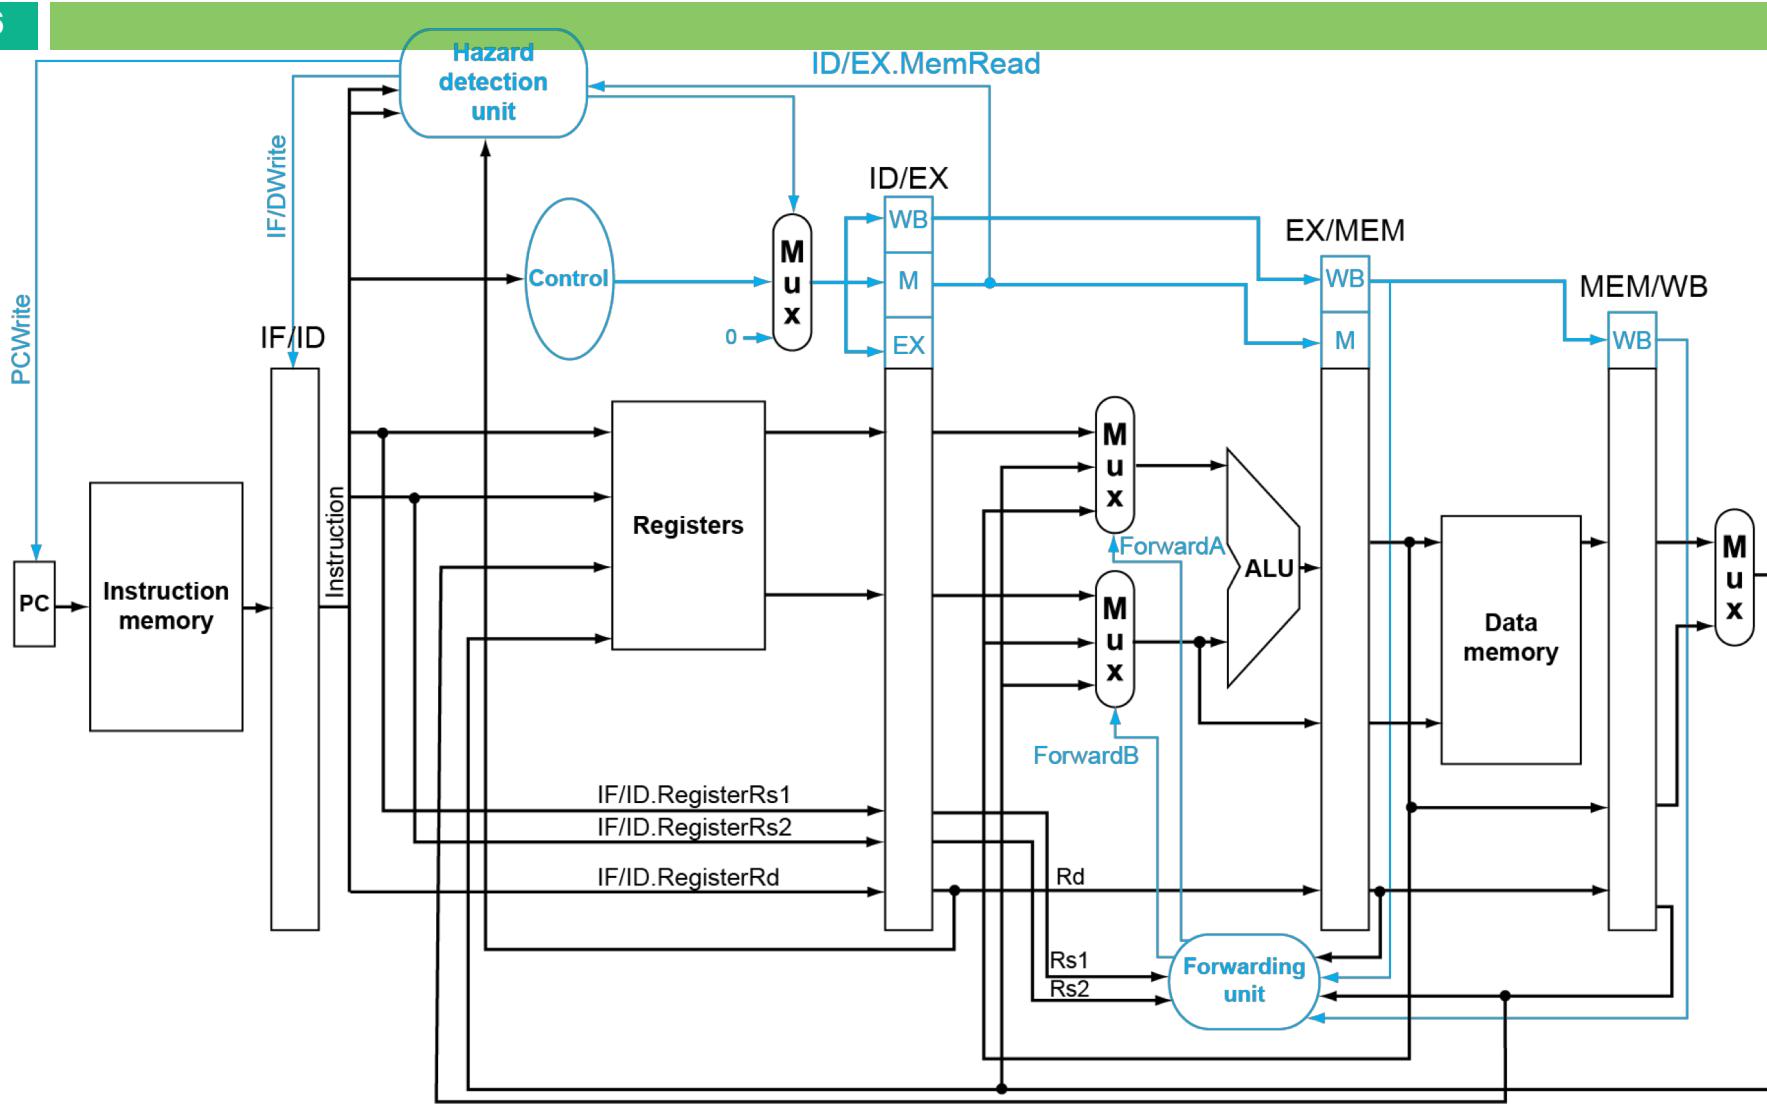
Time (in clock cycles) →  
CC 1 CC 2 CC 3 CC 4 CC 5 CC 6 CC 7 CC 8 CC 9 CC 10

Program  
execution  
order  
(in instructions)



# Datapath with Hazard Detection

76



# Stalls and Performance

77

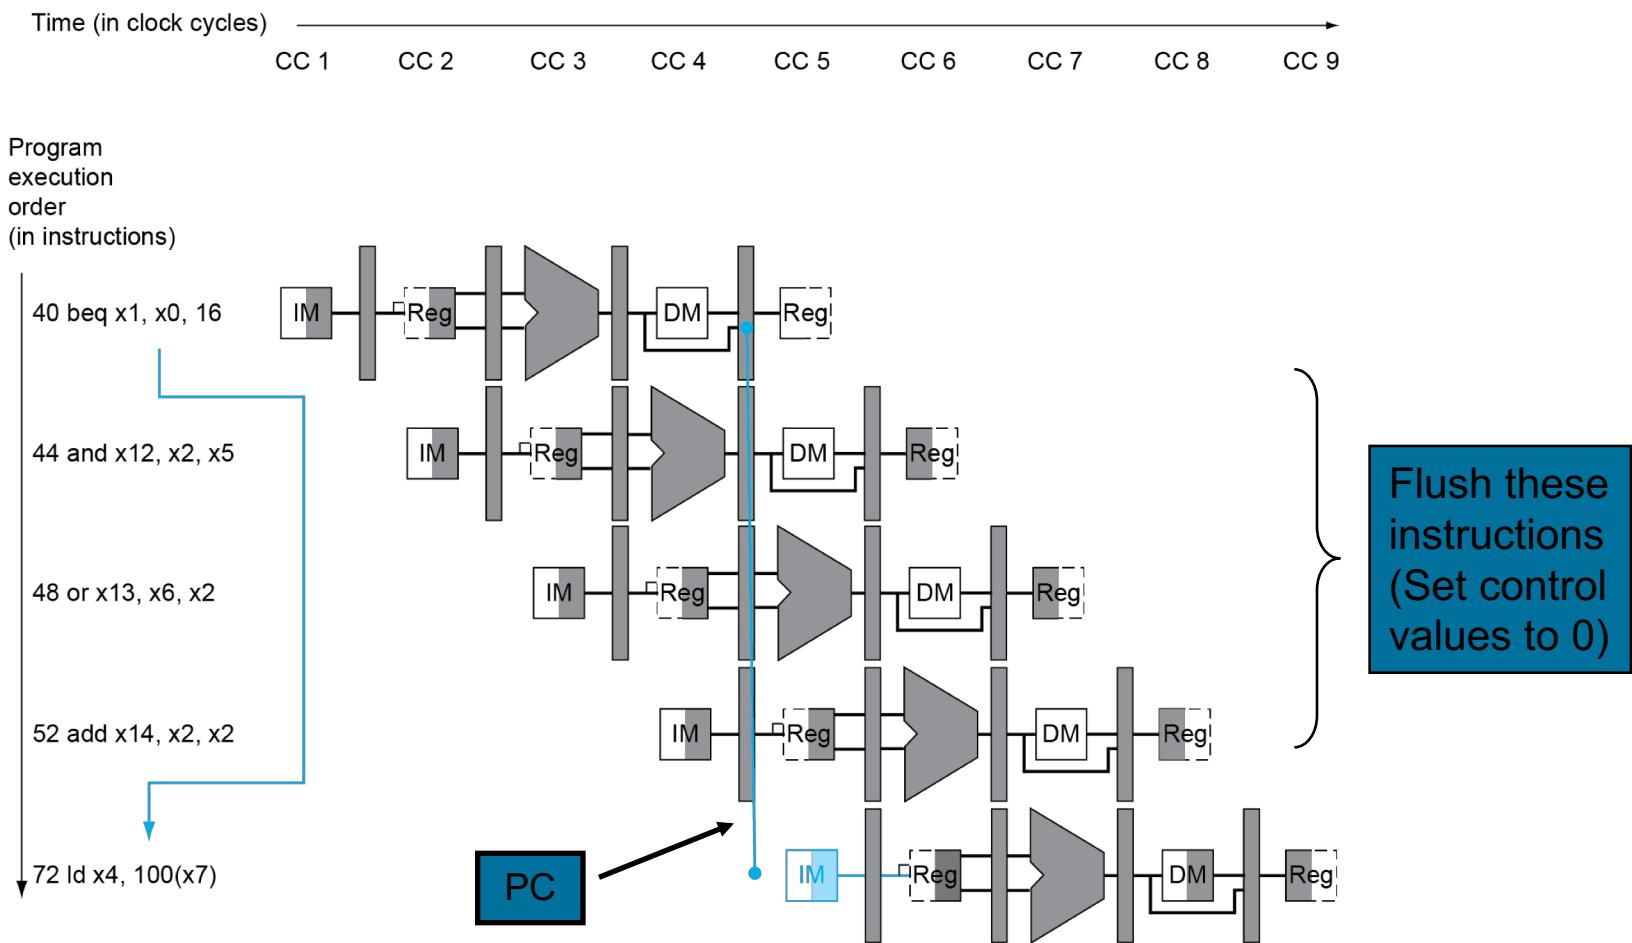
## The BIG Picture

- Stalls reduce performance
  - ▣ But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
  - ▣ Requires knowledge of the pipeline structure

# Branch Hazards

78

## □ If branch outcome determined in MEM



# Reducing Branch Delay

79

- Move hardware to determine outcome to ID stage
  - Target address adder
  - Register comparator
- Example: branch taken

```
36: sub x10, x4, x8
40: beq x1, x3, 16 // PC-relative branch
               // to 40+16*2=72
44: and x12, x2, x5
48: orr x13, x2, x6
52: add x14, x4, x2
56: sub x15, x6, x7
      ...
72: id  x4, 50(x7)
```

# Example: Branch Taken

80

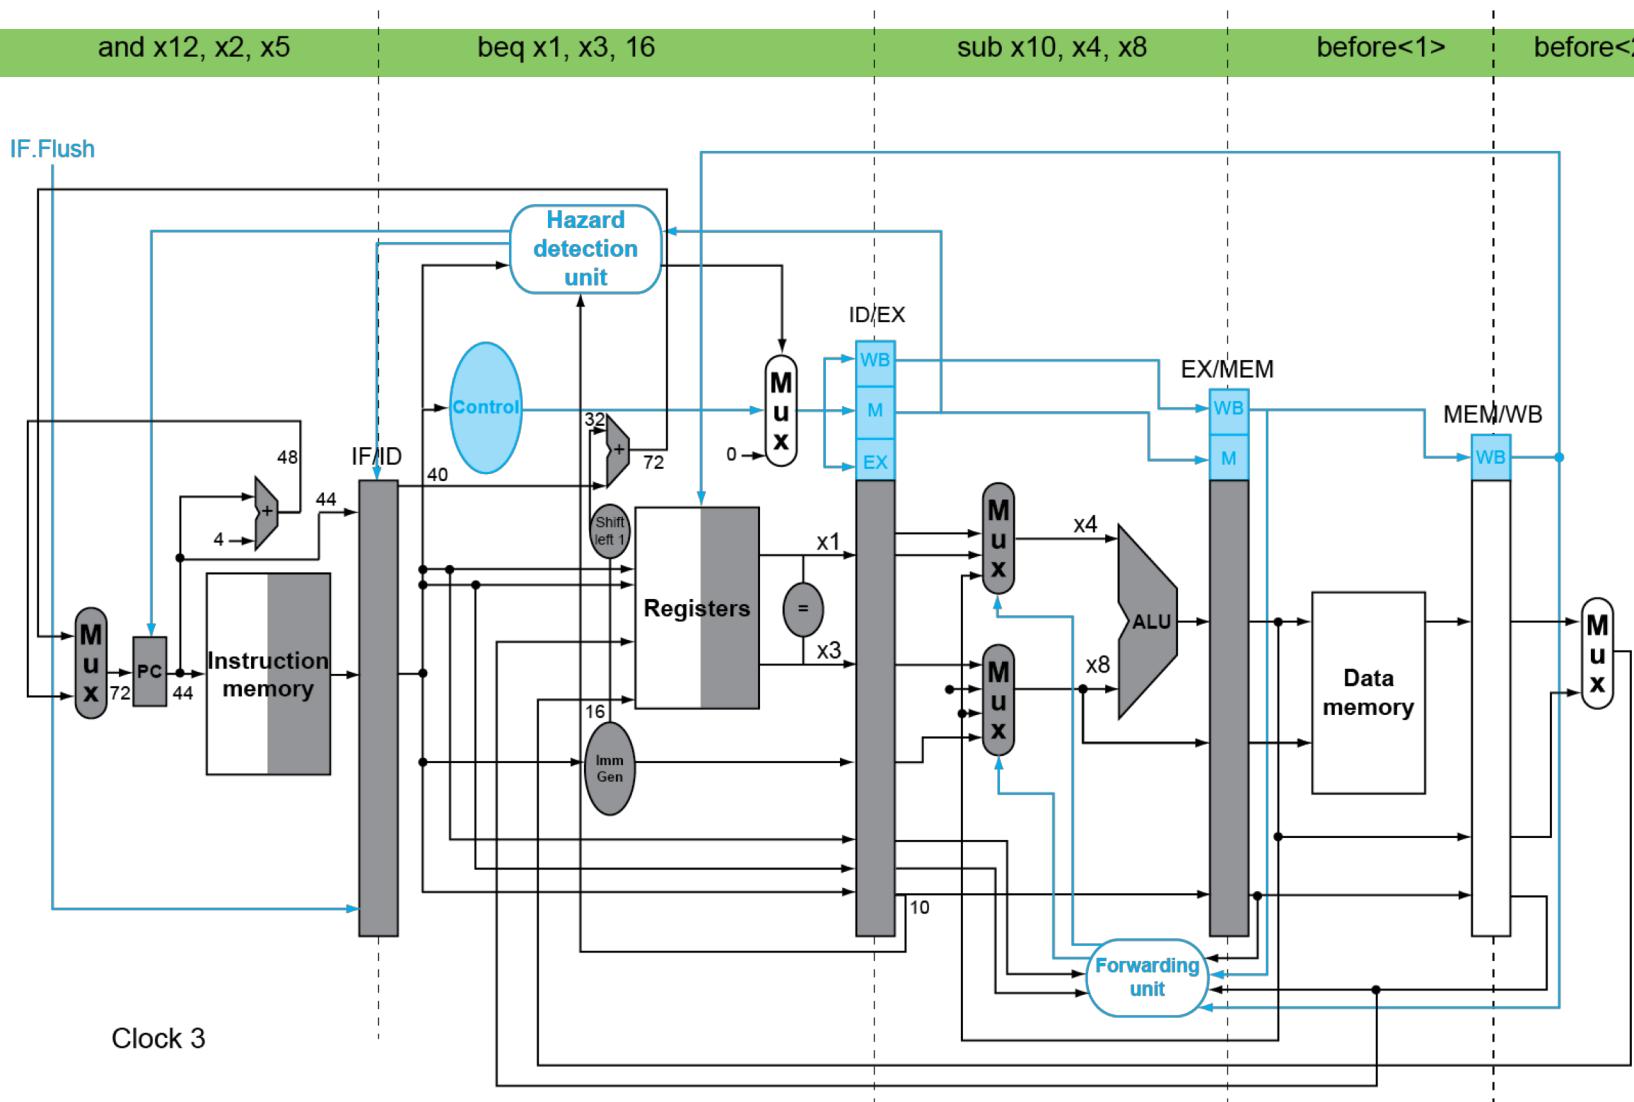
and x12, x2, x5

beq x1, x3, 16

sub x10, x4, x8

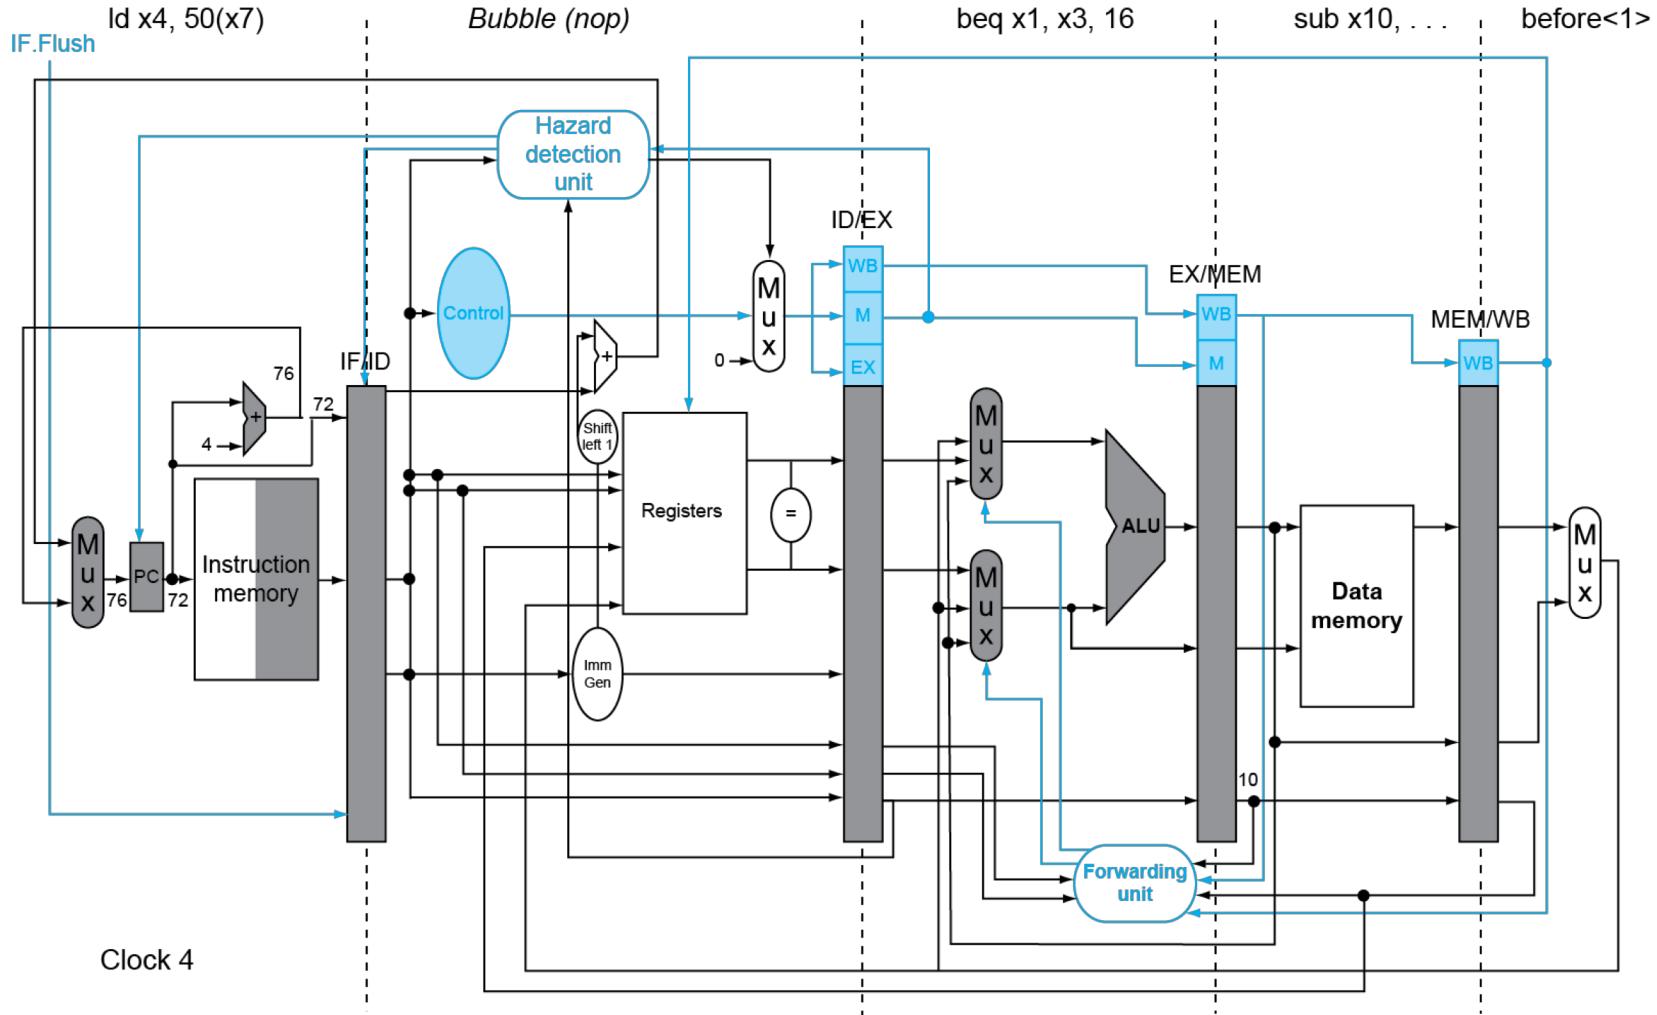
before<1>

before<2>



# Example: Branch Taken

81



# Dynamic Branch Prediction

82

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
  - Branch prediction buffer (aka branch history table)
  - Indexed by recent branch instruction addresses
  - Stores outcome (taken/not taken)
  - To execute a branch
    - Check table, expect the same outcome
    - Start fetching from fall-through or target
    - If wrong, flush pipeline and flip prediction

# 1-Bit Predictor: Shortcoming

83

- Inner loop branches mispredicted twice!

outer: ...



inner: ...



beq ... , ... , inner

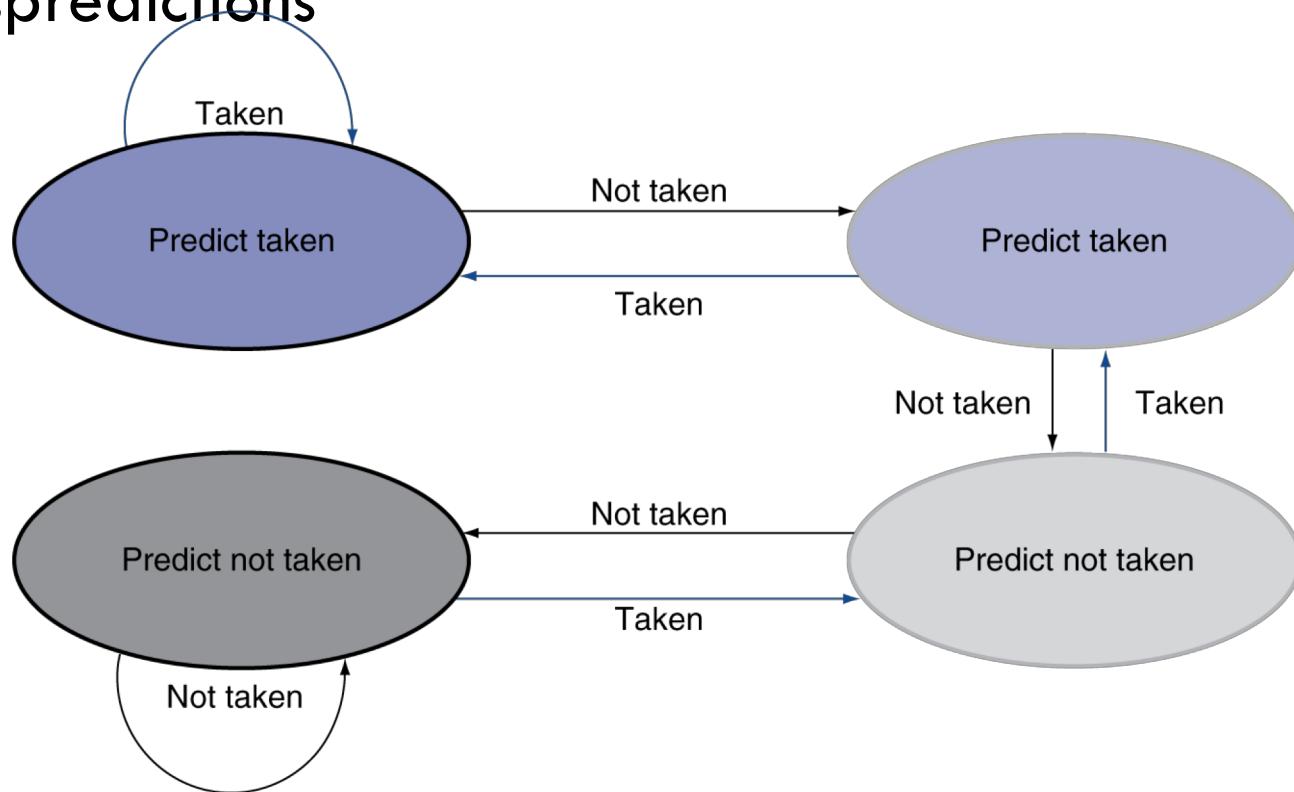
...  
beq ... , ... , outer

- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

# 2-Bit Predictor

84

- Only change prediction on two successive mispredictions



# Calculating the Branch Target

85

- Even with predictor, still need to calculate the target address
  - ▣ 1-cycle penalty for a taken branch
- Branch target buffer
  - ▣ Cache of target addresses
  - ▣ Indexed by PC when instruction fetched
    - If hit and instruction is branch predicted taken, can fetch target immediately

# Exceptions and Interrupts

86

- “Unexpected” events requiring change in flow of control
  - ▣ Different ISAs use the terms differently
- Exception
  - ▣ Arises within the CPU
    - e.g., undefined opcode, syscall, ...
- Interrupt
  - ▣ From an external I/O controller
- Dealing with them without sacrificing performance is hard

# Handling Exceptions

87

- Save PC of offending (or interrupted) instruction
  - ▣ In RISC-V: Supervisor Exception Program Counter (SEPC)
- Save indication of the problem
  - ▣ In RISC-V: Supervisor Exception Cause Register (SCAUSE)
    - ▣ 64 bits, but most bits unused
      - Exception code field: 2 for undefined opcode, 12 for hardware malfunction, ...
- Jump to handler
  - ▣ Assume at 0000 0000 1C09 0000<sub>hex</sub>

# An Alternate Mechanism

88

- Vectored Interrupts
  - Handler address determined by the cause
- Exception vector address to be added to a vector table base register:
  - Undefined opcode 00 0100 0000<sub>two</sub>
  - Hardware malfunction: 01 1000 0000<sub>two</sub>
  - ....
- Instructions either
  - Deal with the interrupt, or
  - Jump to real handler

# Handler Actions

89

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
  - ▣ Take corrective action
  - ▣ use SEPC to return to program
- Otherwise
  - ▣ Terminate program
  - ▣ Report error using SEPC, SCAUSE, ...

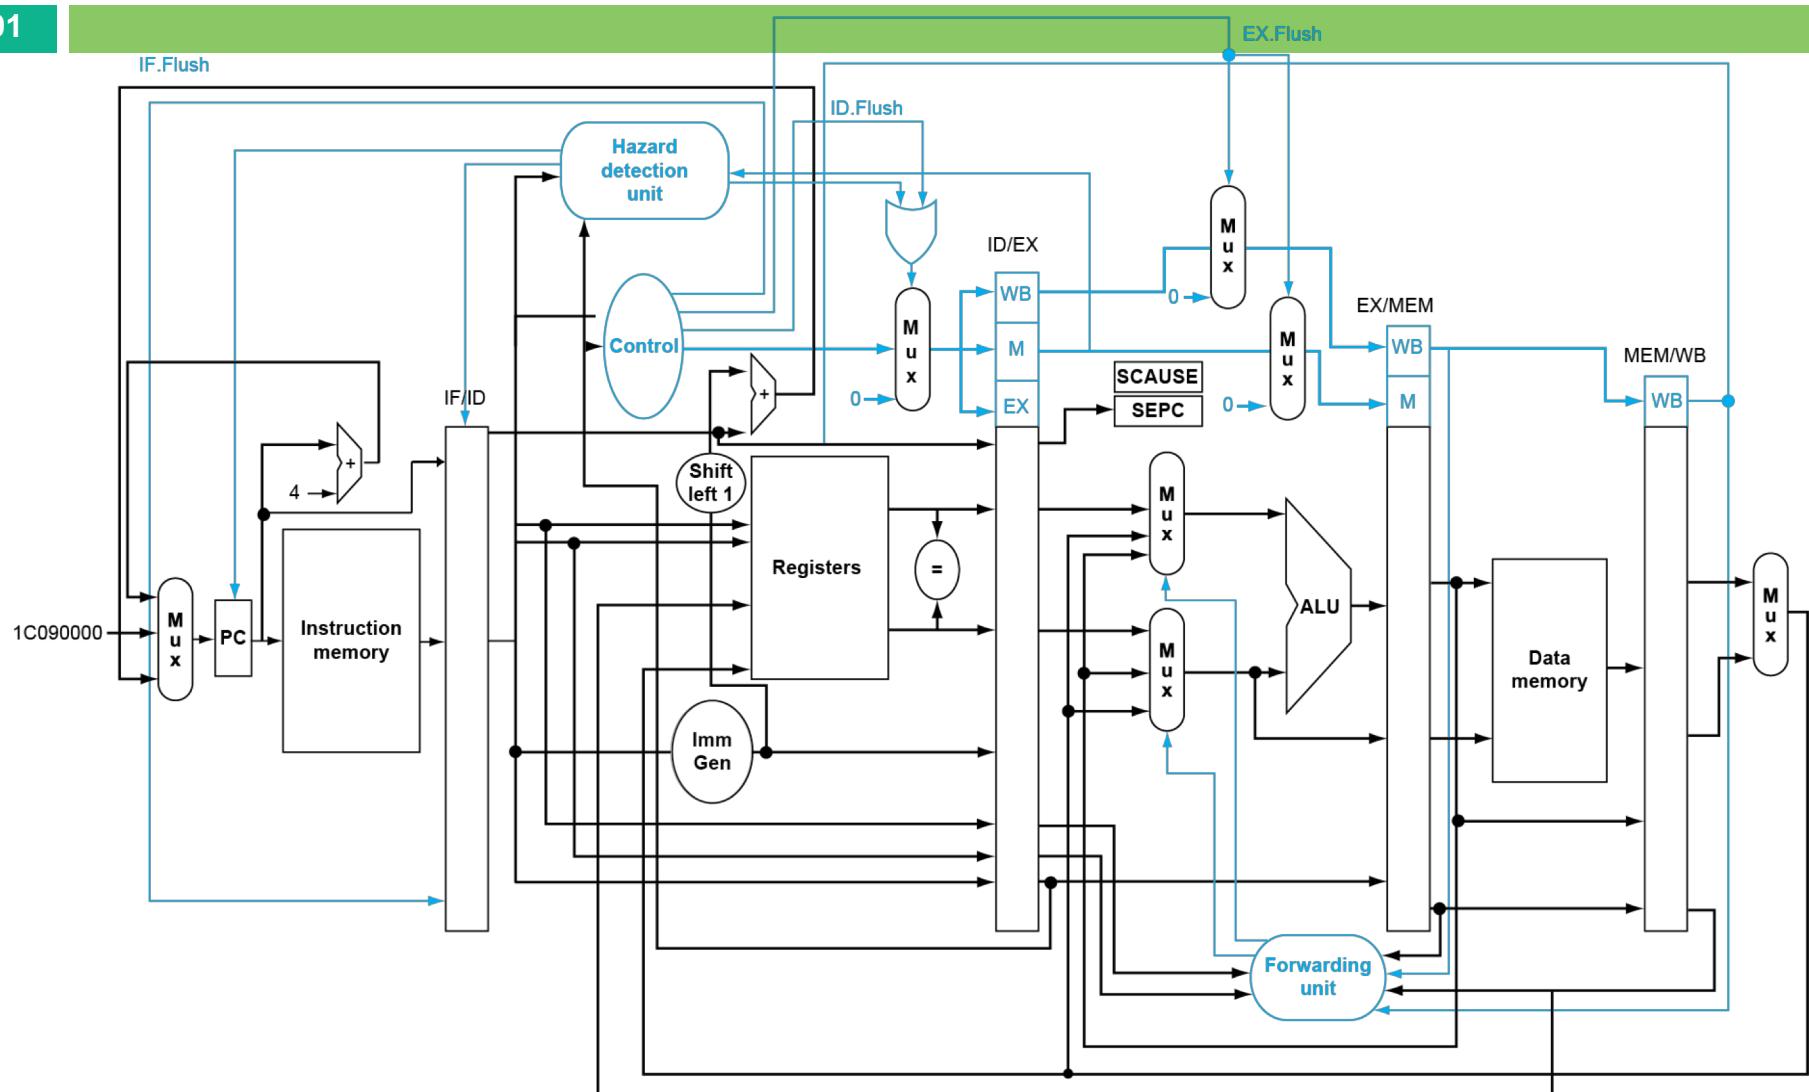
# Exceptions in a Pipeline

90

- Another form of control hazard
- Consider malfunction on add in EX stage
  - add  $x_1, x_2, x_1$ 
    - ▣ Prevent  $x_1$  from being clobbered
    - ▣ Complete previous instructions
    - ▣ Flush add and subsequent instructions
    - ▣ Set SEPC and SCAUSE register values
    - ▣ Transfer control to handler
  - Similar to mispredicted branch
    - ▣ Use much of the same hardware

# Pipeline with Exceptions

91



# Exception Properties

92

- Restartable exceptions
  - Pipeline can flush the instruction
  - Handler executes, then returns to the instruction
    - Refetched and executed from scratch
- PC saved in SEPC register
  - Identifies causing instruction

# Exception Example

93

## □ Exception on **add** in

```
40      sub    x11, x2, x4  
44      and    x12, x2, x5  
48      orr    x13, x2, x6  
4C      add    x1,  x2,  x1  
50      sub    x15, x6, x7  
54      ld     x16, 100(x7)
```

...

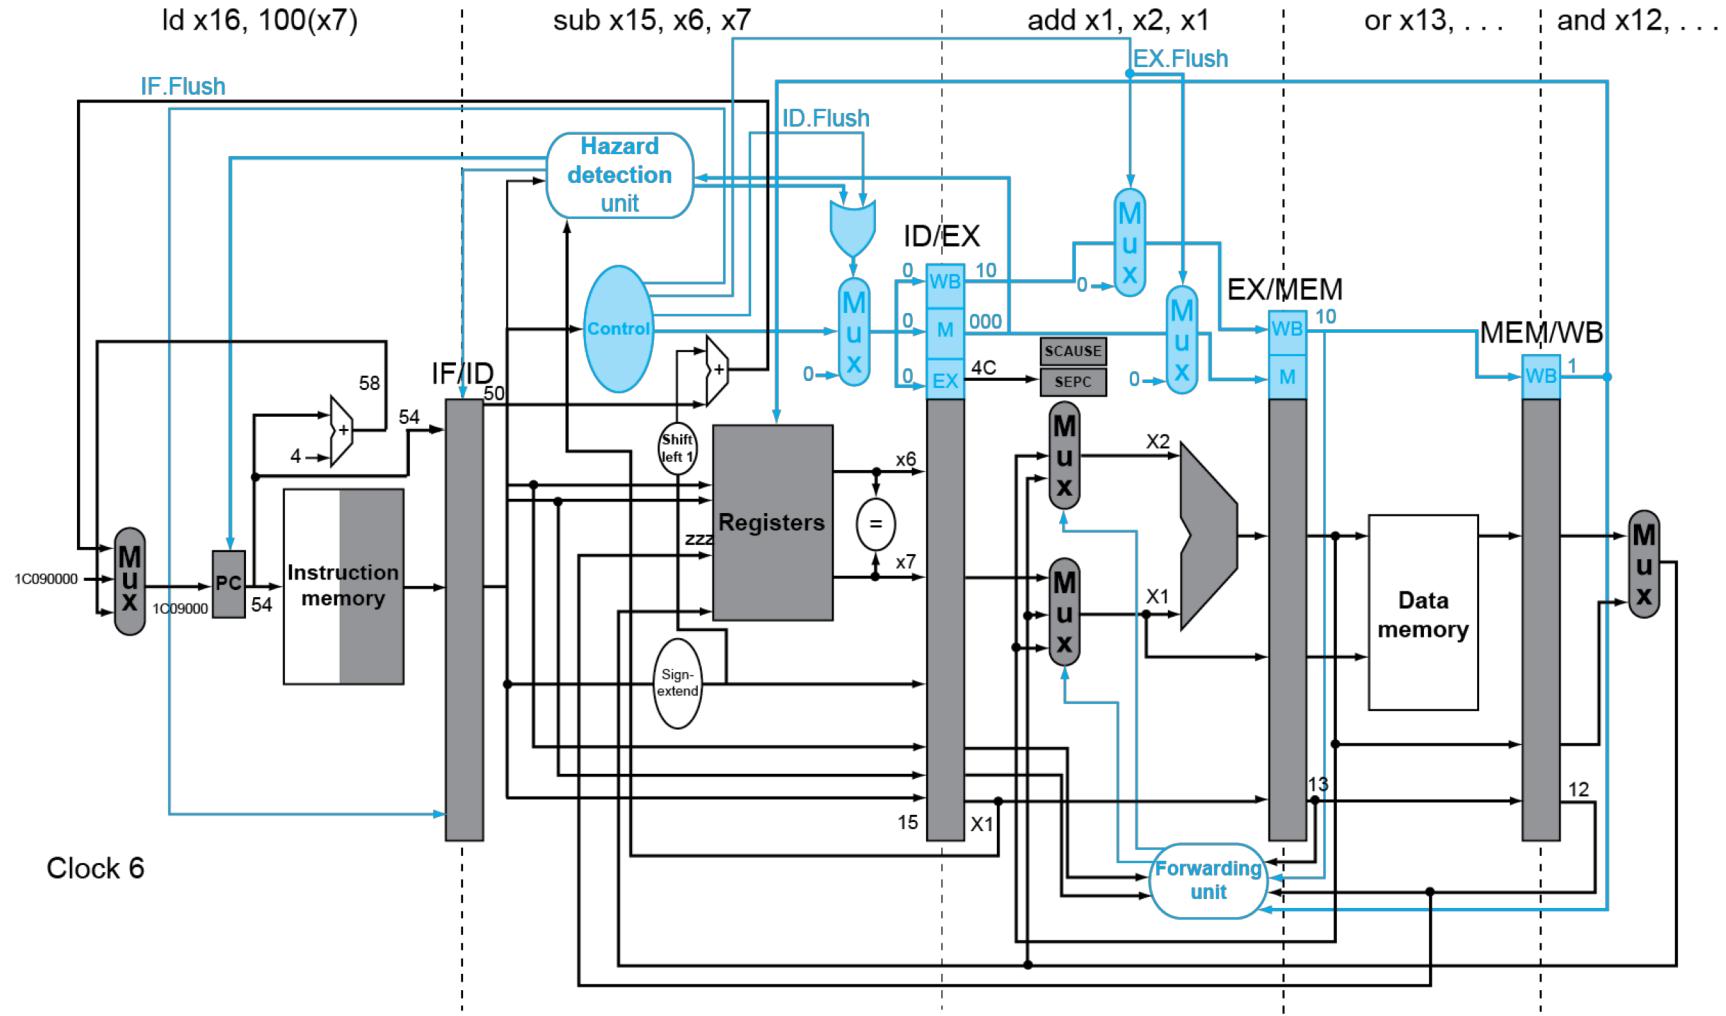
## □ Handler

```
1c090000      sd    x26, 1000(x10)  
1c090004      sd    x27, 1008(x10)
```

...

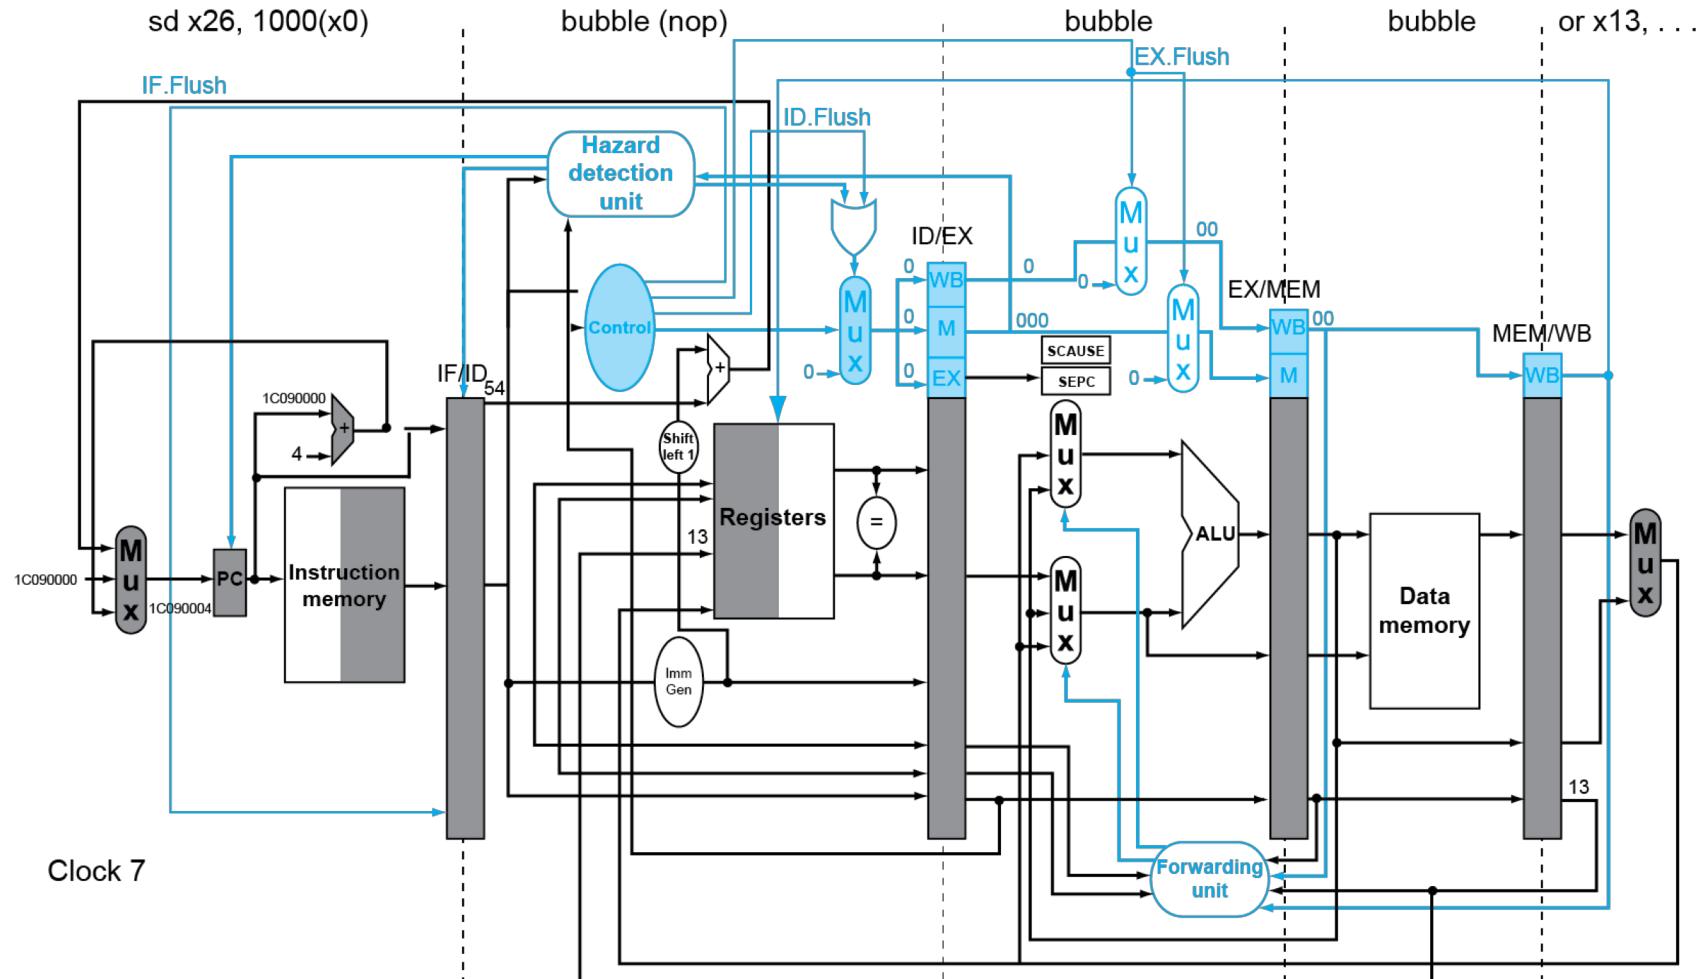
# Exception Example

94



# Exception Example

95



# Multiple Exceptions

96

- Pipelining overlaps multiple instructions
  - Could have multiple exceptions at once
- Simple approach: deal with exception from earliest instruction
  - Flush subsequent instructions
  - “Precise” exceptions
- In complex pipelines
  - Multiple instructions issued per cycle
  - Out-of-order completion
  - Maintaining precise exceptions is difficult!

# Imprecise Exceptions

97

- Just stop pipeline and save state
  - Including exception cause(s)
- Let the handler work out
  - Which instruction(s) had exceptions
  - Which to complete or flush
    - May require “manual” completion
- Simplifies hardware, but more complex handler software
- Not feasible for complex multiple-issue out-of-order pipelines

# Instruction-Level Parallelism (ILP)

98

- Pipelining: executing multiple instructions in parallel
- To increase ILP
  - Deeper pipeline
    - Less work per stage  $\Rightarrow$  shorter clock cycle
  - Multiple issue
    - Replicate pipeline stages  $\Rightarrow$  multiple pipelines
    - Start multiple instructions per clock cycle
    - CPI < 1, so use Instructions Per Cycle (IPC)
    - E.g., 4GHz 4-way multiple-issue
      - 16 BIPS, peak CPI = 0.25, peak IPC = 4
    - But dependencies reduce this in practice

# Multiple Issue

99

- Static multiple issue
  - ▣ Compiler groups instructions to be issued together
  - ▣ Packages them into “issue slots”
  - ▣ Compiler detects and avoids hazards
- Dynamic multiple issue
  - ▣ CPU examines instruction stream and chooses instructions to issue each cycle
  - ▣ Compiler can help by reordering instructions
  - ▣ CPU resolves hazards using advanced techniques at runtime

# Speculation

100

- “Guess” what to do with an instruction
  - Start operation as soon as possible
  - Check whether guess was right
    - If so, complete the operation
    - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
  - Speculate on branch outcome
    - Roll back if path taken is different
  - Speculate on load
    - Roll back if location is updated

# Compiler/Hardware Speculation

101

- Compiler can reorder instructions
  - e.g., move load before branch
  - Can include “fix-up” instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
  - Buffer results until it determines they are actually needed
  - Flush buffers on incorrect speculation

# Speculation and Exceptions

102

- What if exception occurs on a speculatively executed instruction?
  - ▣ e.g., speculative load before null-pointer check
- Static speculation
  - ▣ Can add ISA support for deferring exceptions
- Dynamic speculation
  - ▣ Can buffer exceptions until instruction completion (which may not occur)

# Static Multiple Issue

103

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
  - Specifies multiple concurrent operations
  - ⇒ Very Long Instruction Word (VLIW)

# Scheduling Static Multiple Issue

104

- Compiler must remove some/all hazards
  - Reorder instructions into issue packets
  - No dependencies with a packet
  - Possibly some dependencies between packets
    - Varies between ISAs; compiler must know!
  - Pad with nop if necessary

# RISC-V with Static Dual Issue

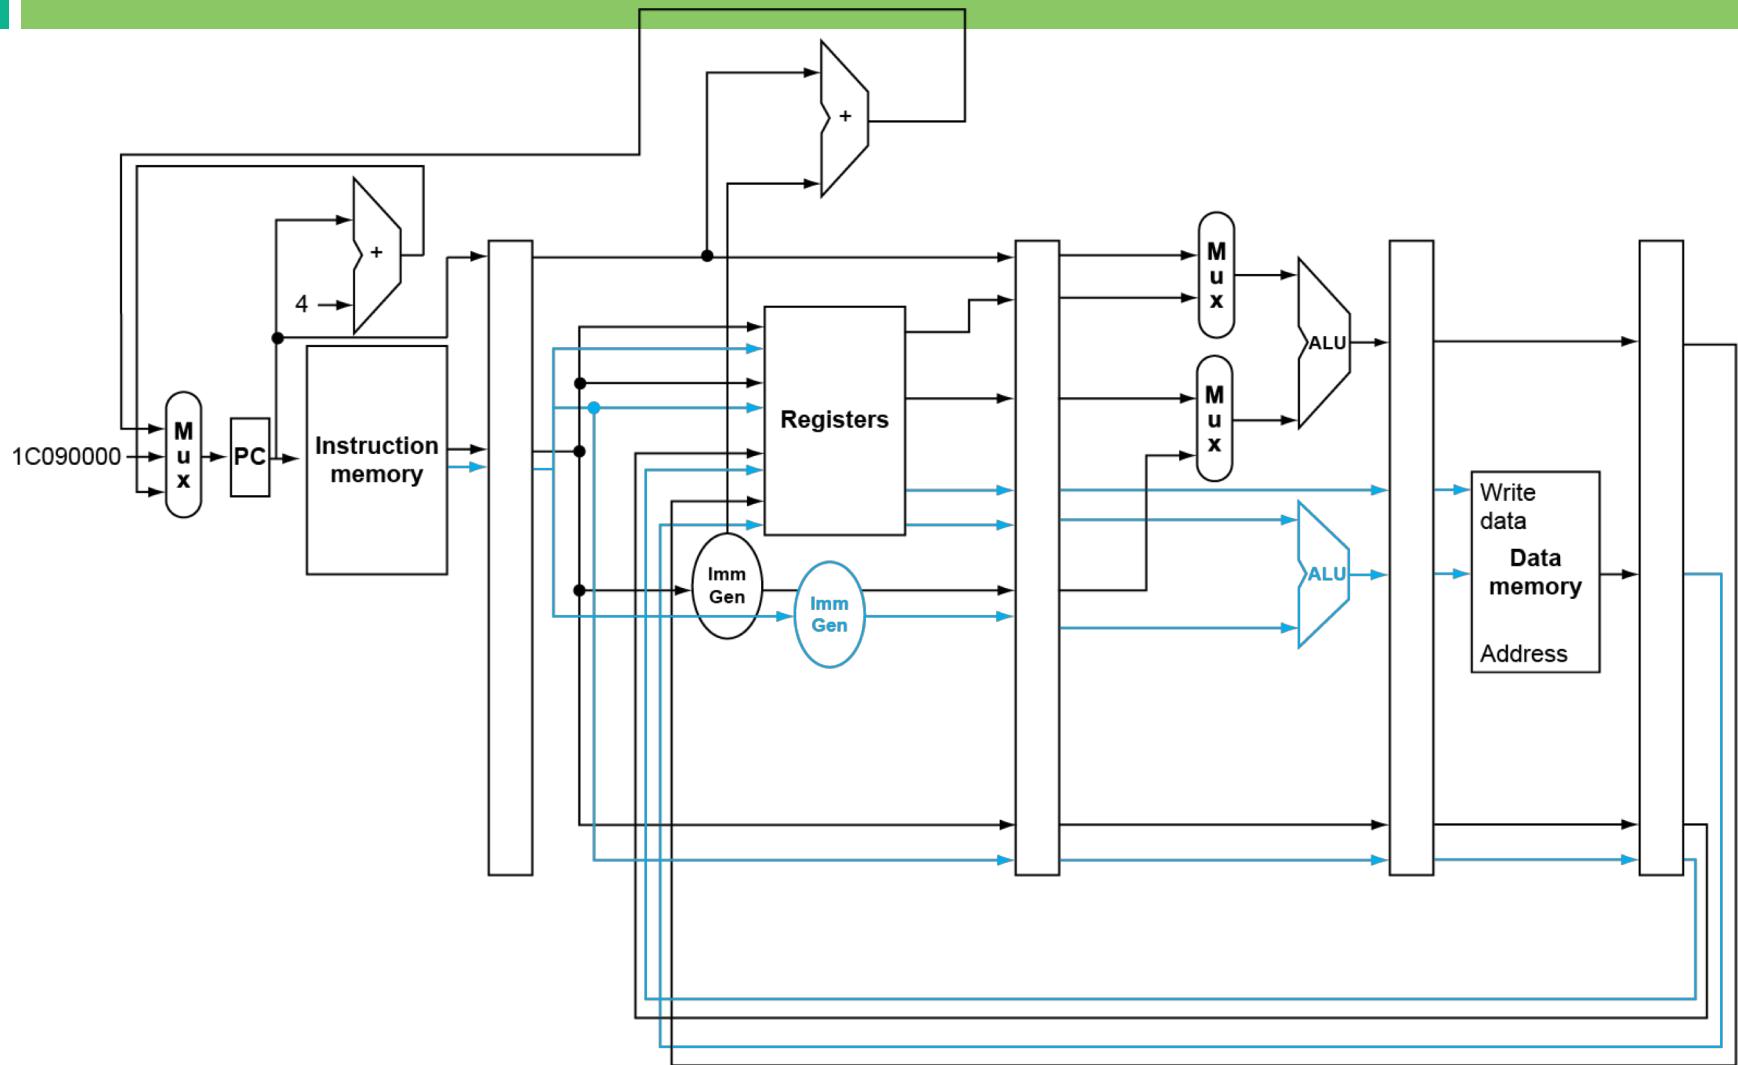
105

- Two-issue packets
  - One ALU/branch instruction
  - One load/store instruction
  - 64-bit aligned
    - ALU/branch, then load/store
    - ~~Pad an unused instruction with nop~~

Address	Instruction type	Pipeline Stages						
		IF	ID	EX	MEM	WB		
n	ALU/branch							
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store			IF	ID	EX	MEM	WB
n + 16	ALU/branch				IF	ID	EX	MEM
n + 20	Load/store				IF	ID	EX	WB

# RISC-V with Static Dual Issue

106



# Hazards in the Dual-Issue RISC-V

107

- More instructions executing in parallel
- EX data hazard
  - Forwarding avoided stalls with single-issue
  - Now can't use ALU result in load/store in same packet
    - add  $x10$ ,  $x0$ ,  $x1$   
 $1d x2, 0(x10)$
    - Split into two packets, effectively a stall
- Load-use hazard
  - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

# Scheduling Example

108

- Schedule this for dual-issue RISC-V

```
Loop: 1d  x31,0(x20)      // x31=array element
        add  x31,x31,x21    // add scalar in x21
        sd   x31,0(x20)      // store result
        addi x20,x20,-8       // decrement pointer
        blt x22,x20,Loop      // branch if x22 < x20
```

	ALU/branch	Load/store	cycle
Loop:	nop	1d  x31,0(x20)	1
	addi x20,x20,-8	nop	2
	add  x31,x31,x21	nop	3
	blt x22,x20,Loop	sd  x31,8(x20)	4

- IPC = 5/4 = 1.25 (c.f. peak IPC = 2)

# Loop Unrolling

109

- Replicate loop body to expose more parallelism
  - Reduces loop-control overhead
- Use different registers per replication
  - Called “register renaming”
  - Avoid loop-carried “anti-dependencies”
    - Store followed by a load of the same register
    - Aka “name dependence”
      - Reuse of a register name

# Loop Unrolling Example

110

- $\text{IPC} = 14/8 = 1.75$
- Closer to 2, but at cost of registers and code size

	ALU/branch	Load/store	cycle
Loop:	addi x20,x20,-32	1d x28, 0(x20)	1
	nop	1d x29, 24(x20)	2
	add x28,x28,x21	1d x30, 16(x20)	3
	add x29,x29,x21	1d x31, 8(x20)	4
	add x30,x30,x21	sd x28, 32(x20)	5
	add x31,x31,x21	sd x29, 24(x20)	6
	nop	sd x30, 16(x20)	7
	blt x22,x20,Loop	sd x31, 8(x20)	8

# Dynamic Multiple Issue

111

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
  - ▣ Avoiding structural and data hazards
- Avoids the need for compiler scheduling
  - ▣ Though it may still help
  - ▣ Code semantics ensured by the CPU

# Dynamic Pipeline Scheduling

112

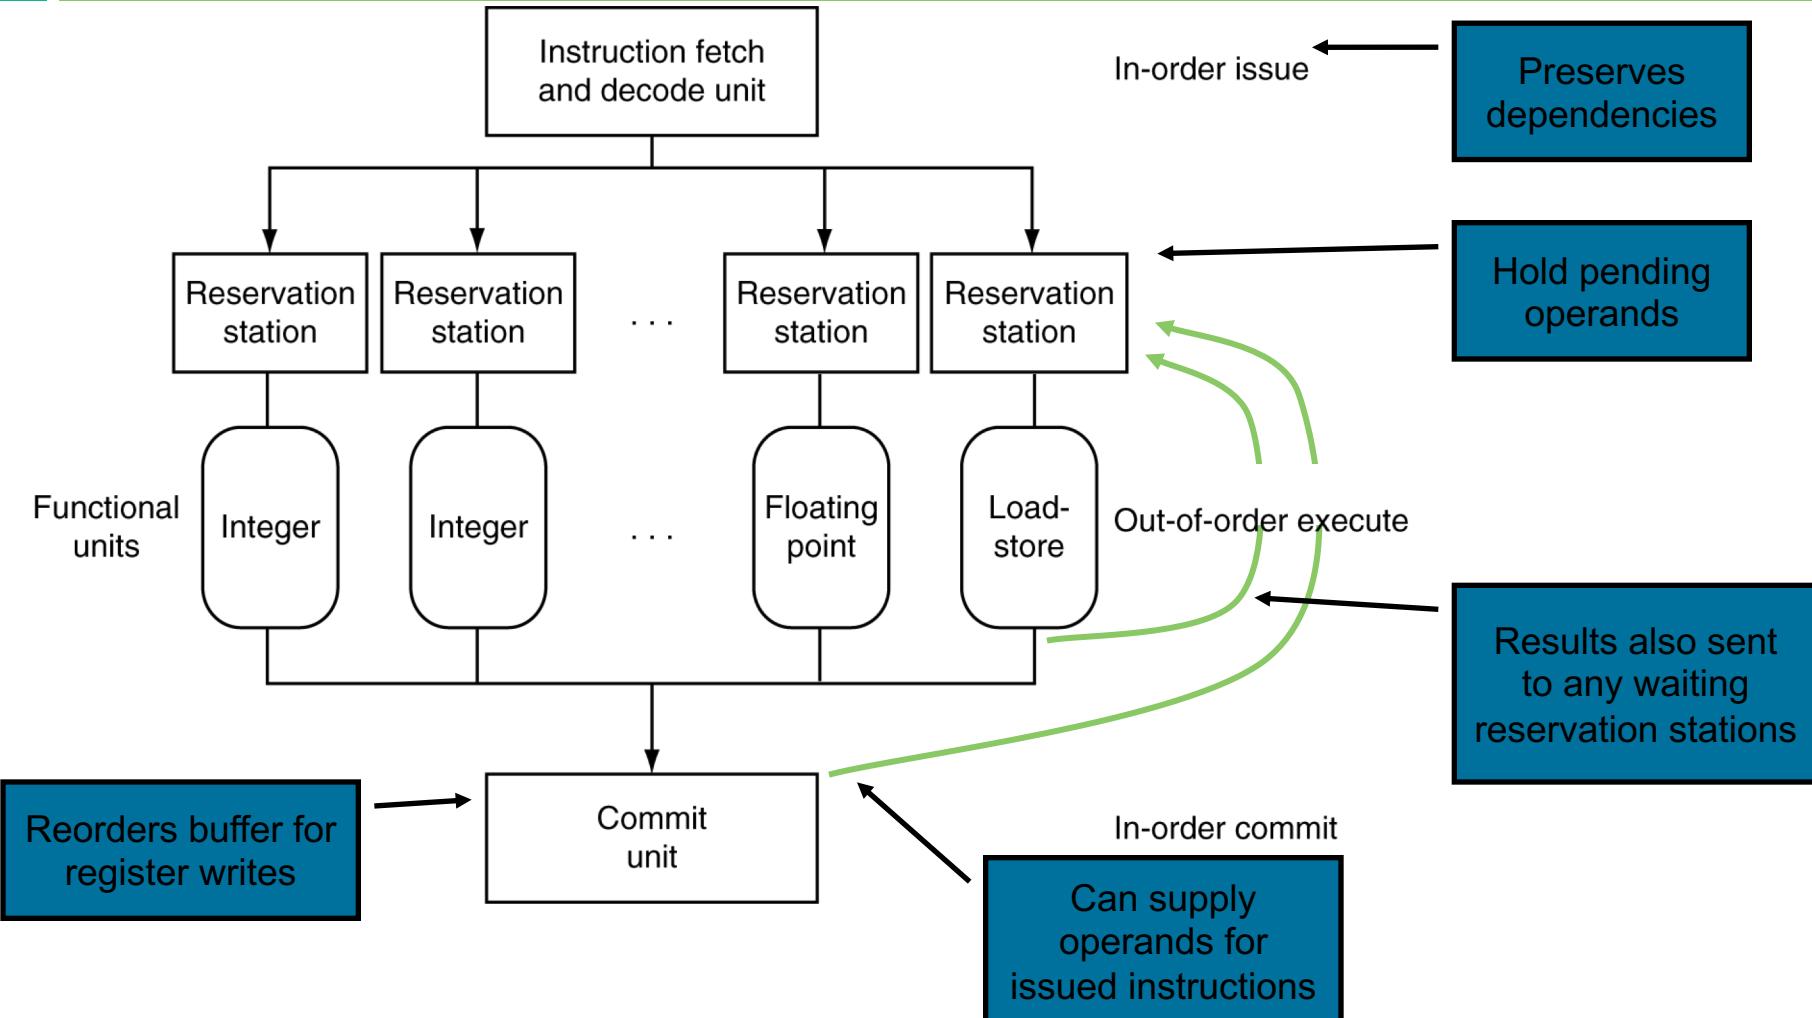
- Allow the CPU to execute instructions out of order to avoid stalls
  - But commit result to registers in order
- Example

```
ld    x31,20(x21)
add   x1,x31,x2
sub   x23,x23,x3
andi  x5,x23,20
```

- Can start sub while add is waiting for ld

# Dynamically Scheduled CPU

113



# Register Renaming

114

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
  - ▣ If operand is available in register file or reorder buffer
    - Copied to reservation station
    - No longer required in the register; can be overwritten
  - ▣ If operand is not yet available
    - It will be provided to the reservation station by a function unit
    - Register update may not be required

# Speculation

115

- Predict branch and continue issuing
  - ▣ Don't commit until branch outcome determined
- Load speculation
  - ▣ Avoid load and cache miss delay
    - Predict the effective address
    - Predict loaded value
    - Load before completing outstanding stores
    - Bypass stored values to load unit
  - ▣ Don't commit load until speculation cleared

# Why Do Dynamic Scheduling?

116

- Why not just let the compiler schedule code?
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

# Does Multiple Issue Work?

117

## The BIG Picture

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
  - ▣ e.g., pointer aliasing
- Some parallelism is hard to expose
  - ▣ Limited window size during instruction issue
- Memory delays and limited bandwidth
  - ▣ Hard to keep pipelines full
- Speculation can help if done well

# Power Efficiency

118

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

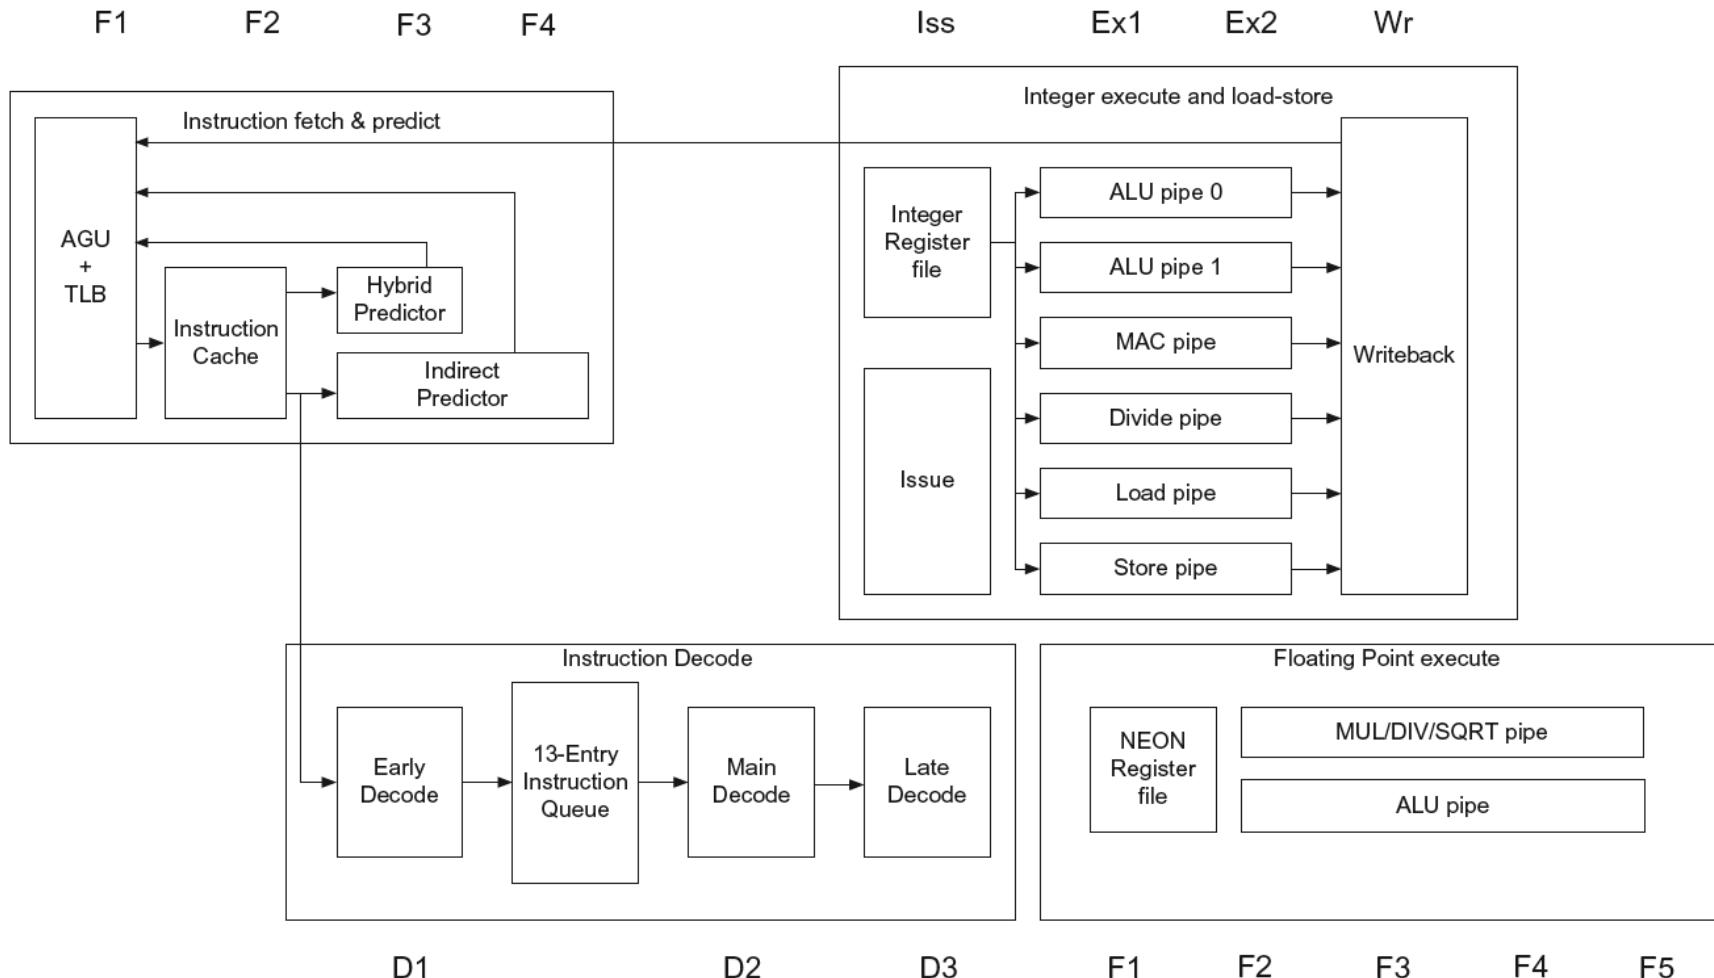
# Cortex A53 and Intel i7

119

Processor	ARM A53	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	100 milliWatts (1 core @ 1 GHz)	130 Watts
Clock rate	1.5 GHz	2.66 GHz
Cores/Chip	4 (configurable)	4
Floating point?	Yes	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	8	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	Hybrid	2-level
1 <sup>st</sup> level caches/core	16-64 KiB I, 16-64 KiB D	32 KiB I, 32 KiB D
2 <sup>nd</sup> level caches/core	128-2048 KiB	256 KiB (per core)
3 <sup>rd</sup> level caches (shared)	(platform dependent)	2-8 MB

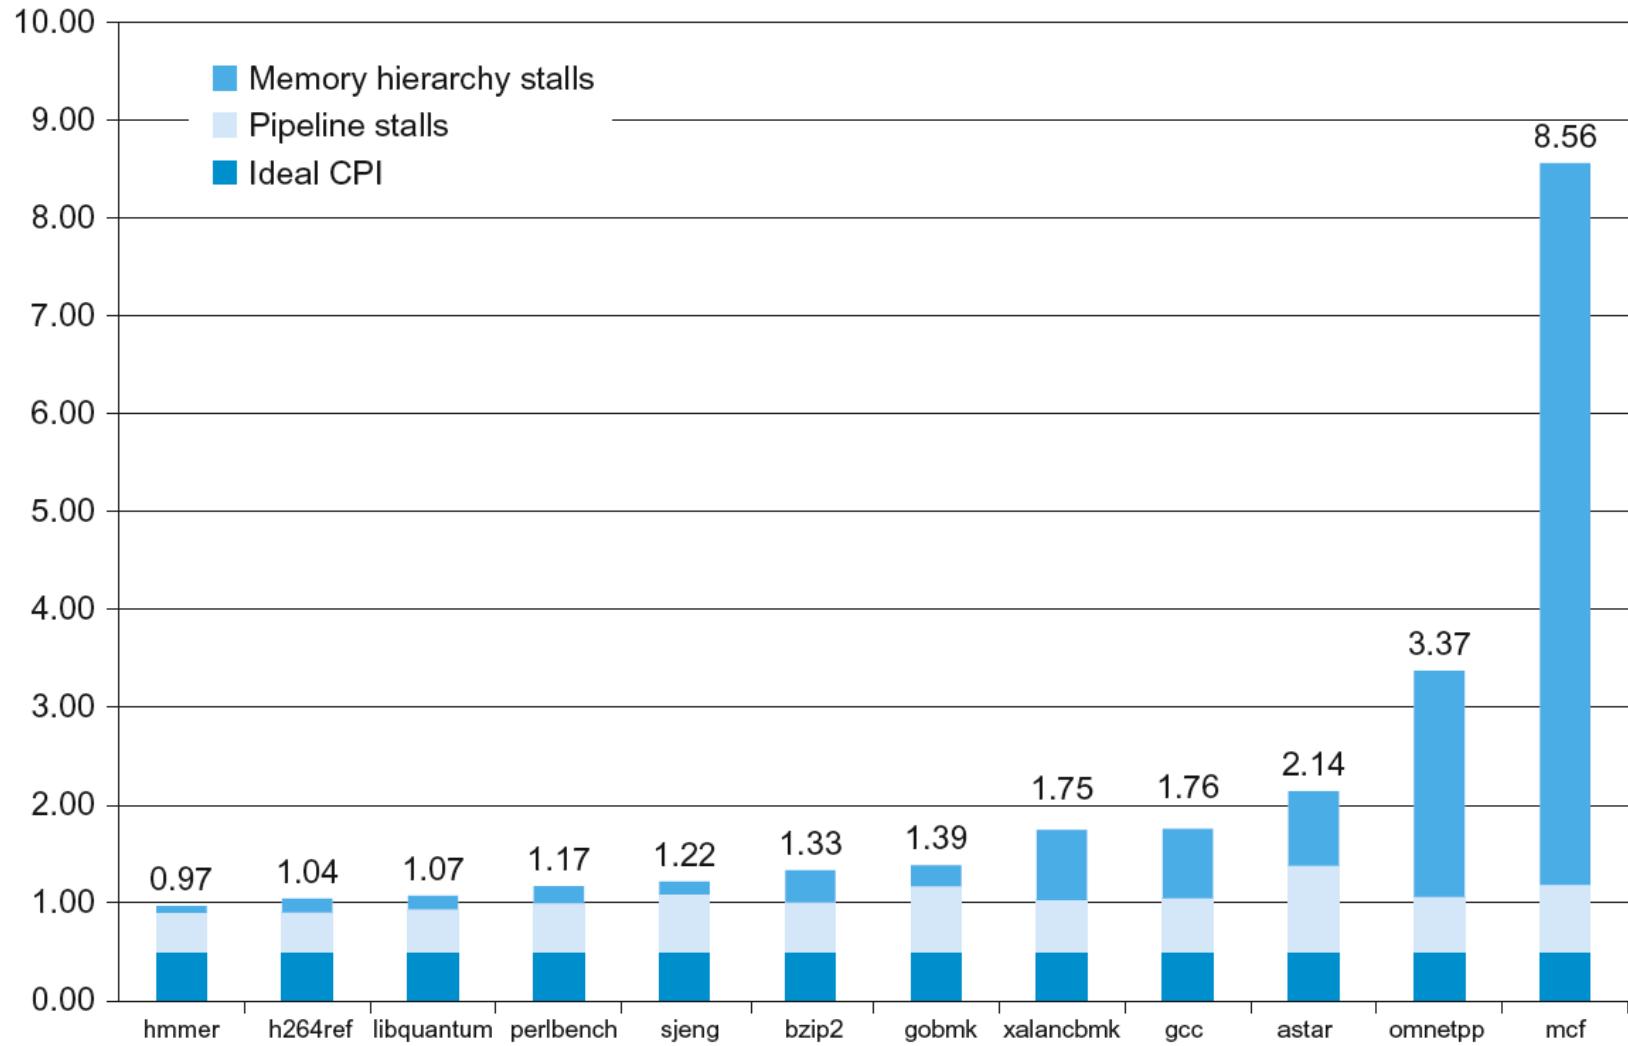
# ARM Cortex-A53 Pipeline

120



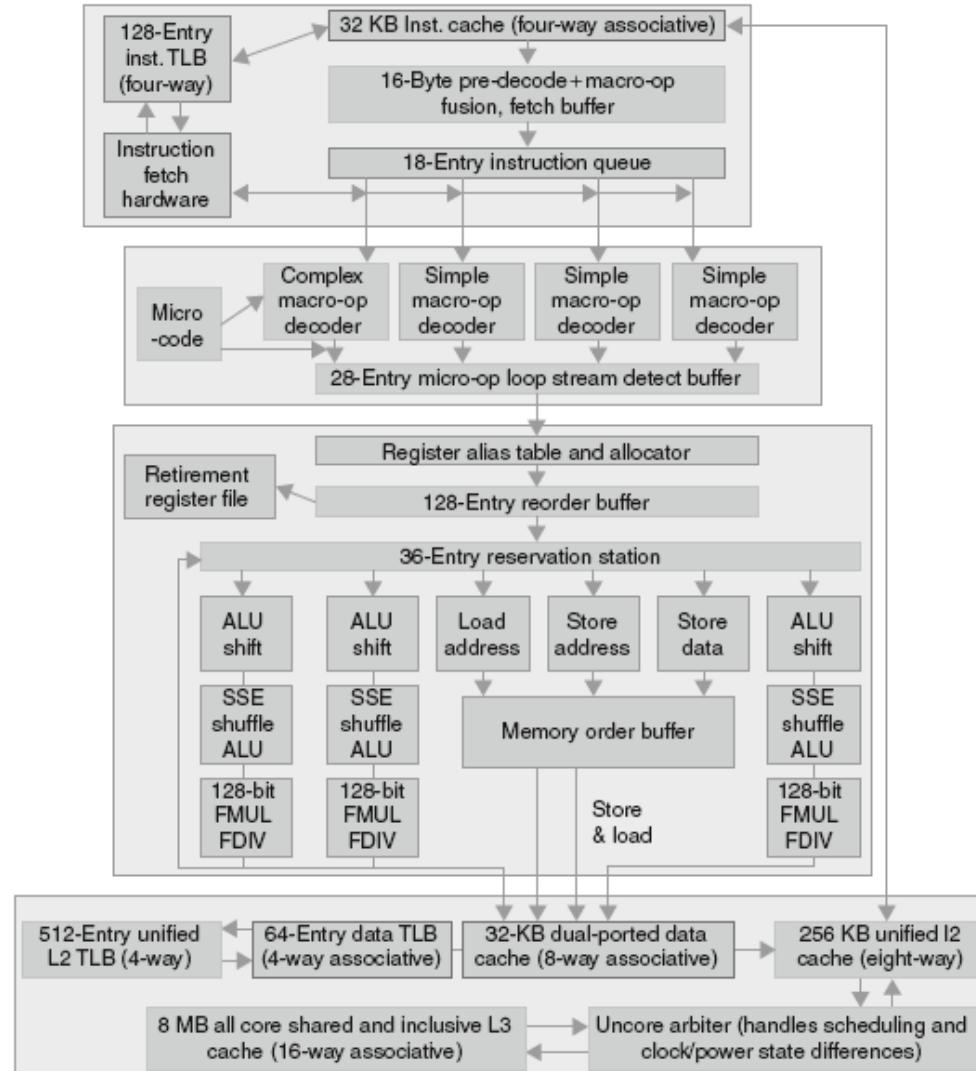
# ARM Cortex-A53 Performance

121



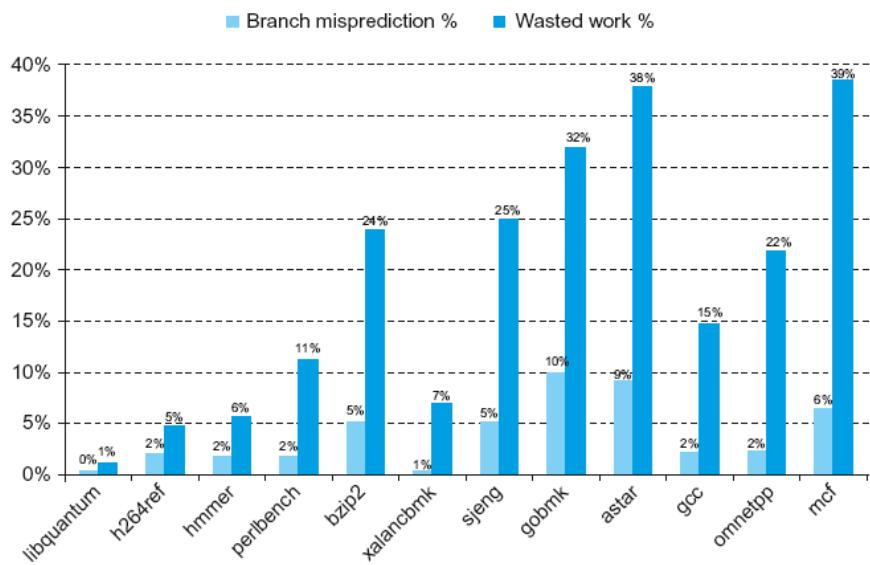
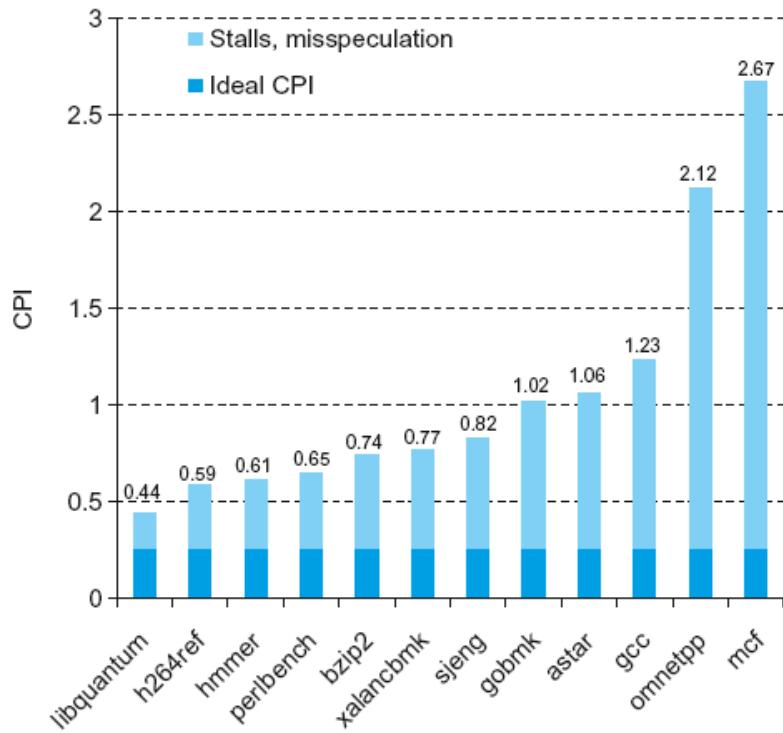
# Core i7 Pipeline

122



# Core i7 Performance

123



# Matrix Multiply

124

## □ Unrolled C code

```
1 #include <x86intrin.h>
2 #define UNROLL (4)
3
4 void dgemm (int n, double* A, double* B, double* C)
5 {
6     for ( int i = 0; i < n; i+=UNROLL*4 )
7         for ( int j = 0; j < n; j++ ) {
8             __m256d c[4];
9             for ( int x = 0; x < UNROLL; x++ )
10                c[x] = _mm256_load_pd(C+i+x*4+j*n);
11
12            for( int k = 0; k < n; k++ )
13            {
14                __m256d b = _mm256_broadcast_sd(B+k+j*n);
15                for (int x = 0; x < UNROLL; x++)
16                    c[x] = _mm256_add_pd(c[x],
17                                         _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
18            }
19
20            for ( int x = 0; x < UNROLL; x++ )
21                _mm256_store_pd(C+i+x*4+j*n, c[x]);
22        }
23 }
```

# Matrix Multiply

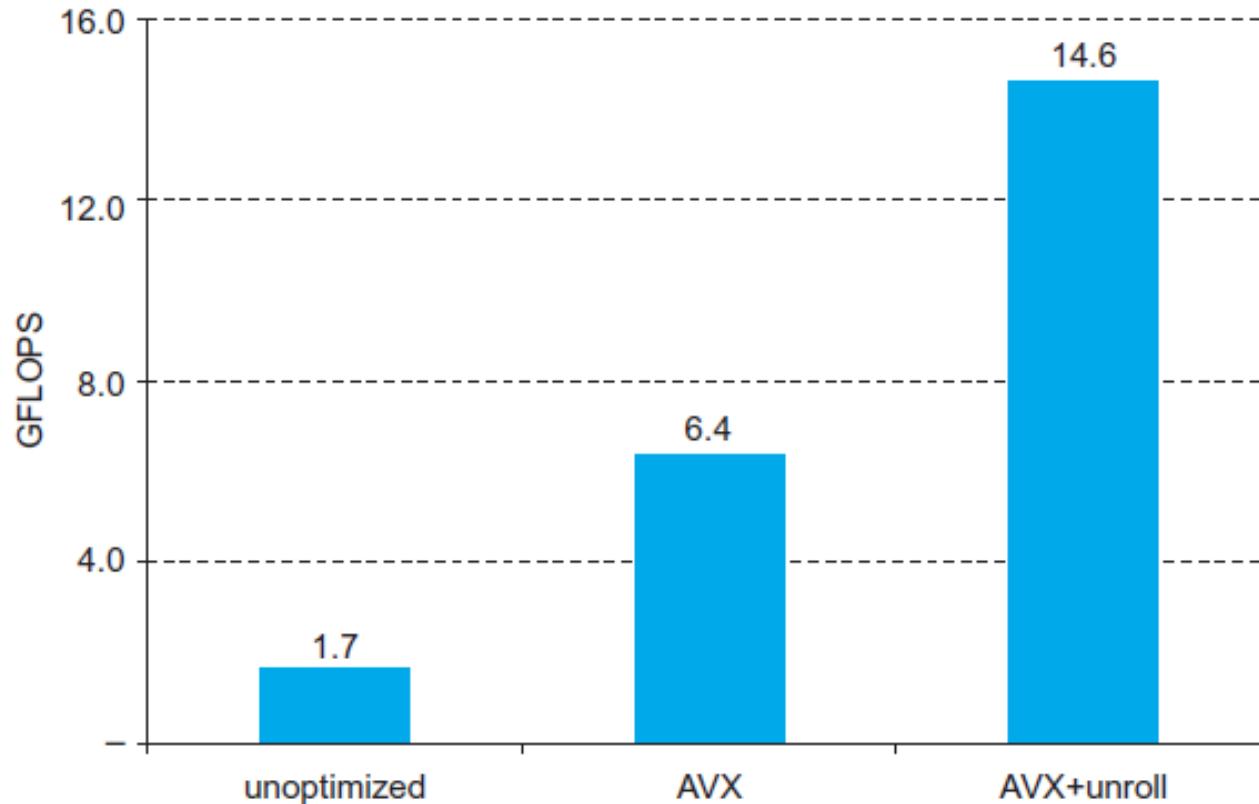
125

## □ Assembly code:

```
1 vmovapd (%r11),%ymm4          # Load 4 elements of C into %ymm4
2 mov %rbx,%rax                # register %rax = %rbx
3 xor %ecx,%ecx                # register %ecx = 0
4 vmovapd 0x20(%r11),%ymm3      # Load 4 elements of C into %ymm3
5 vmovapd 0x40(%r11),%ymm2      # Load 4 elements of C into %ymm2
6 vmovapd 0x60(%r11),%ymm1      # Load 4 elements of C into %ymm1
7 vbroadcastsd (%rcx,%r9,1),%ymm0 # Make 4 copies of B element
8 add $0x8,%rcx # register %rcx = %rcx + 8
9 vmulpd (%rax),%ymm0,%ymm5    # Parallel mul %ymm1,4 A elements
10 vaddpd %ymm5,%ymm4,%ymm4     # Parallel add %ymm5, %ymm4
11 vmulpd 0x20(%rax),%ymm0,%ymm5 # Parallel mul %ymm1,4 A elements
12 vaddpd %ymm5,%ymm3,%ymm3     # Parallel add %ymm5, %ymm3
13 vmulpd 0x40(%rax),%ymm0,%ymm5 # Parallel mul %ymm1,4 A elements
14 vmulpd 0x60(%rax),%ymm0,%ymm0 # Parallel mul %ymm1,4 A elements
15 add %r8,%rax                # register %rax = %rax + %r8
16 cmp %r10,%rcx                # compare %r8 to %rax
17 vaddpd %ymm5,%ymm2,%ymm2     # Parallel add %ymm5, %ymm2
18 vaddpd %ymm0,%ymm1,%ymm1     # Parallel add %ymm0, %ymm1
19 jne 68 <dgemm+0x68>         # jump if not %r8 != %rax
20 add $0x1,%esi                # register %esi = %esi + 1
21 vmovapd %ymm4,(%r11)          # Store %ymm4 into 4 C elements
22 vmovapd %ymm3,0x20(%r11)      # Store %ymm3 into 4 C elements
23 vmovapd %ymm2,0x40(%r11)      # Store %ymm2 into 4 C elements
24 vmovapd %ymm1,0x60(%r11)      # Store %ymm1 into 4 C elements
```

# Performance Impact

126



# Fallacies

127

- Pipelining is easy (!)
  - The basic idea is easy
  - The devil is in the details
    - e.g., detecting data hazards
- Pipelining is independent of technology
  - So why haven't we always done pipelining?
  - More transistors make more advanced techniques feasible
  - Pipeline-related ISA design needs to take account of technology trends
    - e.g., predicated instructions

# Pitfalls

128

- Poor ISA design can make pipelining harder
  - ▣ e.g., complex instruction sets (VAX, IA-32)
    - Significant overhead to make pipelining work
    - IA-32 micro-op approach
  - ▣ e.g., complex addressing modes
    - Register update side effects, memory indirection
  - ▣ e.g., delayed branches
    - Advanced pipelines have long delay slots

# Concluding Remarks

129

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
  - ▣ More instructions completed per second
  - ▣ Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
  - ▣ Dependencies limit achievable parallelism
  - ▣ Complexity leads to the power wall